

*Prof. O. Hilliges**(Advanced Interactive Technologies Lab)**Prof. S. Coros (Computational Robotics Lab)**A. Cohen, J. Song, B. Hepp**V. Megaro, R. Roveri, J. Bern, L. Helminger**C. Öztireli***Final Exam**

?? February 2018

First and Last name: _____

ETH number: _____

Signature: _____

General Remarks

- At first, please check that your exam questionnaire is complete (there are 26 pages total).
- Remove all material from your desk which is not allowed by examination regulations.
- Fill in your first and last name and your ETH number and sign the exam. Place your student ID in front of you.
- You have 3 hours for the exam. There are 10 questions, where you can earn a total of 180 points. You don't have to score all points to obtain a very good grade.
- Start each question on a separate sheet. Put your name and ETH number on top of each sheet. Only write on the question sheet where explicitly stated.
- Please do not use a pencil or red color pen to write your answers.
- You may provide at most one valid answer per question. Invalid solutions must be canceled out clearly.

	Topic	Max. Points	Points Achieved	Visum
1	Principal Component Analysis	15		
2	Filtering	20		
3	Fourier Transform	15		
4	Optical Flow	25		
5	Miscellaneous	15		
6	OpenGL and Rendering	30		
7	Light, Color and Ray Tracing	24		
8	Transformations	22		
9	Rigid Bodies	12		
Total		180		

Grade:

Question 1: PCA (15 pts.)

- a) Given a raw data matrix \mathbf{D} , what are the pre-processing steps you have to perform before applying Principal Component Analysis (PCA)? **2 pts.**

[illegible]

Center the data. **1 pts.**

Normalize the data. **1 pts.**

[illegible]

- b) Given a properly pre-processed data matrix \mathbf{D} . How do you compute the PCA? **2 pts.**

[illegible]

Compute the covariance matrix Σ of \mathbf{D} . 1 pts.

Perform an Eigen Decomposition of Σ . 1 pts.

The resulting eigen vectors give directions that describe the variation in the data independent of each other.

The corresponding eigenvalue describes how large the variation along this direction is. 1
pts.

Note: 0.5 pts if only eigenvector or only eigenvalue is mentioned.

[illegible]

- c) When computing a PCA, the resulting components are orthogonal to each other and the corresponding eigenvalues are real. Explain why this is always the case. **2 pts.**

Question 2: Filtering (20 pts.)

- a) Given a high-pass filter, we can get the low-pass filtered image by first filtering it using the high-pass filter and subtracting it from the original image.

Alternatively, following the same idea, we can generate a low-pass filter kernel from a high-pass filter kernel.

Compute a low-pass filter kernel using the high-pass filter kernel given below. **3 pts.**

-0.07	-0.12	-0.07
-0.12	0.76	-0.12
-0.07	-0.12	-0.07

[illegible]

<i>Impulse kernel (image)</i>	<i>High-pass kernel</i>	<i>Low-pass kernel</i>
$ > m4ex > m4ex > m4ex $ <i>height0</i>	0	0
0	1	0
0	0	0

i m4ex

i m4ex

i m4ex

$$\begin{array}{ccc} -0.07 & -0.12 & -0.07 \\ -0.12 & 0.76 & -0.12 \\ -0.07 & -0.12 & -0.07 \end{array}$$

i m5ex

i m5ex

	0.07	0.12	0.07
	0.11	0.24	0.11
	0.07	0.12	0.07

2 pts.

$$G2 = \begin{bmatrix} -1/6 & 0 & 1/6 \\ -1/6 & 0 & 1/6 \\ -1/6 & 0 & 1/6 \end{bmatrix}.$$

Question 3: Fourier transform (15 pts.)

a) The fourier transform of an infinite continuous signal g is defined as

$$F(g)(u) = \int_{\mathbb{R}} g(x) e^{-j2\pi ux} dx \quad (1)$$

The inverse fourier transform is defined as

$$g(x) = \int_{\mathbb{R}} e^{j2\pi ux} F(g)(u) du \quad (2)$$

The fourier transform of the dirac delta function is given as

$$F(\delta(x - x_0))(u) = e^{-j2\pi ux_0} \quad (3)$$

Euler's formula states

$$e^{jx} = \cos(x) + j \sin(x) \quad (4)$$

If q is a bidimensional continuous signal, the fourier transform is defined as

$$F(g)(u, v) = \int \int_{\mathbb{R}^2} g(x, y) e^{-j2\pi(ux+vy)} dx dy \quad (5)$$

Similarly the fourier transform of 2D discrete and finite signals (i.e. images) can be defined by using zero padding and summation instead of integration as:

$$F(g)(u, v) = \frac{1}{N} \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} g(x, y) e^{-j2\pi(\frac{ux+vy}{N})} \quad (6)$$

- i) Derive the Fourier Transform of the infinite continuous signal $f_1(x) = \cos(2\pi k_0 x)$ using Equation 1 (Your solution must not contain any integrals). **4 pts.**
- ii) Given a function $f_2(x)$ whose Fourier Transform is F_2 and a function $f_3(x)$ which is the product of $f_1(x)$ and $f_2(x)$, compute the Fourier Transform of $f_3(x)$ in terms of F_2 . **2 pt.**

[illegible][illegible]

- b) You are given a generic normalized low pass filter kernel. Which filter would be obtained by subtracting 1 from its central element? Explain your answer. **4 pts.** ANSWER

Let f be a generic low pass filter and a be a generic image, the result of convolving the new filter $(f - \delta)$ with a is:

$$(f - \delta) * a = f * a - \delta * a = -(a - f * a)$$

To derive this expression, we need to use the linearity properties of the convolution and the fact that the convolution with a delta does not alter the original signal. In the end, it results in a high pass filter which inverts the sign of the image (i.e., its phase is delayed by 180 degrees). (high pass filter 1pt) (incomplete 2pts) (some errors 3pts) (inverted high pass filter + procedure 4pts)

[illegible]

- f) Explain how you would use optical flow for image stabilization. What other application can you think of for optical flow? **3 pts.**

[illegible]

Estimate the flow between the frames, and warp the image using the same flow over all the pixels so that the flow is close 0. (2 points) Name another application (e.g. video compression, slow motion, etc) (1 point)

[illegible]

Question 5: Miscellaneous (15 pts)

- a) What are the three properties a feature detector should have in order to be robust? Is the Harris Corner detector robust to all three of them? **3 pts.** ANSWER ANSWER

1pt per item.

- *Invariant to brightness change*
- *Invariant to shift and rotation*
- *Invariant to scale (this is where Harris Corner detector fails)*

ANSWER

- b) How would you exploit temporal redundancy for video compression? Describe the three types of frames used in this method. **4 pts.** ANSWERANSWERANSWERANSWERANSWER

1pt: The idea is to predict the current frame based on previously coded frames: Temporal redundancy reduction 3pts: 3 types of frames:

- *I-frame: intra-coded frame, coded independently of all other frames.*
- *P-frame: coded frame, coded based on previously coded frame (based on previous I and P frames)*
- *B-frame: Bi-directionally predicted frame, coded based on both previous and future coded frames.*

ANSWER

- c) When may temporal redundancy reduction be ineffective? **2 pts.** ANSWER ANSWER

When there are many scene changes (1pt) or/and high motion (1pt)

ANSWER

- d) Which technique can be used when temporal redundancy reduction fails? Briefly describe the main 2 steps of the practical approach for this technique. **4 pts.**

The technique is motion-compensated prediction (1pt.) Practical approach: Block-matching Motion Estimation. Partition each frame into blocks and describe motion of each block by finding the best matching block in the reference frame.(3pts)

ANSWER

- e) Name the two disadvantages of the block-matching algorithm. **2 pts.** ANSWER

1pt per disadvantage. 1) It assumes translational motion model which breaks down for more complex motion. 2) often produces blocking artifacts.

ANSWER

Question 6: OpenGL and Rendering (30 pts.)

a) General Questions

- i) Two objects, A and B, are rendered in OpenGL. A is placed in front of B, and is not fully opaque. However, B is not visible through A.

Give two possible explanations for this apparent bug.

2 pts. ANSWER ANSWER

- *blending not enabled*
- *objects drawn in wrong order.*

ANSWER ANSWER

- ii) Given the following CPP code:

```
glEnable(GL_BLEND);  
glBlendEquation(GL_FUNC_ADD);  
glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA);
```

An object A with color $c_A = (1.0, 0.0, 0.0, 1.0)$ was already rendered. A second object B with color $c_B = (0.0, 1.0, 0.0, 0.6)$ is now placed in front of A. What is the RGB value of the overlapping region.

2 pts. ANSWER ANSWER ANSWER

final = (0.4, 0.6, 0.0)

ANSWER ANSWER

- iii) Why does z-fighting happen?

1 pts. ANSWER ANSWER ANSWER ANSWER ANSWER

Because two objects have the same value in the depth buffer. / floating point unprecision

ANSWER ANSWER

- iv) Name 3 stages of the graphics pipeline and order them according to the step they occur.

1 pts. ANSWER

Vertex shader, rasterization, Fragment Shader

ANSWER ANSWER

- v) Name two implicit and two explicit representations in geometry. In addition, name two pros and two cons of implicit representations.

3 pts. ANSWER ANSWER ANSWER

Implicit: algebraic surfaces, constructive solid geometry, level set methods, blobby surfaces, fractals

Explicit: triangle meshes, polygon meshes, subdivision surfaces, NURBS, point clouds

Implicit: Pros: description can be very compact (e.g., a polynomial), easy to determine if a point is inside/outside (just plug it in!), other queries may also be easy (e.g., distance to surface), for simple shapes, exact description/no sampling error, easy to handle changes in topology (e.g., fluid). Cons: expensive to find all points in the shape (e.g., for drawing) very difficult to model complex shapes

b) Vertex Shader

We consider the following vertex shader.

```
#version 440

uniform mat4 projectionMatrix;
uniform mat4 modelviewMatrix;
uniform mat3 normalMatrix;

in vec3 inPosition;
in vec3 inNormal;

out vec3 outPosition;
out vec3 outNormal;

void main()
{
    outPosition = inPosition;
    outNormal = normalMatrix * inNormal;

    gl_Position = projectionMatrix * modelviewMatrix * vec4(inPosition, 1.0f);
}
```

- i) What is the purpose of `projectionMatrix` and `modelviewMatrix`? 4 pts.

2pts for each.

[illegible]

- ii) The input normal is transformed using `normalMatrix`, unlike the input position. Why is that? **2 pts.**

[illegible]

Because if the modelviewMatrix contains shear/scale, the normal doesn't remain orthogonal.

[illegible]

c) Fragment Shader

- i) Consider the following fragment shader. Complete the missing lines of code to implement the Phong reflection model, where I_a is the ambient color, I_p is the color for the diffuse and specular components, k_a , k_d and k_s are the material parameters for the ambient, diffuse and specular component, respectively, and α is the exponent of the specular component. The geometry vectors are shown in Figure 2. **8 pts.**




Figure 3: Geometry vectors for the reflection model.

```
#version 440

in vec3 outNormal;
in vec3 outPosition;

uniform vec3 eyePosition;
uniform vec3 lightDir;

out vec4 gl_FragColor;

void main()
{
    vec3 normal = normalize(outNormal);
    lightDir = normalize(lightDir);

    vec3 Ia = vec3(1.0, 0.0, 0.0);
    vec3 Ip = vec3(1.0, 1.0, 1.0);

    float ka = 0.2;
    float kd = 0.2;
    float ks = 0.6;
    float alpha = 2.0;

    // Ambient
    vec3 ambient = . . . . .

    // Diffuse
    vec3 diffuse = . . . . .

    // Specular
    vec3 R = . . . . .
    vec3 V = . . . . .

    R = normalize(R);
    V = normalize(V);

    vec3 specular = . . . . .

    vec3 color = . . . . .

    gl_FragColor = vec4(color, 1.0);
}
```

ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER ANSWER

```
#version 440

in vec3 outNormal;
in vec3 outPosition;
in vec3 eyePosition;

out vec4 gl_FragColor;

void main()
{
```


- ii) Given are two possible values for α ($\alpha = 2.0$ and $\alpha = 20.0$), and two possible vectors I_a ($I_a = (0.0, 0.0, 1.0)$ and $I_a = (1.0, 0.0, 0.0)$). For each subfigure (a,b,c,d) in Figure 3 obtained with the above fragment shader, decide which values for α and I_a were used. **2 pts.**

a) $\alpha = \dots\dots I_a = \dots\dots\dots$

b) $\alpha = \dots\dots I_a = \dots\dots\dots$

c) $\alpha = \dots\dots I_a = \dots\dots\dots$

d) $\alpha = \dots\dots I_a = \dots\dots\dots$




Figure 4: Different configurations of the phong shading model.

ANSWER ANSWER

- a) $\alpha = 2 \ I_a = (1.0, 0.0, 0.0)$
b) $\alpha = 20 \ I_a = (1.0, 0.0, 0.0)$
c) $\alpha = 2 \ I_a = (0.0, 0.0, 1.0)$
d) $\alpha = 20 \ I_a = (0.0, 0.0, 1.0)$

ANSWER ANSWER

- iii) The Blinn-Phong reflection model is a modification to the Phong reflection model. The Blinn-Phong model is the same as the Phong model, except for the specular component which is computed as $(H \cdot N)^{\alpha'}$, where H is the normalized halfway vector between light and viewing direction, as shown in Figure 2. In this exercise, we assume that $\alpha' = \alpha$. Complete the following lines of code to compute the specular component of the Blinn-Phong reflection model. **5 pts.**

```
vec3 H = .....
.....
vec3 specular = .....
```

[illegible]

```
vec3 halfDir = normalize(lightDir + eyePosition);
float specAngle = max(dot(halfDir, normal), 0.0);
vec3 specular = vec3(pow(specAngle, alpha));
```

[illegible]

Question 7: Light, Color and Ray Tracing (24 pts.)

- a) Why, in the YIQ color space, less bandwidth is required for Q than for I? Where is the YIQ color applied? **3 pts.**

[illegible]

The YIQ system is intended to take advantage of human color-response characteristics. The eye is more sensitive to changes in the orange-blue (I) range than in the purple-green range (Q). Therefore less bandwidth is required for Q than for I. YIQ is the color space used by the NTSC color TV system.

[illegible]

- b) Explain a main difference between the RGB color space and the CMY color space. **2 pts.**

[illegible]

RGB is additive, while CMY is subtractive.


[illegible]

Figure 5: Color Matching Functions results from the CIE Experiment.

- c) In 1931 CIE performed an experiment, where the observers had to trim the brightness of each of three primary beams in order to match a reference color. Figure 4 shows the resulting color matching functions, describing how the primaries were mixed. Why does one of them have some negative values? **2 pts.**

[illegible]

Not all test colors could be matched using the original experiment technique. When this was the case, a variable amount of one of the primaries could be added to the test color, and a match with the remaining two primaries was carried out with the variable color spot. For these cases, the amount of the primary added to the test color was considered to be a negative value. In this way, the entire range of human color perception could be covered.

[illegible]

d) Consider the 3 following primaries given in the CIE xyY color space:

	x	y	Y
c_1	0.3	0.2	20
c_2	0.5	0.25	10
c_3	0.35	0.5	10

Figure 5 is an empty CIE xy Chromaticity diagram.

Plot c1, c2 and c3 on the diagram.

2 pts.

[illegible]

Simply draw the x and y coordinates of the three points in the CIE plot.


[illegible]

Figure 6: CIE xy Chromaticity diagram

i) What is an advantage and a disadvantage of an octree compared to a K-D tree? **2**
pts.

Question 8: Transformations (XX pts.)

a) Transformations → Drawings


We will be working with the following homogeneous transformation matrices.

Note: This problem is in 2D, so a point $\begin{pmatrix} x \\ y \end{pmatrix}$ in Cartesian space would be written $\begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$ in homogeneous coordinates.

$$\mathbf{T} = \begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

$$\mathbf{R} = \begin{pmatrix} \cos 120^\circ & -\sin 120^\circ & 0 \\ \sin 120^\circ & \cos 120^\circ & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

Consider the following equilateral triangle, with vertices labeled with arcs.




In each of the following questions, you will be given some transformation \mathbf{A} .

Please draw the result of applying \mathbf{A} to the original triangle. Drawings do **not** need to be precise, just a rough sketch is perfect.

Note: To be clear, if \mathbf{p} is some point on the triangle written in homogeneous coordinates, then the transformed point would be \mathbf{Ap} .

Example Problem: $\mathbf{A} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$


Example Solution:



- | | |
|---------------------------------|--------|
| i) $\mathbf{A} = \mathbf{T}$ | 1 pts. |
| ii) $\mathbf{A} = \mathbf{R}$ | 1 pts. |
| iii) $\mathbf{A} = \mathbf{TR}$ | 1 pts. |
| iv) $\mathbf{A} = \mathbf{RT}$ | 1 pts. |

b) **Drawings** \rightarrow **Transformations**

Now consider the following square, again with vertices labeled with arcs.




In each of the following questions, you will be given a picture. This picture is the result of applying some transformation \mathbf{B} to the original square.

Note: To be clear, if \mathbf{q} is some point on the square written in homogeneous coordinates, then the transformed point would be \mathbf{Bq} .

Please write down \mathbf{B} for each picture. Your answers do not need to be in any particular form, the product of a few matrices is fine, as is a single matrix. As long as your answer would produce the given figure, it is perfect.


Example Problem:



Example Solution: $\mathbf{B} = \begin{pmatrix} 1 & 0 & 0.5 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$


i) What is **B**?

1 pts.




ii) What is **B**?

1 pts.



iii) What is **B**?


1 pts.



Question 9: Rigid Bodies (XX pts.)

a) Rigid Body Transforms


i) Is this a rigid body transformation?



- ☐ Yes
☐ No

1 pt.


ii) Is this a rigid body transformation?



- ☐ Yes
☐ No

1 pt.


iii) Is this a rigid body transformation?



- ☐ Yes
☐ No

1 pt.


iv) Is this a rigid body transformation?



- ☐ Yes
- ☐ No

1 pt.

v) Is this a rigid body transformation?



- ☐ Yes
- ☐ No

1 pt.

b) **Bouncing Balls**

We consider a very simple collision example.

Consider two spheres.

They both have radius $1m$ (one meter).

The spheres are constrained to the x -axis.

They both have the same mass.

We neglect gravity, air resistance, etc.

Sphere A has initial center position $x_0^A = -0.5m$.

Sphere B has initial center position $x_0^B = 0.5m$.

Sphere A has initial velocity $\dot{x}_0^A = 1.0m/s$.

Sphere B has initial velocity $\dot{x}_0^B = -1.0m/s$

Notation: x_k^A means the position of Sphere A at timestep k .

- i) Draw a 2D diagram of the initial state of the system. Include arrows for velocities. Don't forget to draw the origin and x -axis. **1 pts.**
- ii) Given sphere positions x^A, x^B , write down a condition for whether they are currently colliding (intersecting). **1 pts.**
- iii) Evolve the system forward in time one step using forward Euler (explicit Euler) with timestep $h = 1.0s$. What are x_1^A, x_1^B ? **1 pts.**

$$\begin{pmatrix} x_{k+1}^A \\ x_{k+1}^B \end{pmatrix} = \begin{pmatrix} x_k^A \\ x_k^B \end{pmatrix} + h \begin{pmatrix} \dot{x}_k^A \\ \dot{x}_k^B \end{pmatrix}$$

- iv) The spheres should now be colliding (if not please check your math). Apply the following update rule (simple inelastic collision for spheres of equal mass) with $\varepsilon = 0.25$. What are \dot{x}_1^A, \dot{x}_1^B ? **1 pts.**

$$(\dot{x}_1^A, \dot{x}_1^B) = (\varepsilon \dot{x}_0^B, \varepsilon \dot{x}_0^A)$$

- v) Draw a diagram of the system after this first time step. Include arrows for velocities. Don't forget to draw the origin and x -axis. **1 pts.**
- vi) Follow the same steps to compute x_2^A, x_2^B . The spheres should still be colliding (if not please check your math). **1 pts.**
- vii) What would happen if we naively applied the inelastic collision update rule **again**, to compute \dot{x}_2^A, \dot{x}_2^B ? What does this approach mean for the behavior of the spheres in the long run? **1 pts.**