



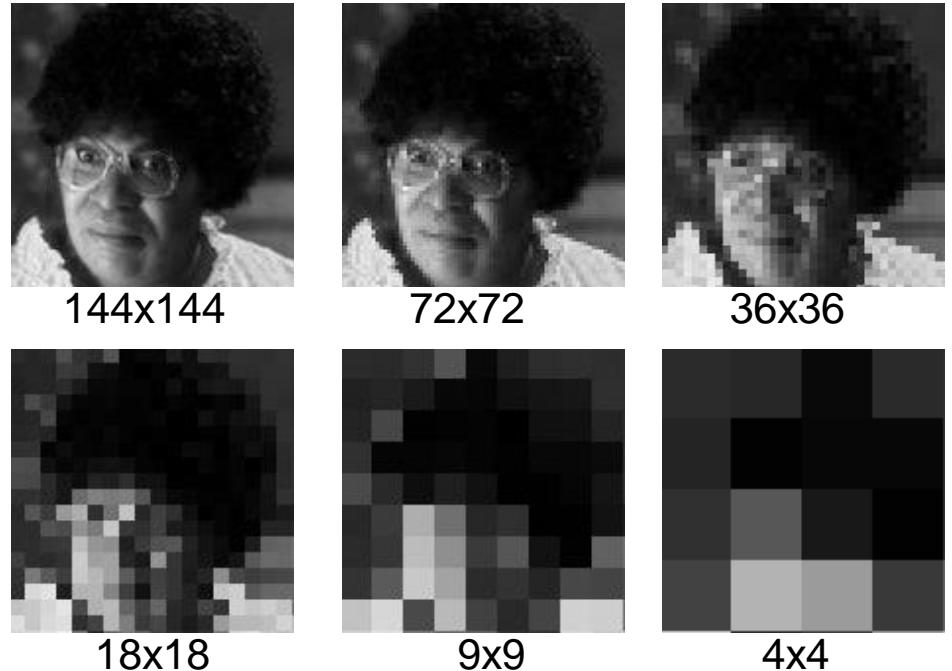
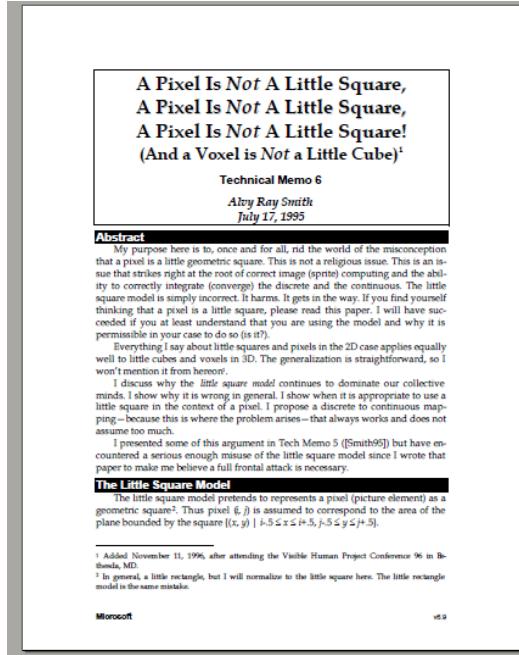
©1994 Encyclopaedia Britannica, Inc.

Visual Computing: Image Segmentation

Prof. Marc Pollefeys

Recap

Geometric resolution



Radiometric resolution



Aliasing and SNR

- What is the disadvantage of low sampling resolution?
- What is the disadvantage of high sampling resolution?
- Lossless vs. Lossy
 - Name some formats?

Unassessed Assignment

Use python to change the geometric and radiometric quantization resolution in one of your images. For each level of sampling and quantization, plot the image function, as in slides 71 & 72, and compare the approximations to the true intensity function that you get at each level.

Usual quantization intervals

- Grayscale image
 - 8 bit = 2^8 = 256 grayvalues
- Color image RGB (3 channels)
 - 8 bit/channel = 2^{24} = 16.7M colors
- 12bit or 16bit from some sensors
- Nonlinear, for example log-scale



Photo: Paulo Barcellos Jr.

Image Noise

- A common model is *additive Gaussian noise*:

$$I(x, y) = f(x, y) + c$$

where $c \sim N(0, \sigma^2)$. So that $p(c) = (2\pi\sigma^2)^{-1} e^{-c^2/2\sigma^2}$

- Poisson noise:

(shot noise) $p(k) = \frac{\lambda^k e^{-\lambda}}{k!}$

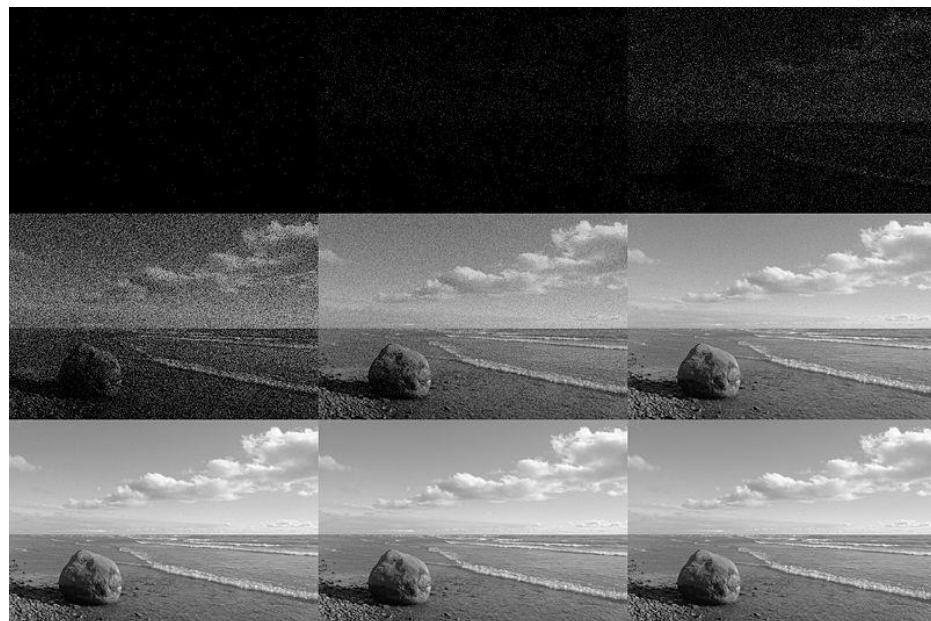
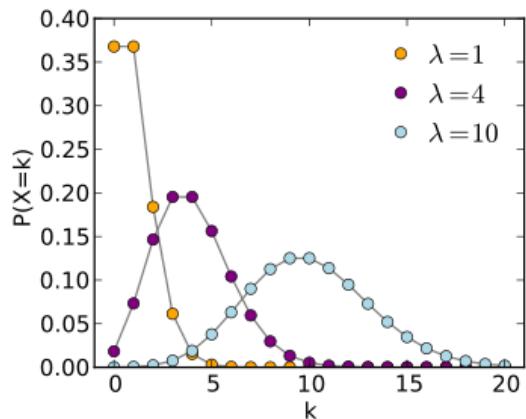
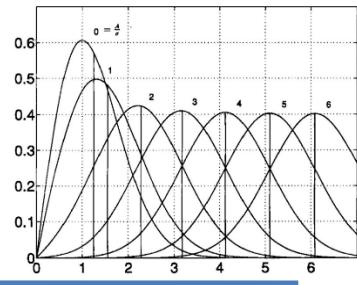


Image Noise



- Rician noise:
(appears in MRI)

$$p(I) = \frac{I}{\sigma^2} \exp\left(-\frac{(I^2 + f^2)}{2\sigma^2}\right) I_0\left(\frac{If}{\sigma^2}\right)$$

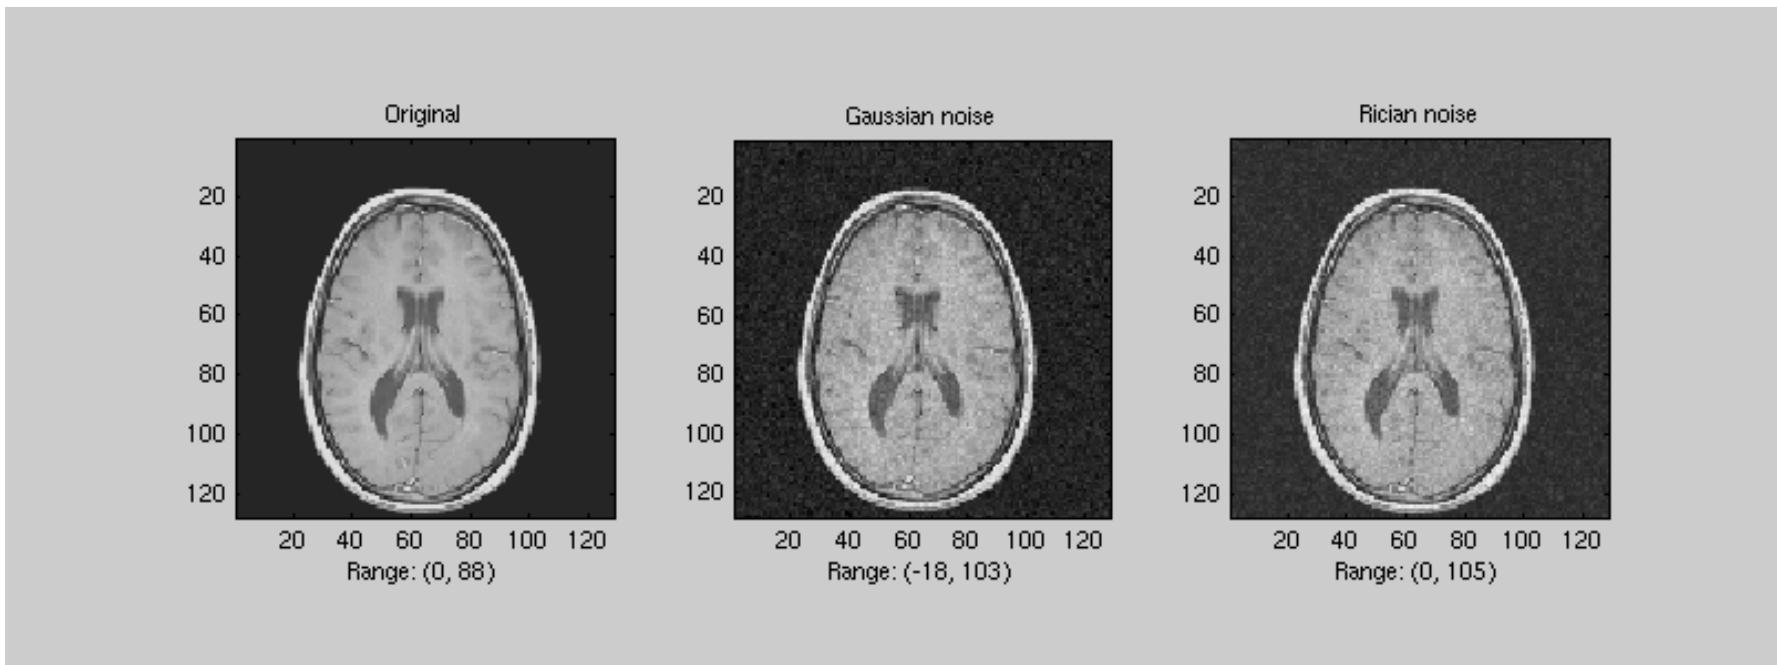


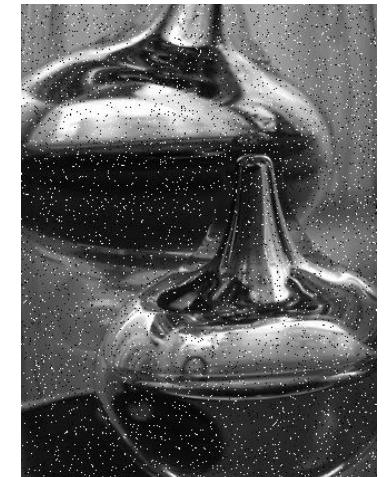
Image Noise

- Multiplicative noise:

$$I = f + fc$$

- Quantization errors
- Impulse “salt-and-pepper” noise

- The *signal to noise ratio (SNR)* s is an index of image quality



$$s = \frac{F}{\sigma}, \text{ where } F = \frac{1}{XY} \sum_{x=1}^X \sum_{y=1}^Y f(x, y)$$

Often used instead: *Peak Signal to Noise Ratio (PSNR)* $s_{peak} = \frac{F_{\max}}{\sigma_{10}}$

Colour Images

R



G



B

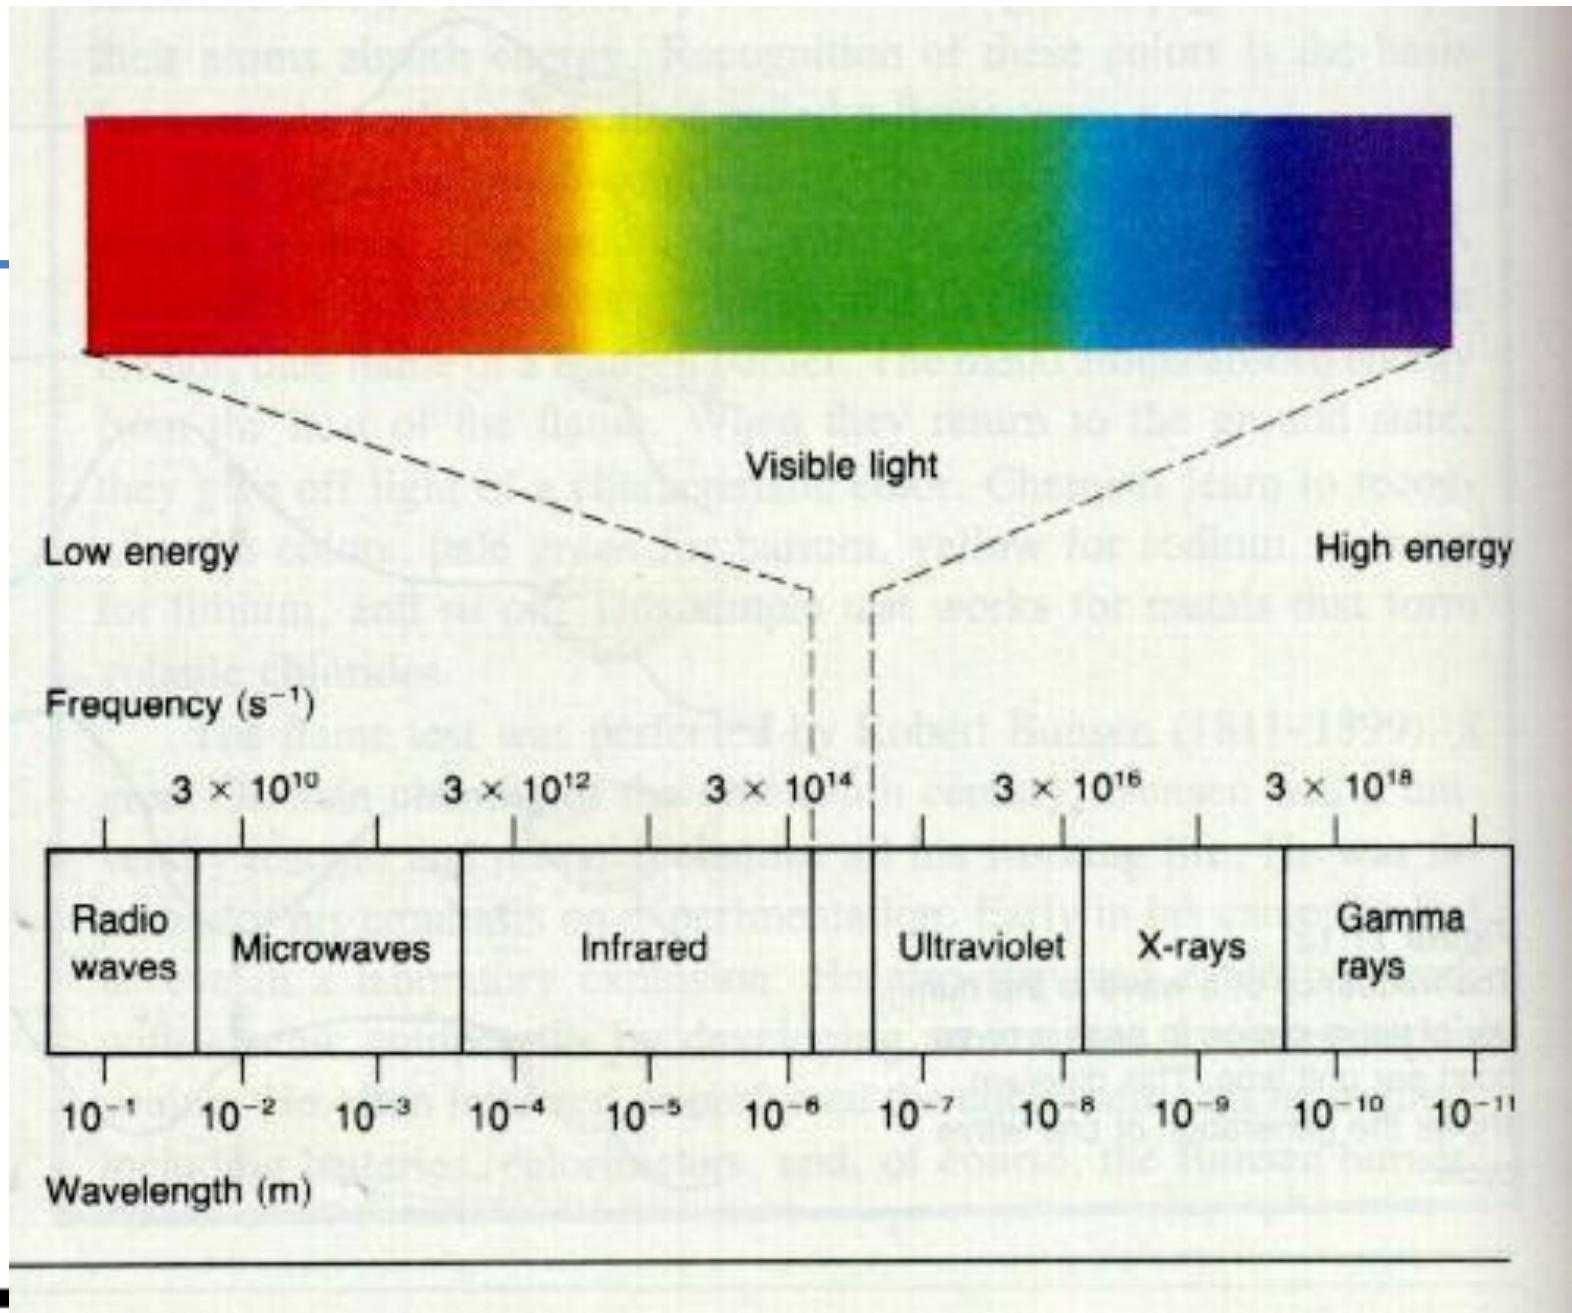


+

+

=





Color cameras

We consider 3 concepts:

1. Prism (with 3 sensors)
2. Filter mosaic
3. Filter wheel

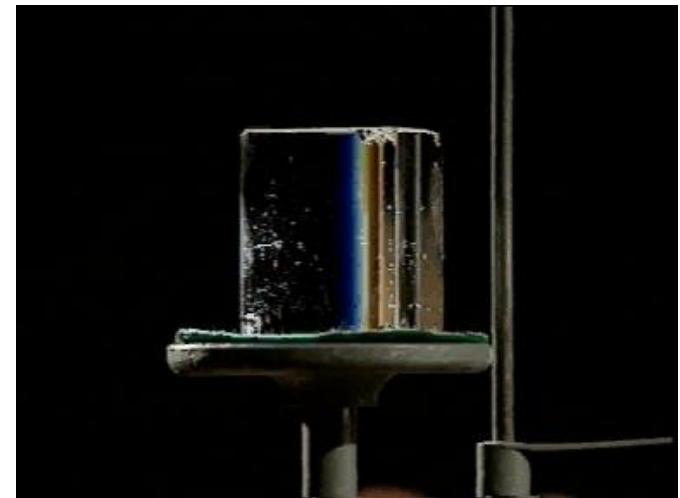
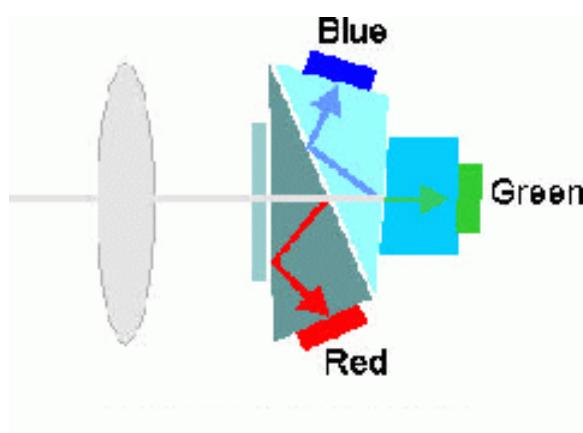
... and X3

Prism color camera

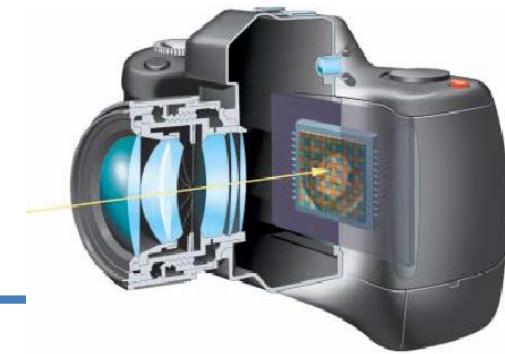
Separate light in 3 beams using dichroic prism

Requires 3 sensors & precise alignment

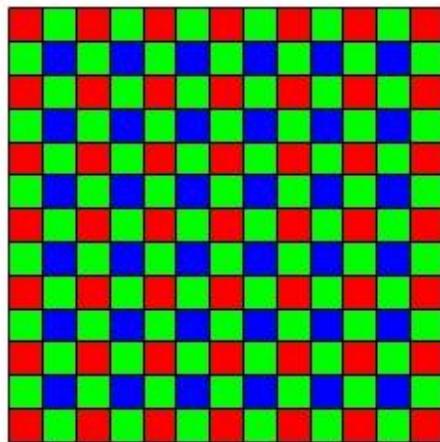
Good color separation



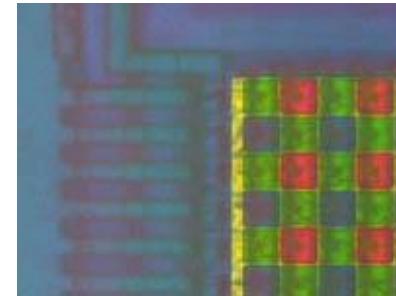
Filter mosaic



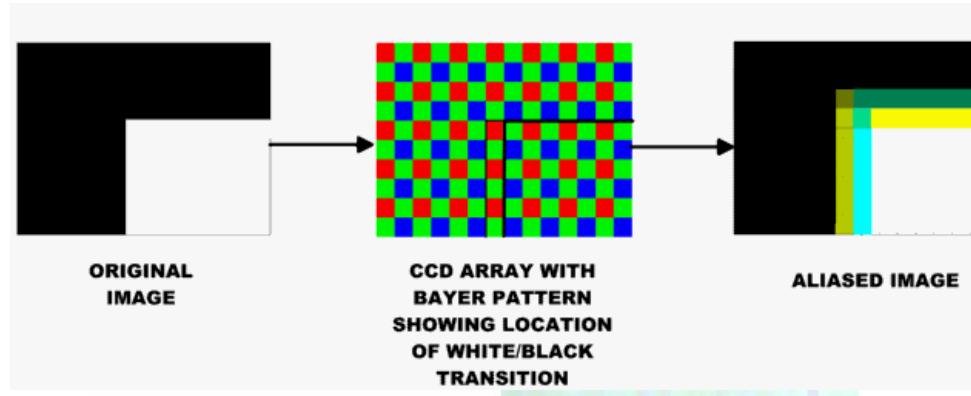
Coat filter directly on sensor



Bayer filter



Demosaicing (obtain full colour & full resolution image)



More colors:

R	E	R	E
G	B	G	B
R	E	R	E
G	B	G	B

Filter wheel

Rotate multiple filters in front of lens

Allows more than 3 colour bands



Only suitable for static scenes

Prism vs. mosaic vs. wheel

<u>approach</u>	<u>Prism</u>	<u>Mosaic</u>	<u>Wheel</u>
# sensors	3	1	1
Separation	High	Average	Good
Cost	High	Low	Average
Framerate	High	High	Low
Artefacts	Low	Aliasing	Motion
Bands	3	3	3 or more

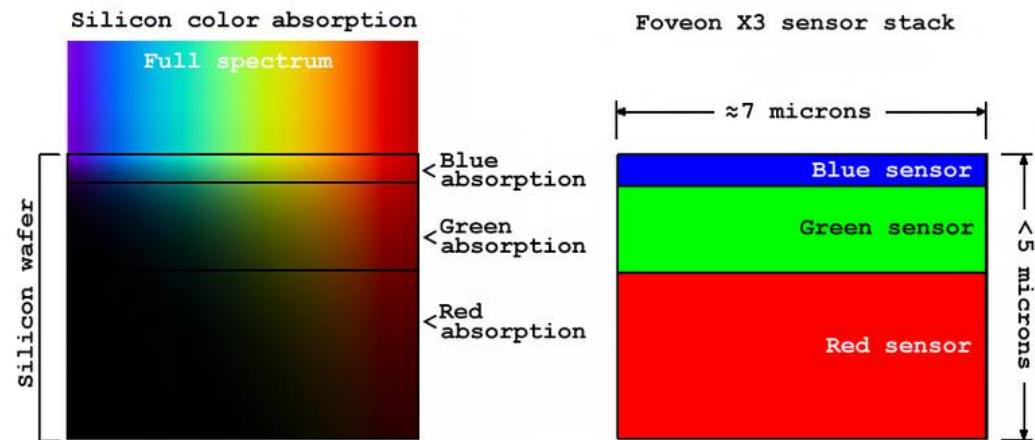
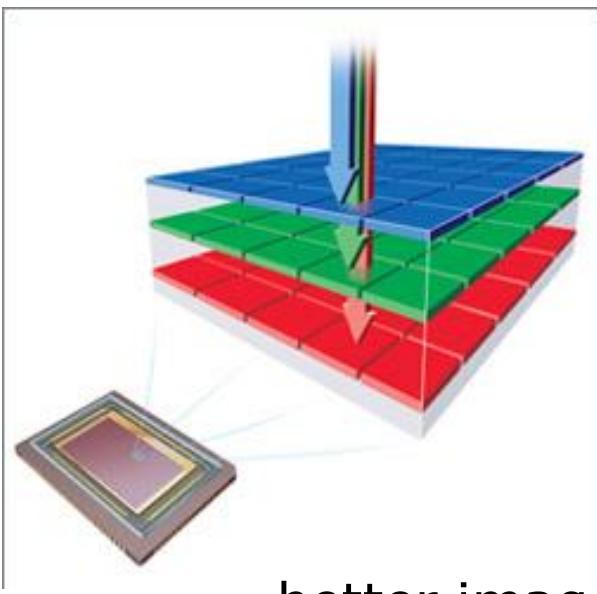
High-end
cameras

Low-end
cameras

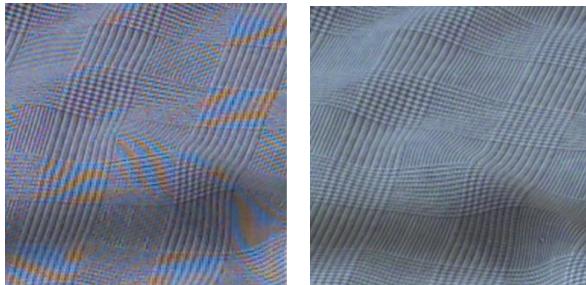
Scientific
applications

color CMOS sensor

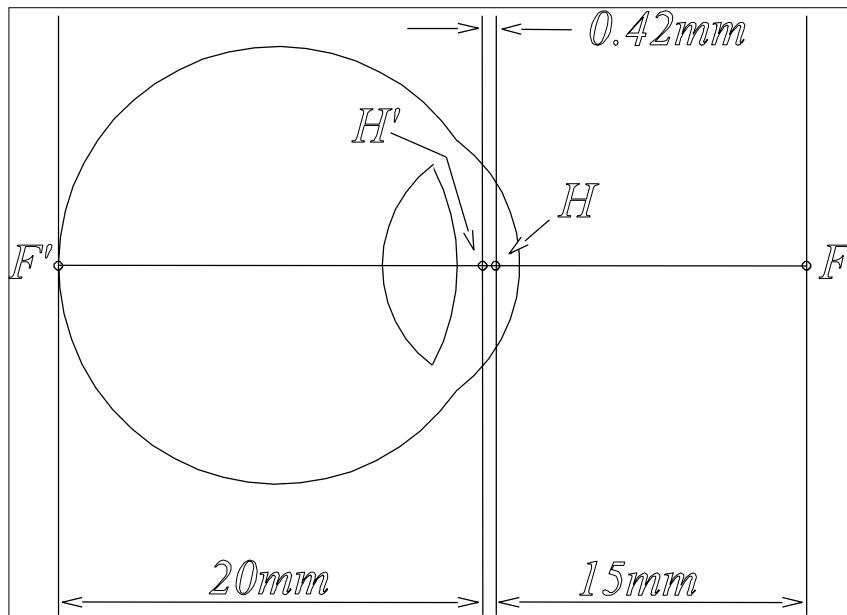
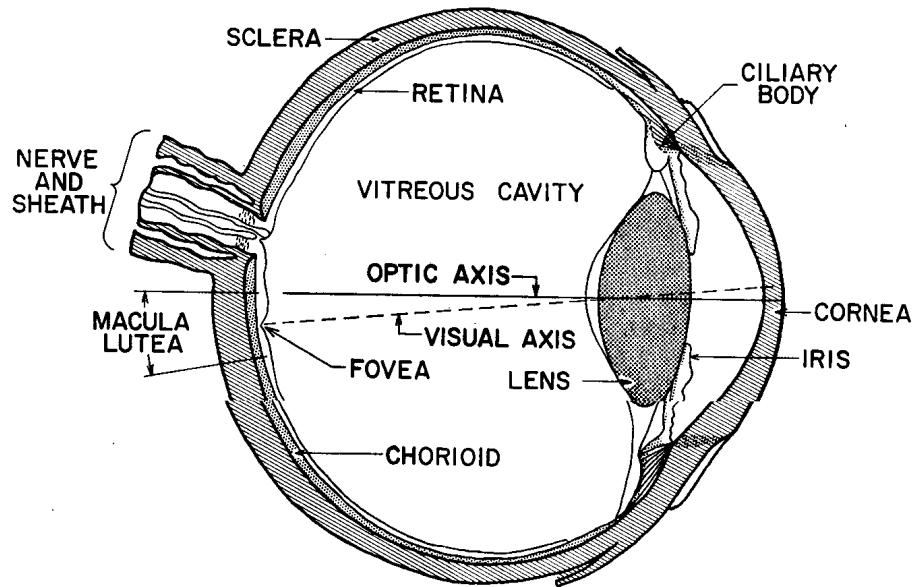
Foveon's X3



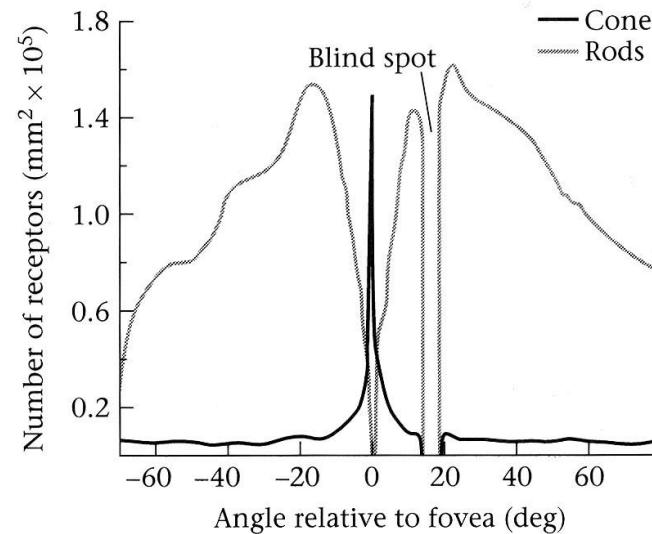
better image quality



The Human Eye

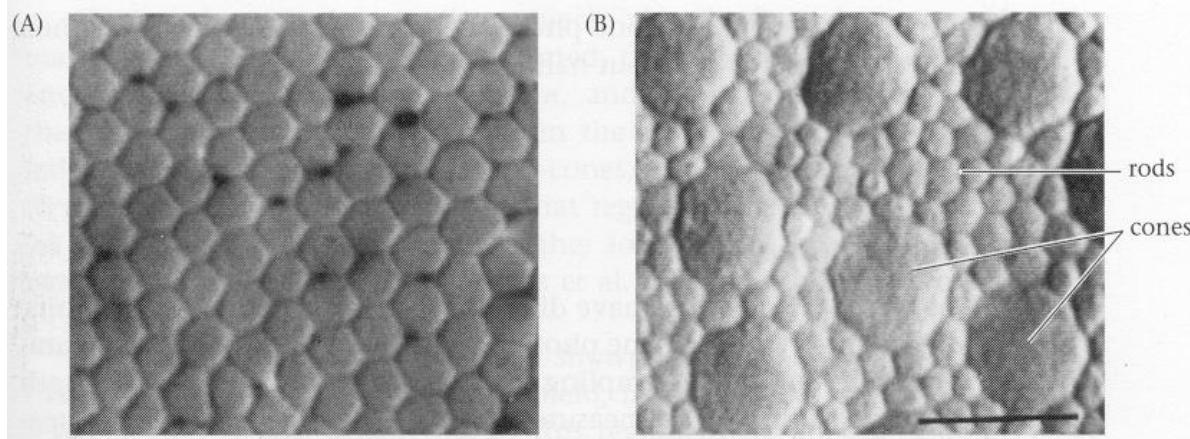


The distribution of rods and cones across the retina



Reprinted from Foundations of Vision, by B. Wandell, Sinauer Associates, Inc., (1995). © 1995 Sinauer Associates, Inc.

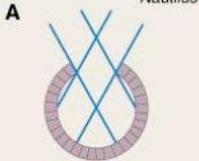
Cones in the fovea



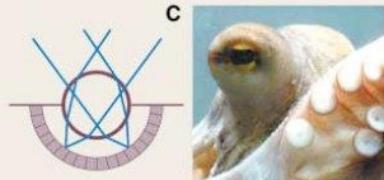
Reprinted from Foundations of Vision, by B. Wandell, Sinauer Associates, Inc., (1995). © 1995 Sinauer Associates, Inc.

More eyes in nature...

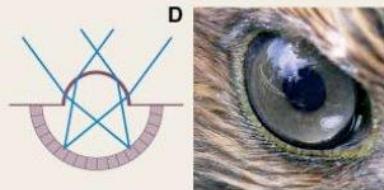
Chambered eyes



Nautilus



Octopus

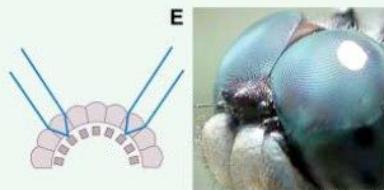


Red-tailed hawk

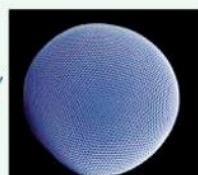
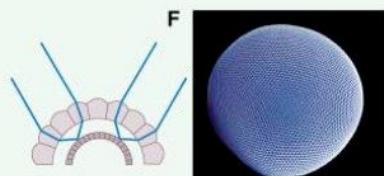
Compound eyes



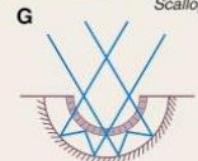
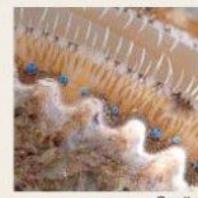
Sea fan



Dragonfly

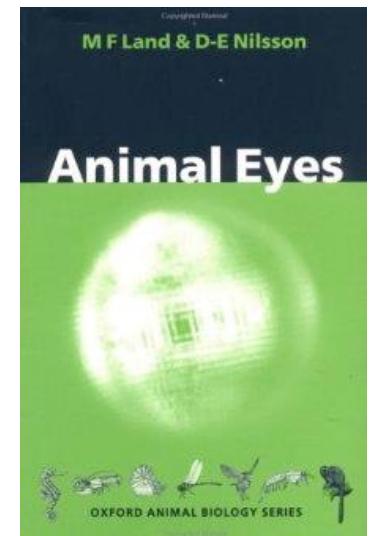


Krill eye



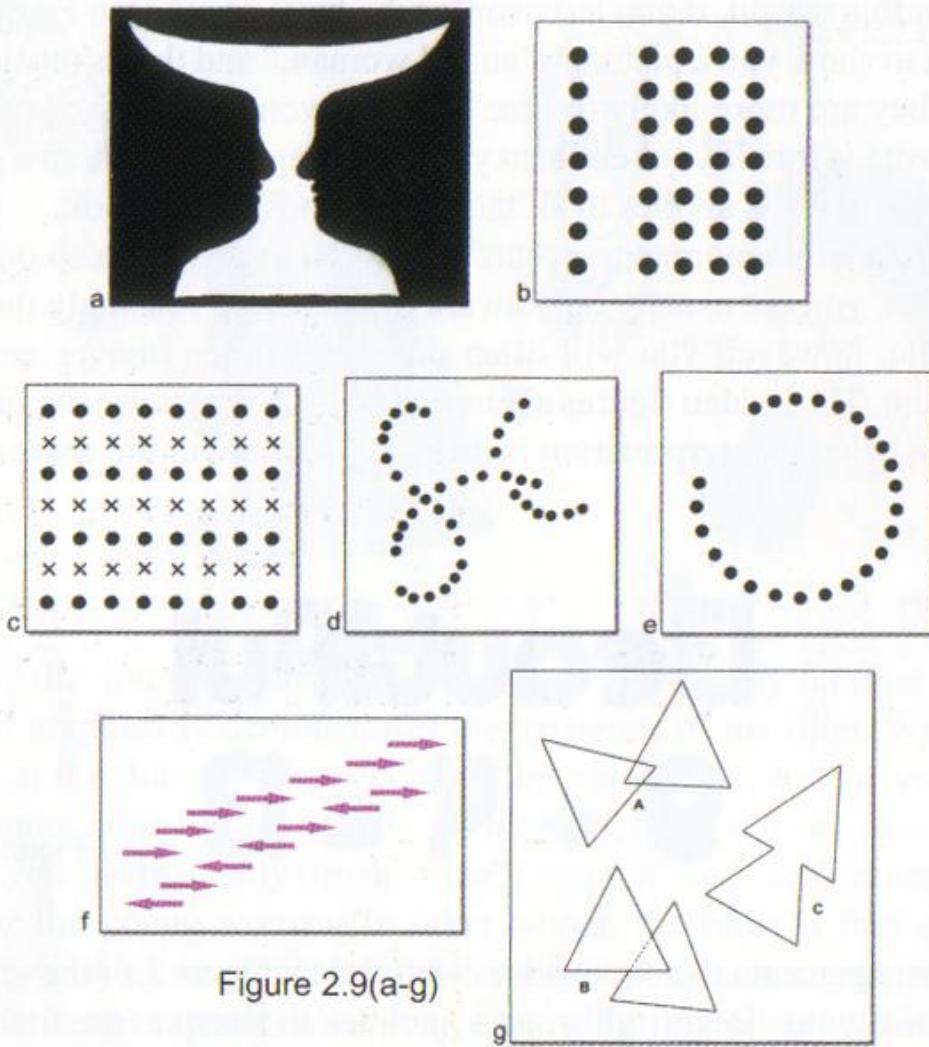
Scallop

More info:



ET...

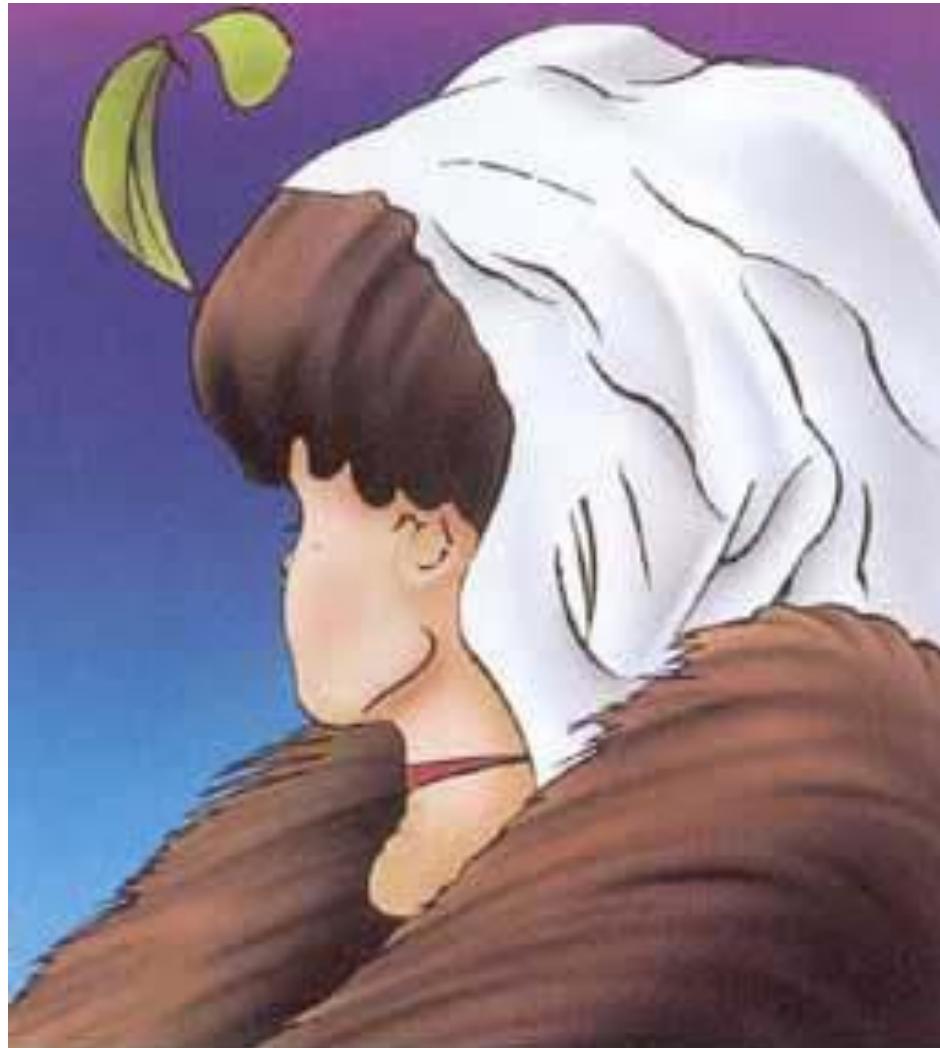
Fernald, R. D. 2006. Casting a Genetic Light on the Evolution of Eyes. Science 313, 1914-1918



Gestalt Phenomena

- Figure-ground
- Proximity
- Similarity
- Continuation
- Closure
- Common fate
- Symmetry

Young Lady or An Old Hag?



What is Image Segmentation?

Segmentation is the ultimate classification problem.

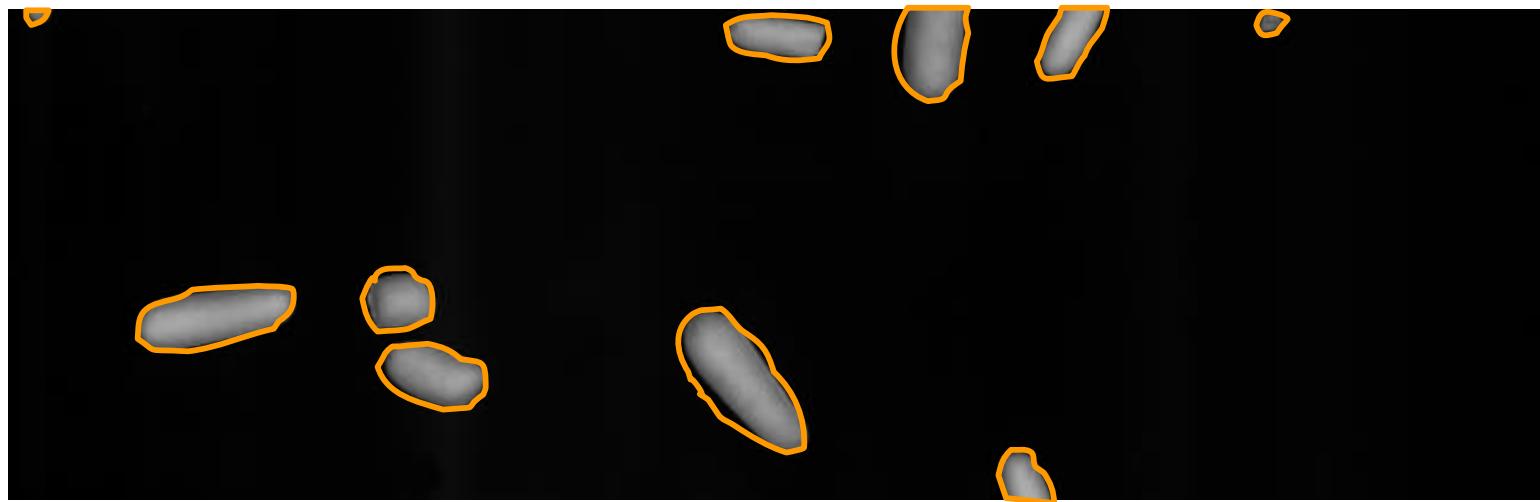
Once solved, Computer Vision is solved.

What is Image Segmentation?

- It partitions an image into regions of interest
- The first stage in many automatic image analysis systems
- A *complete segmentation* of an image I is a finite set of regions R_1, \dots, R_N , such that

$$I = \bigcup_{i=1}^N R_i \text{ and } R_i \cap R_j = \emptyset \forall i \neq j.$$

How should I segment this?



Exclude dark pixels?

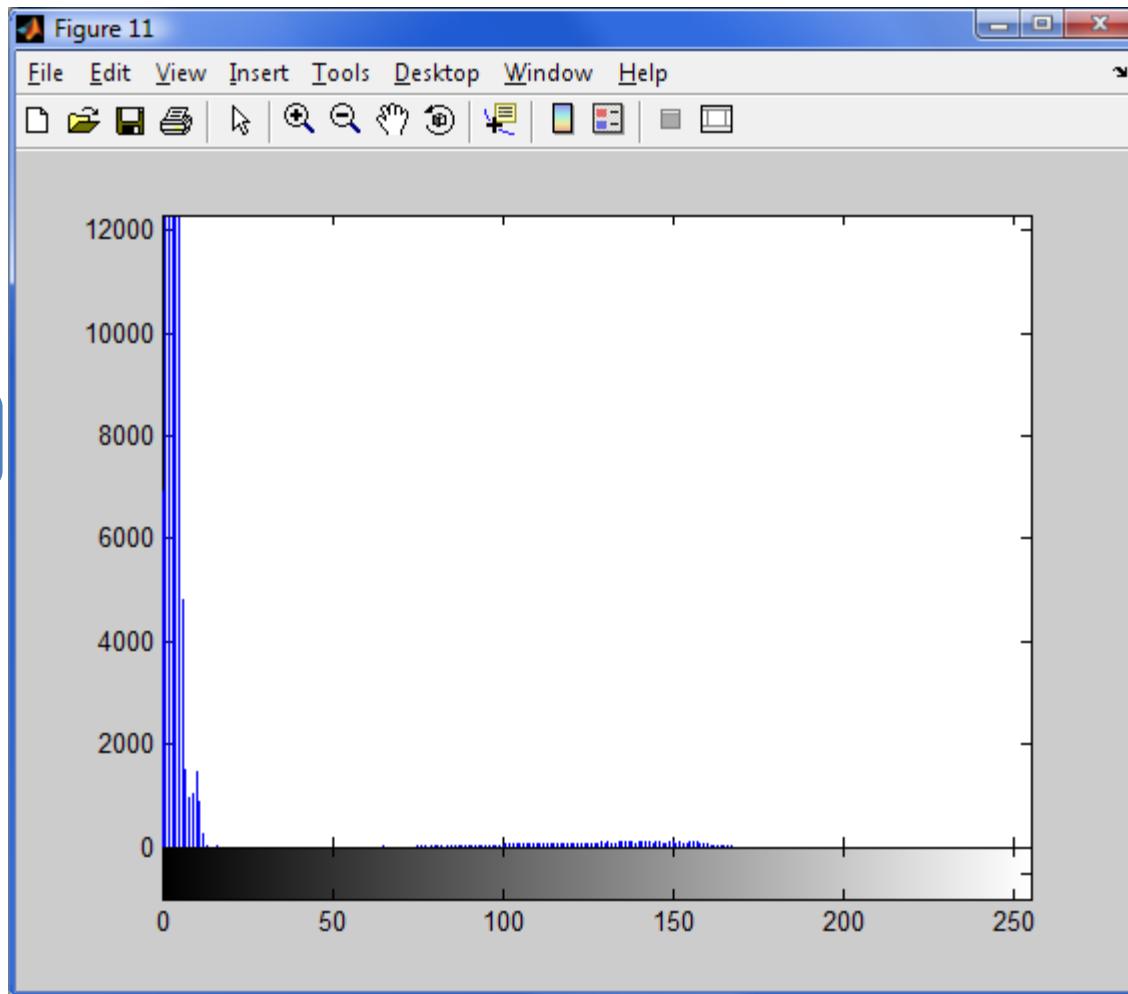
```
img = cv2.imread('BlobsIP.png')
cv2.imshow('BlobsIP', img)
cv2.waitKey(0)

img.shape      --> [ 244    767    3  ]

hist = np.histogram(img, bins=256)
cv2.imshow('Histogram', hist)
cv2.waitKey(0)

cv2.imshow('Mask', img[:, :, 1] > 20)
cv2.waitKey(0)
```

Histogram



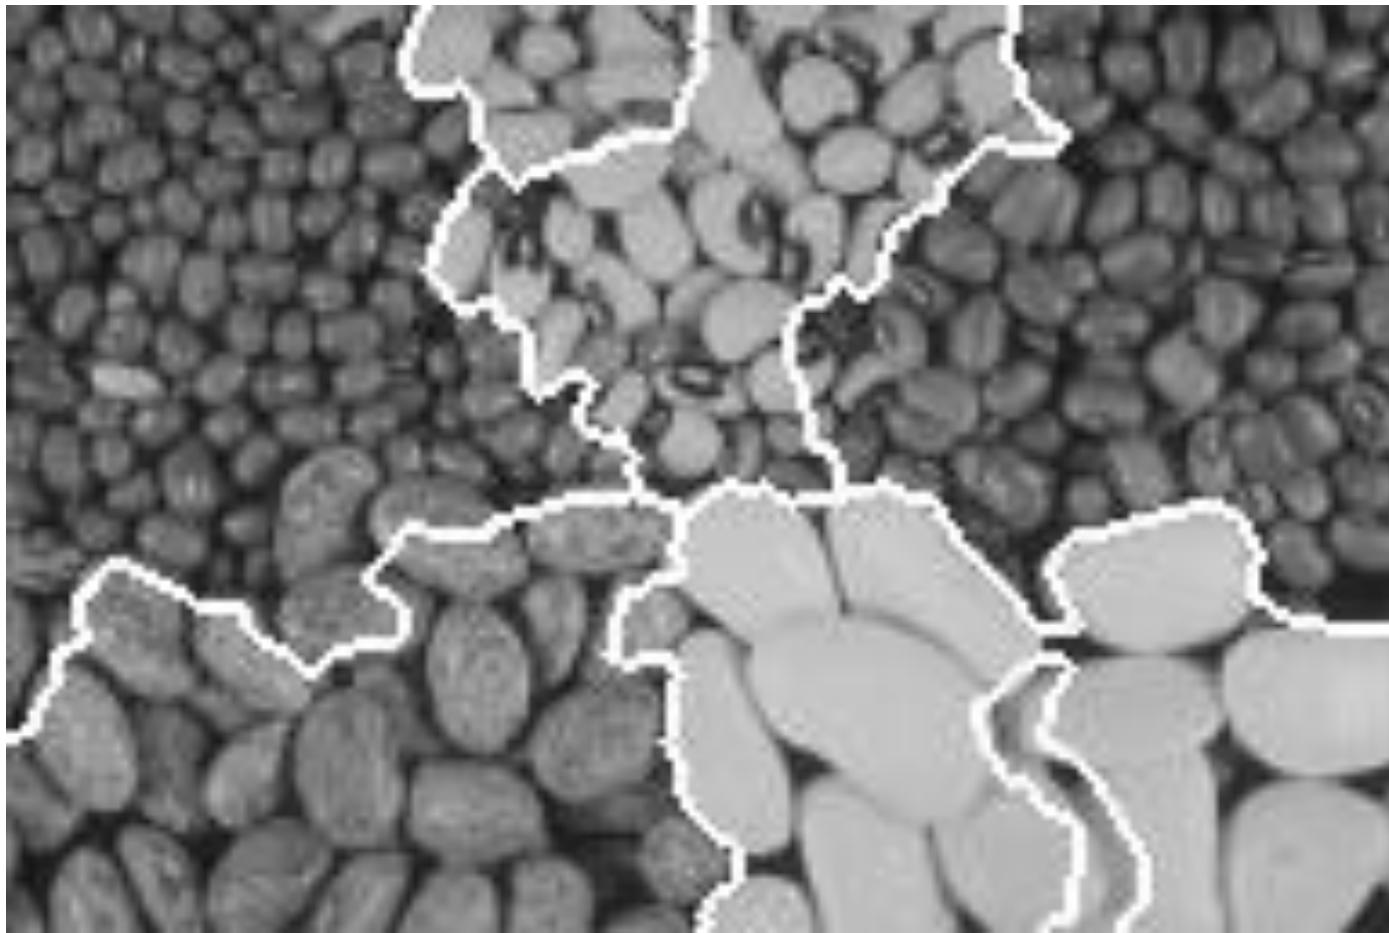
How should I segment this?



Segmentation Quality

- The quality of a segmentation depends on what you want to do with it.
 - Segmentation algorithms must be chosen and evaluated with an application in mind.
-

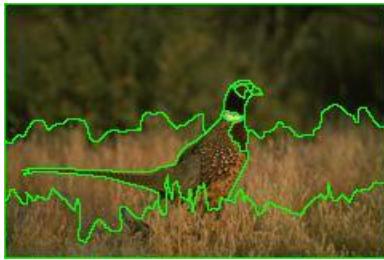
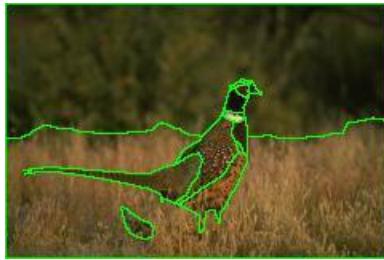
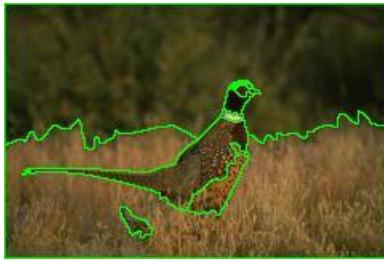
Segmentation example



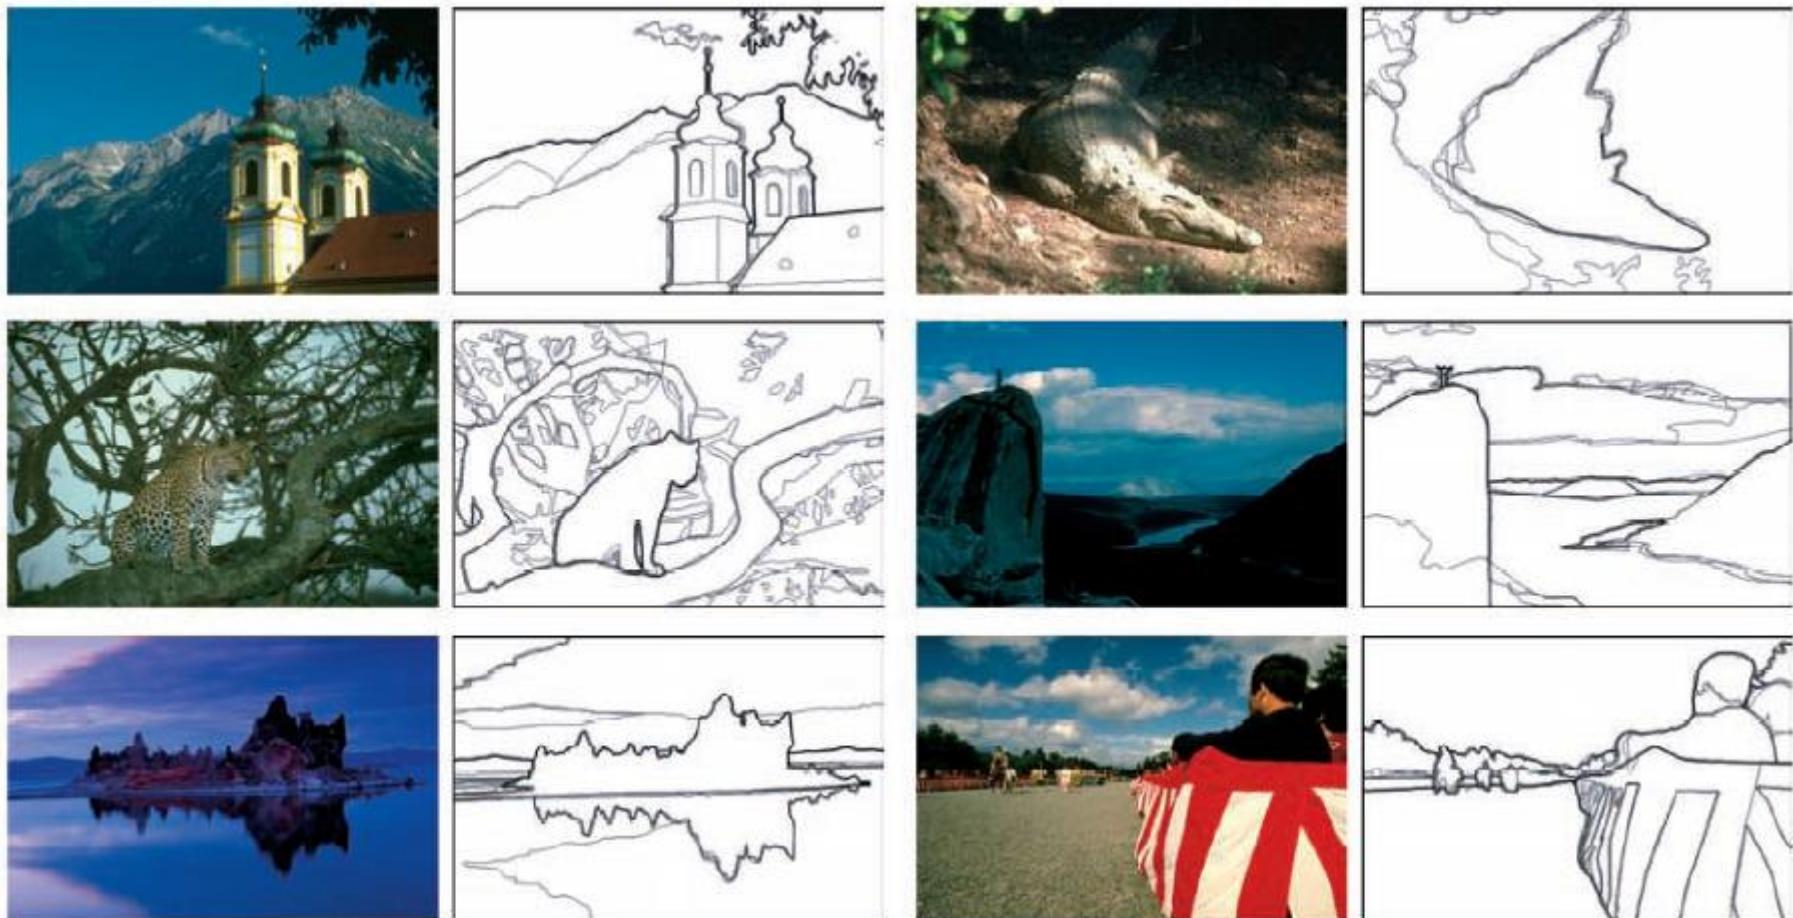
Berkeley Segmentation Database and Benchmark



<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/resources.html>



Berkeley Segmentation Database and Benchmark



Martin et al. PAMI 2004

Thresholding

- Thresholding is a simple segmentation process.
- Thresholding produces a binary image B .
- It labels each pixel **in** or **out** of the region of interest by comparison of the greylevel with a threshold T :

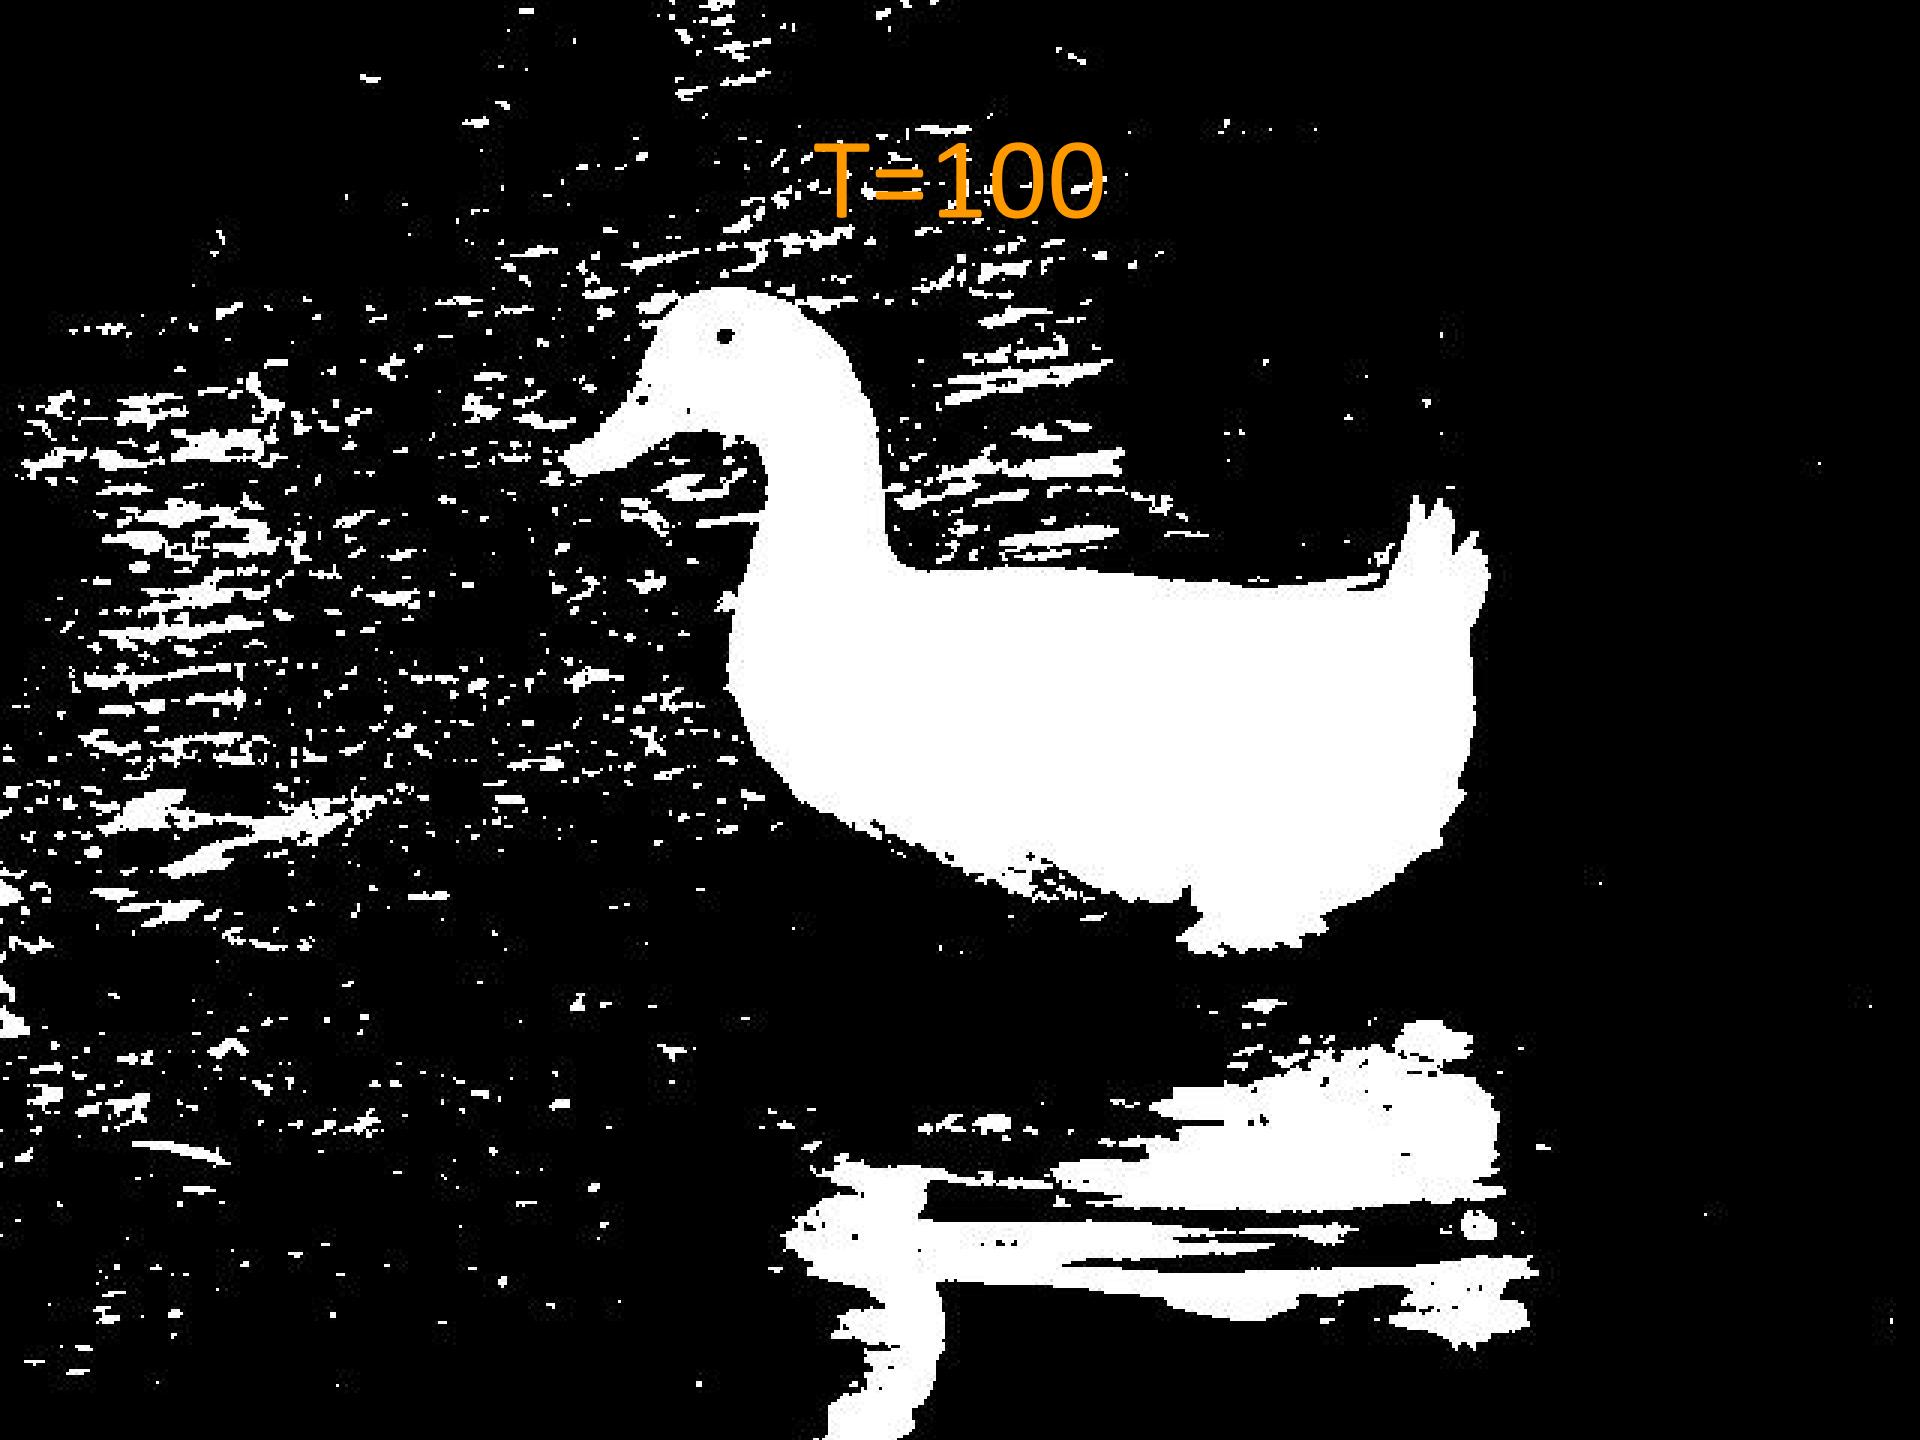
$$B(x, y) = \begin{cases} 1 & \text{if } I(x, y) \geq T \\ 0 & \text{if } I(x, y) < T. \end{cases}$$

Thresholding example

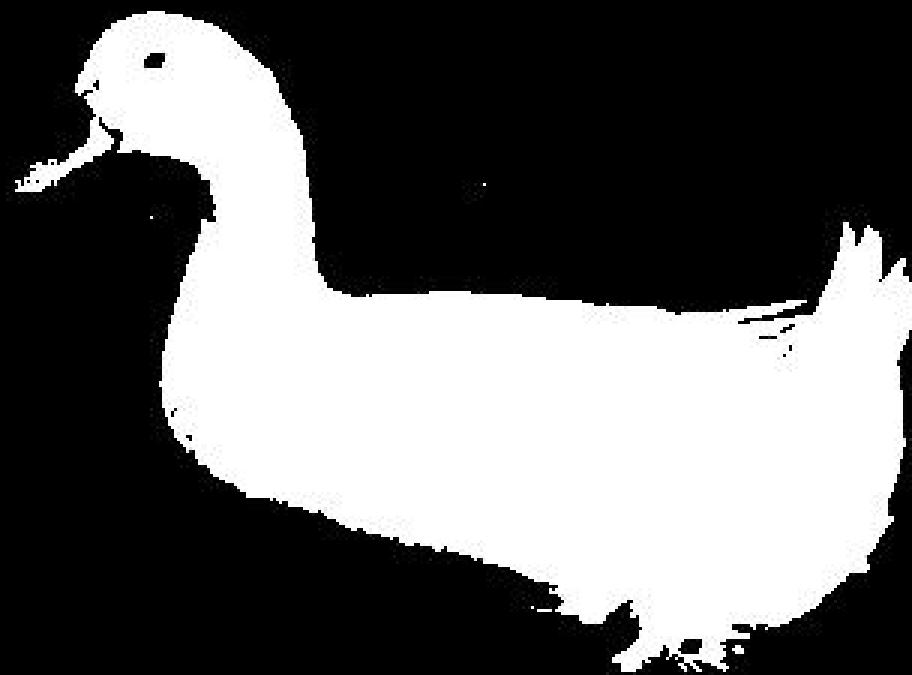


T=50

T=100



T=150



$T=200$

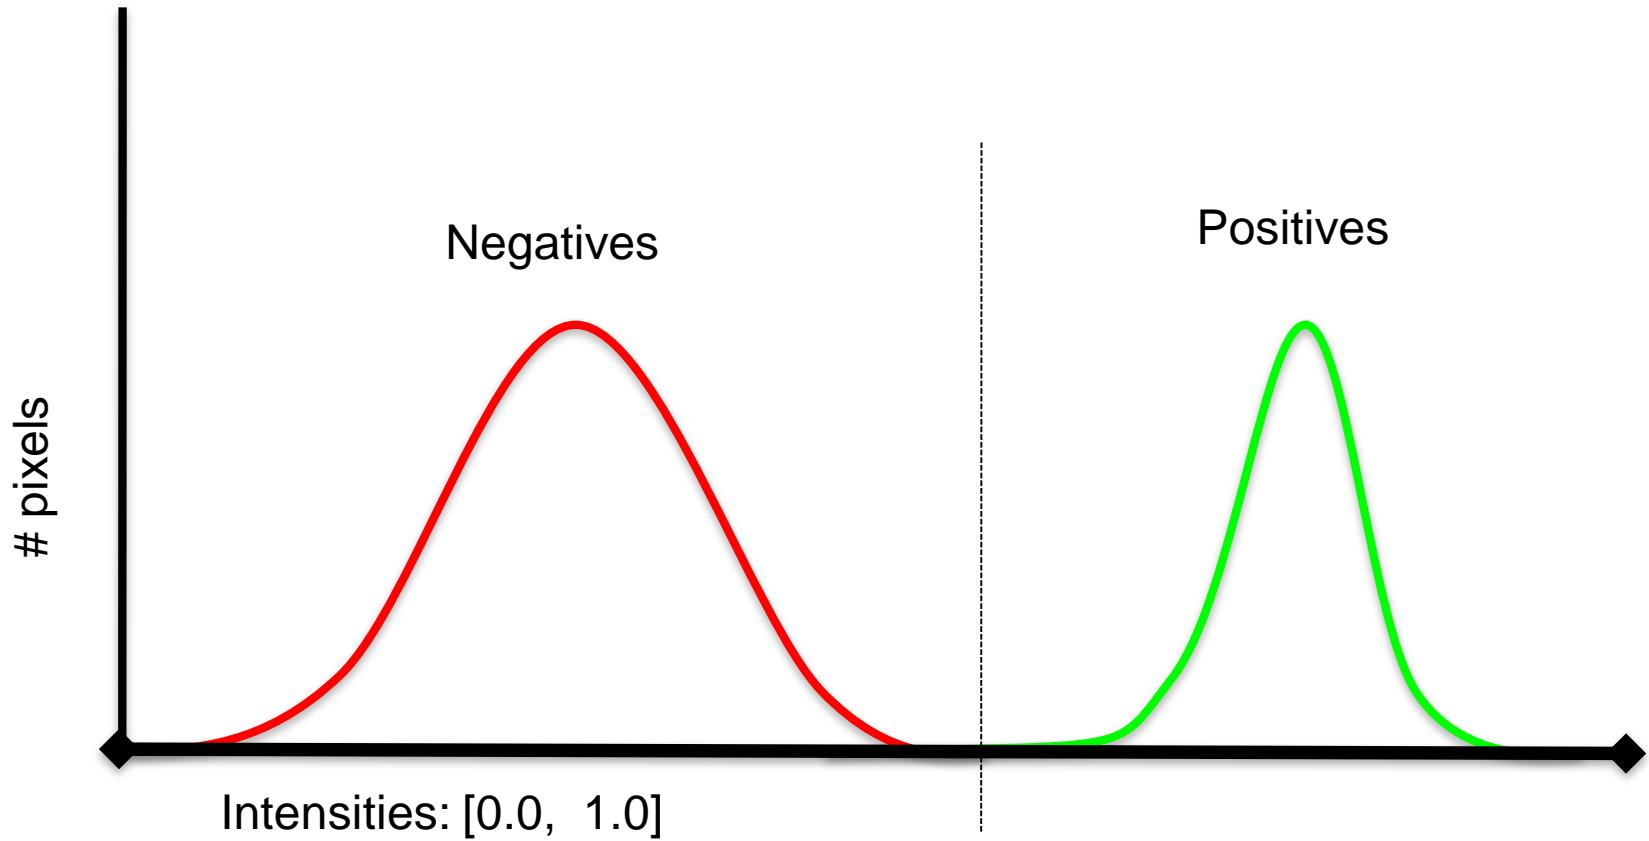


How do we choose T?

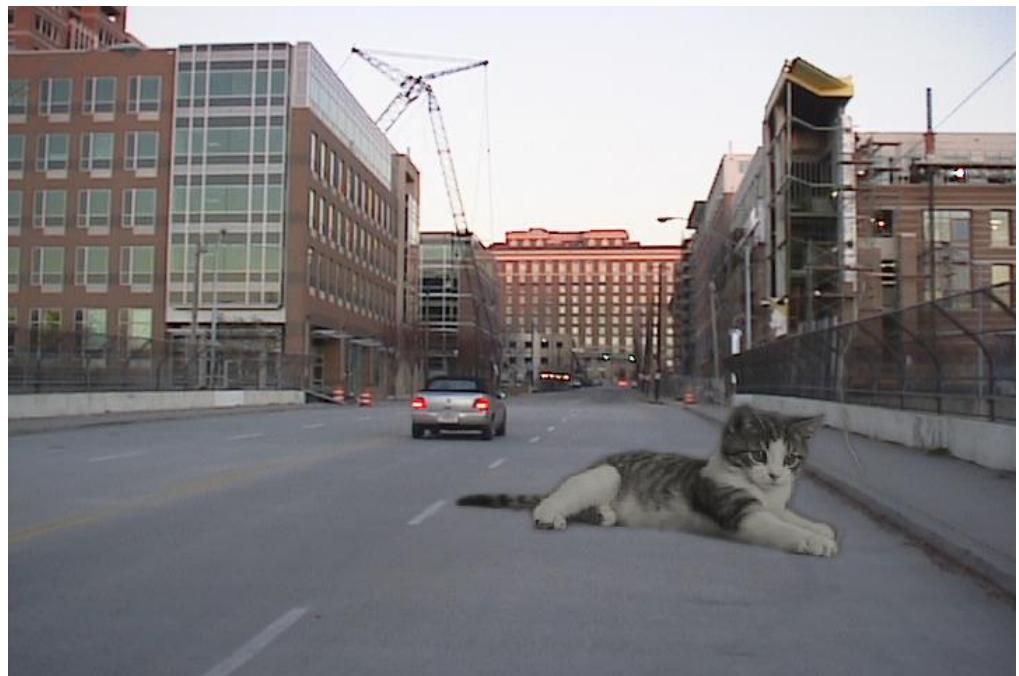
- Trial and error
- Compare results with ground truth
- Automatic methods*

*= We'll discuss ROC curves later

Wouldn't it be nice...



Planning to Segment?



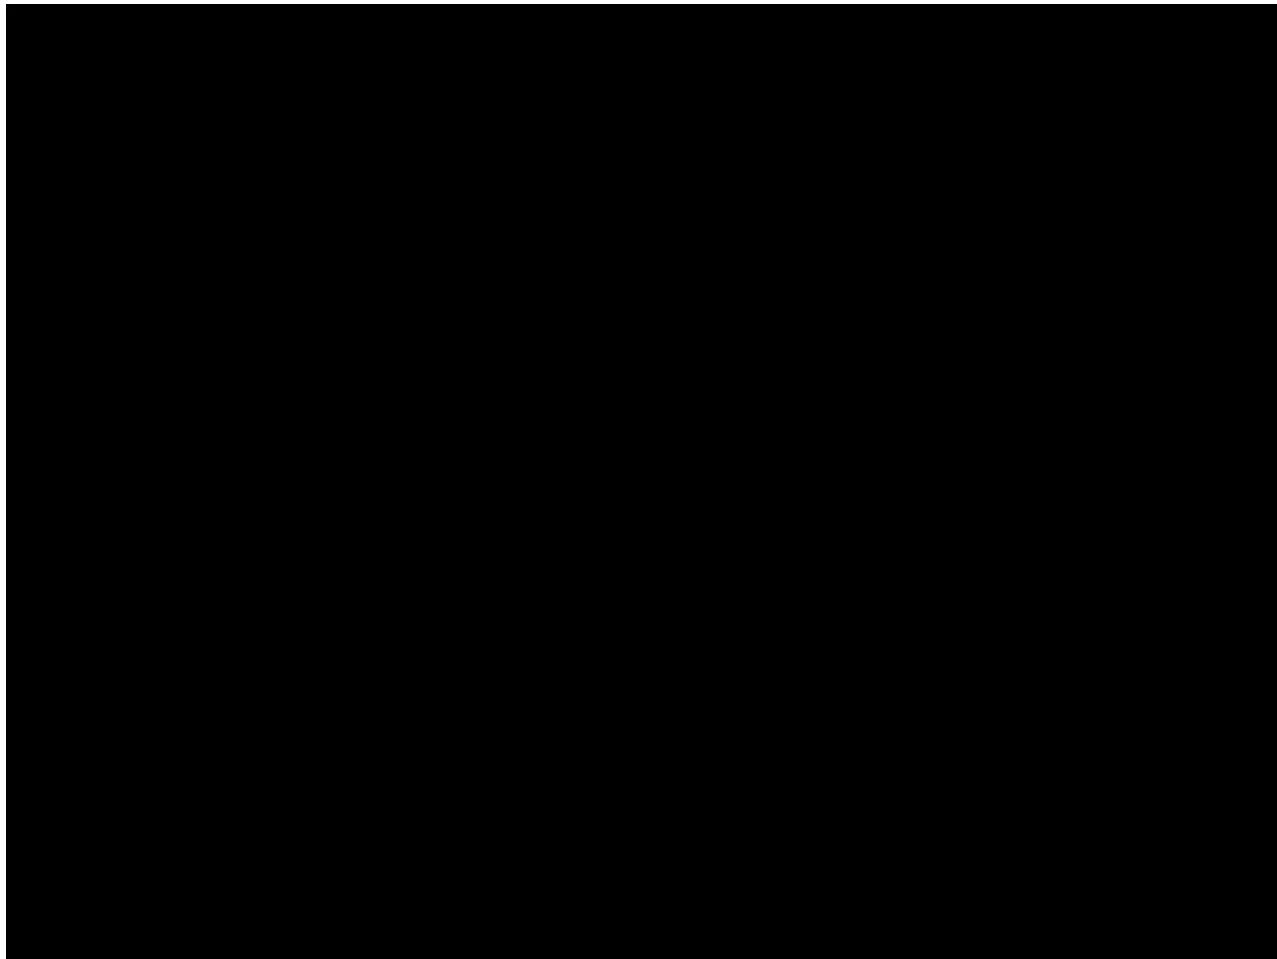
Shraybman, HW1, 2003

Chromakeying: Control Lighting!



generalspecialist.com

Chrome keying example



Chromakeying

- “Plain” distance measure

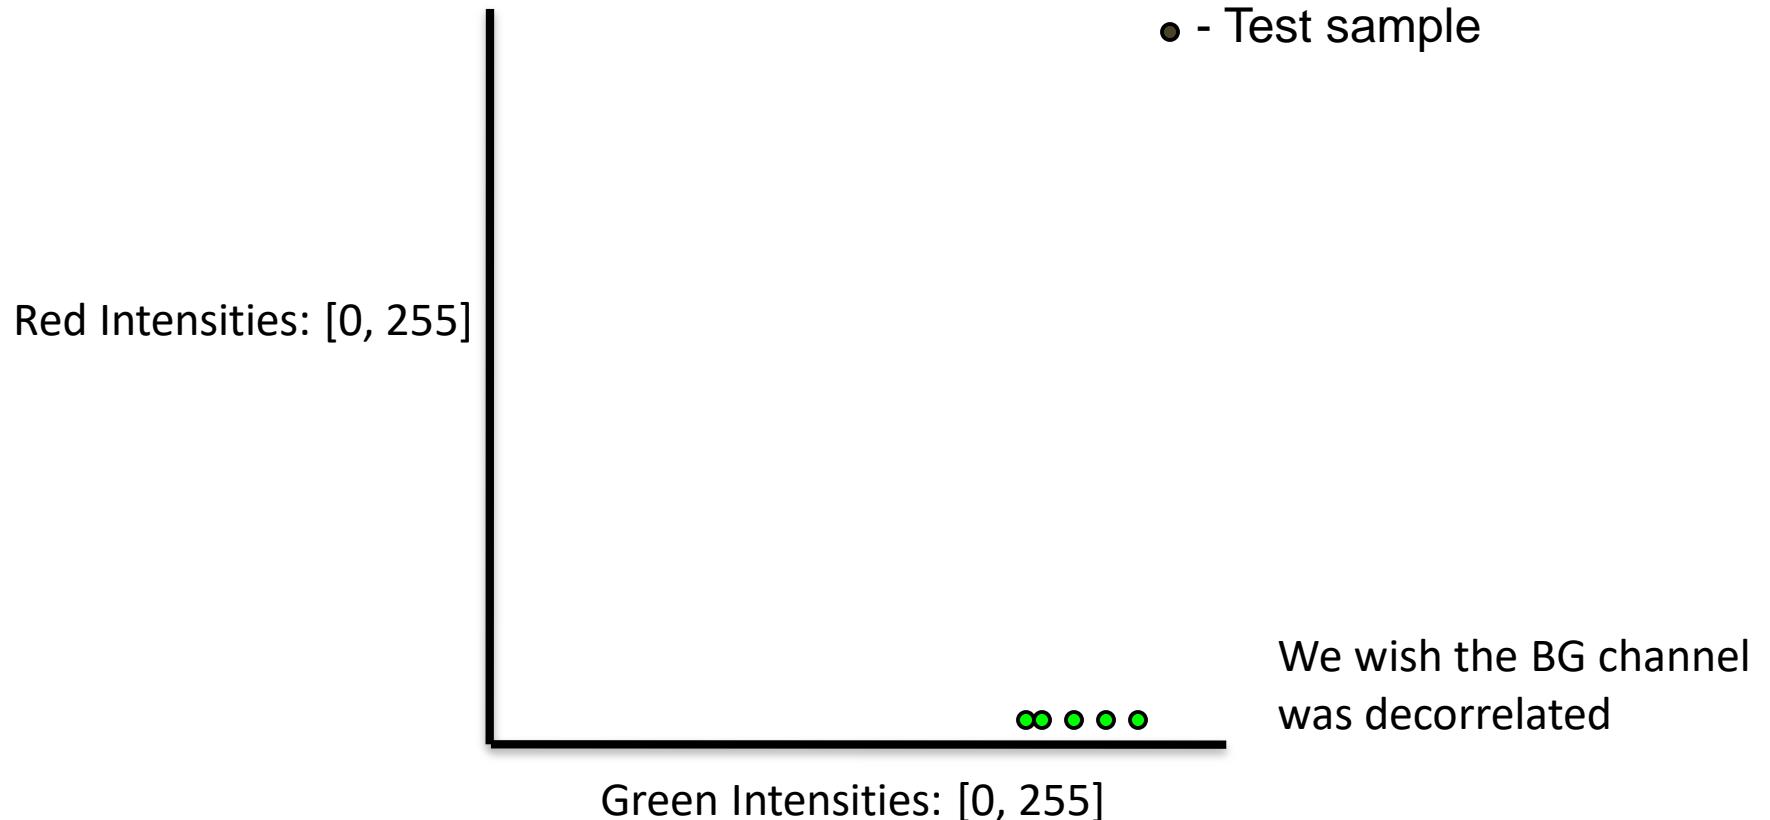
$$I_{\alpha} = |I - g| > T$$

$$T = \sim 20$$

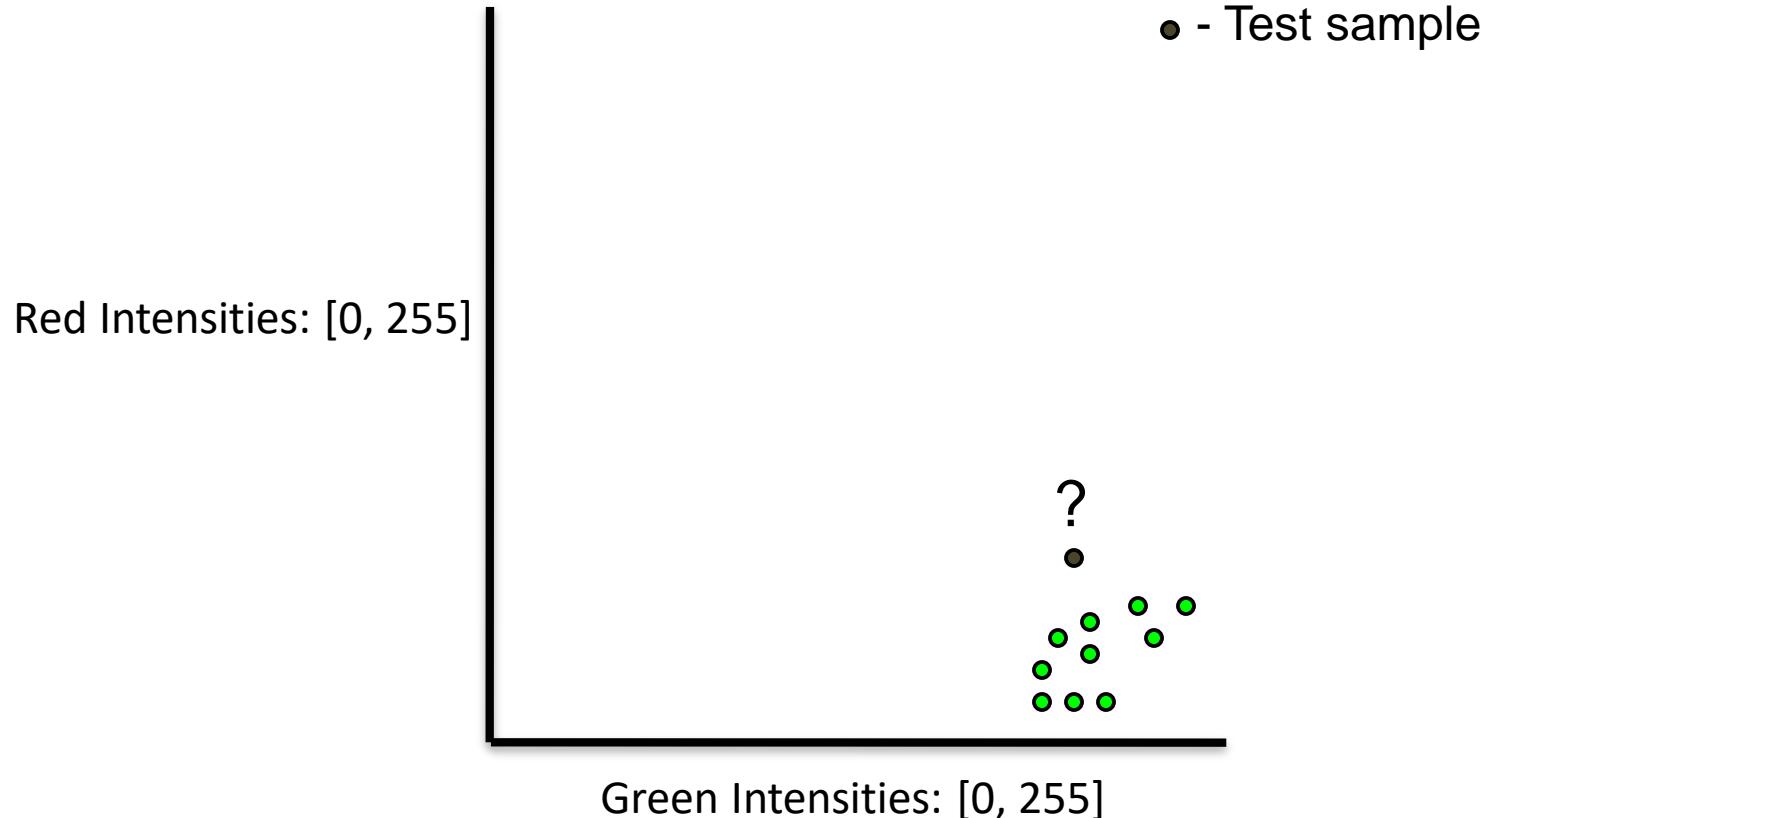
$$g = (0 \ 255 \ 0) \quad (\text{for example})$$

- Problems:
 - Variation is NOT same in all 3 channels
 - Hard alpha mask: $I_{\text{comp}} = I_{\alpha} I_a + (1-I_{\alpha}) I_b$

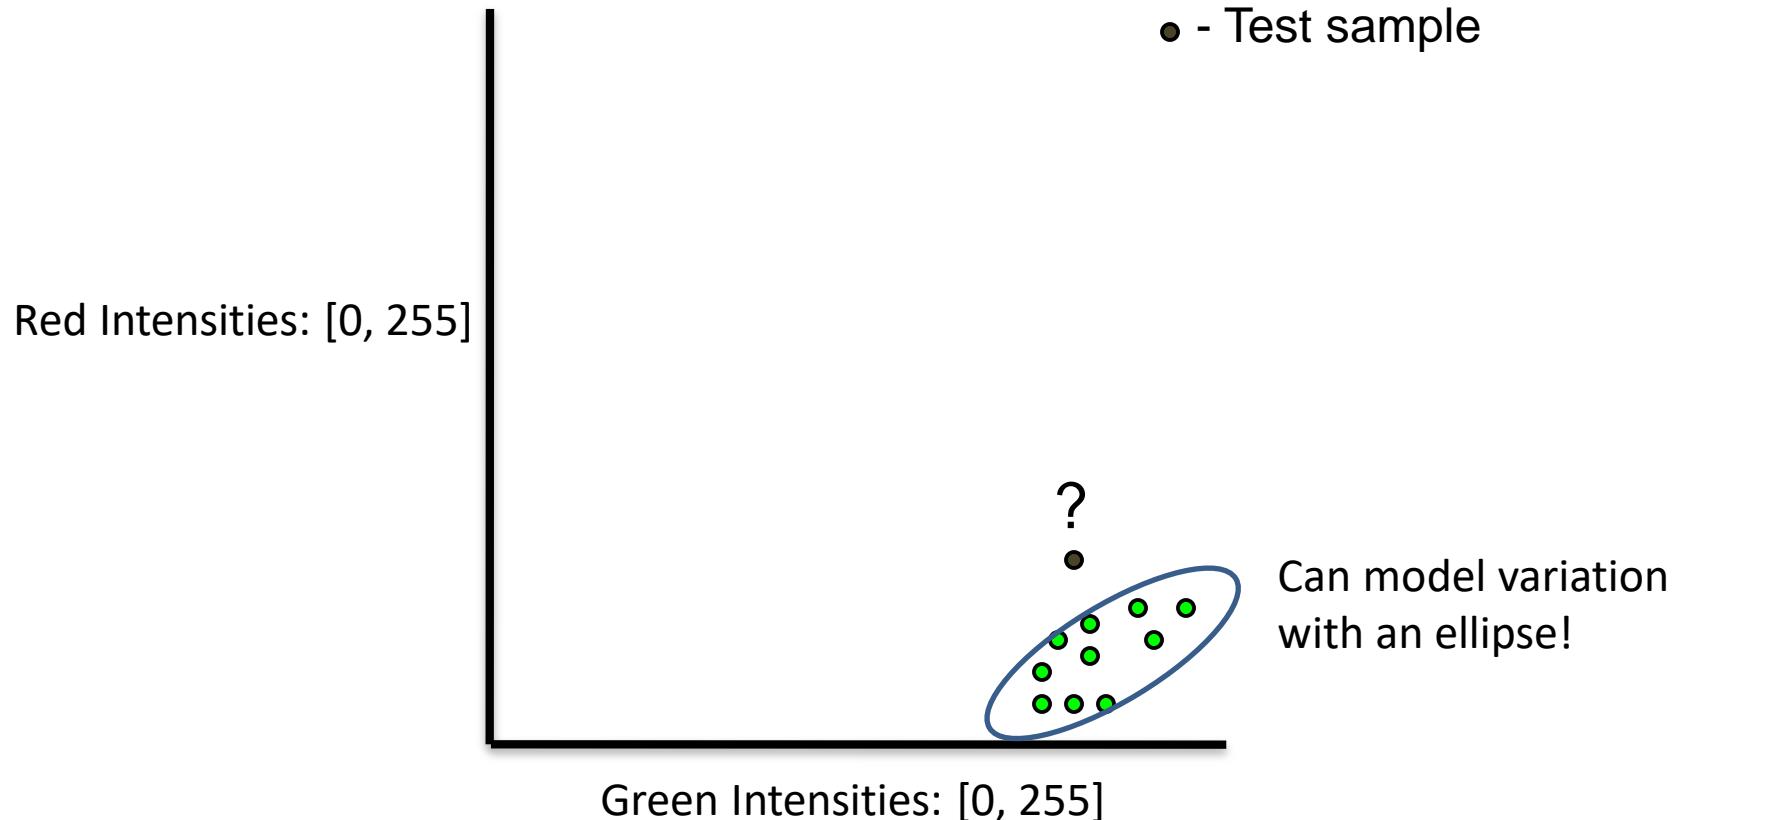
Background Color Variation



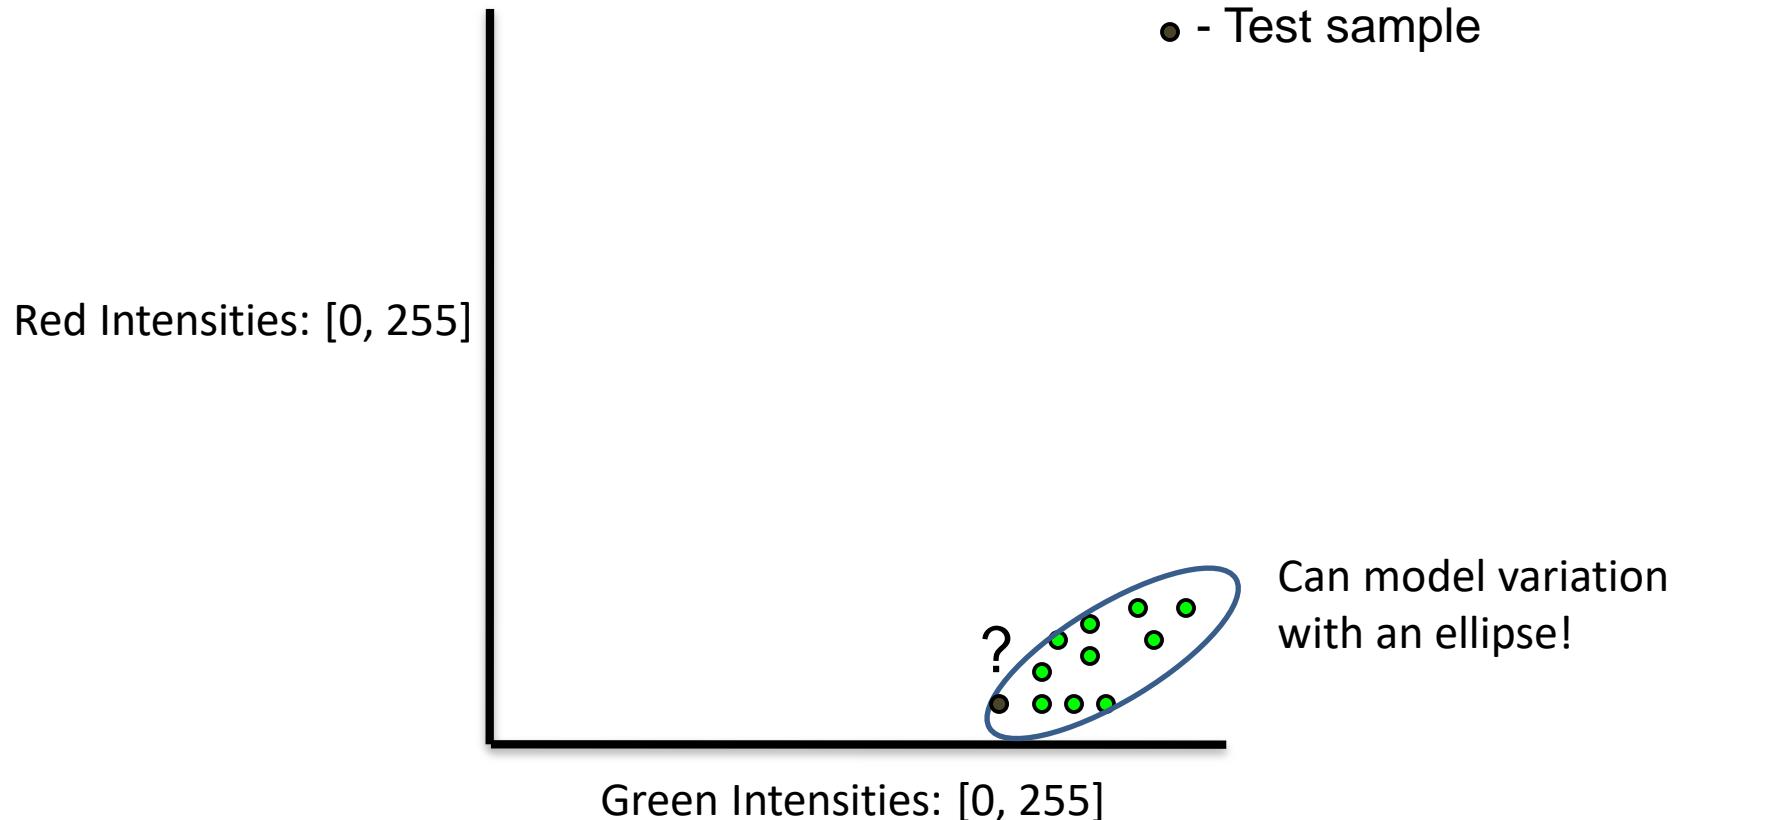
Background Color Variation



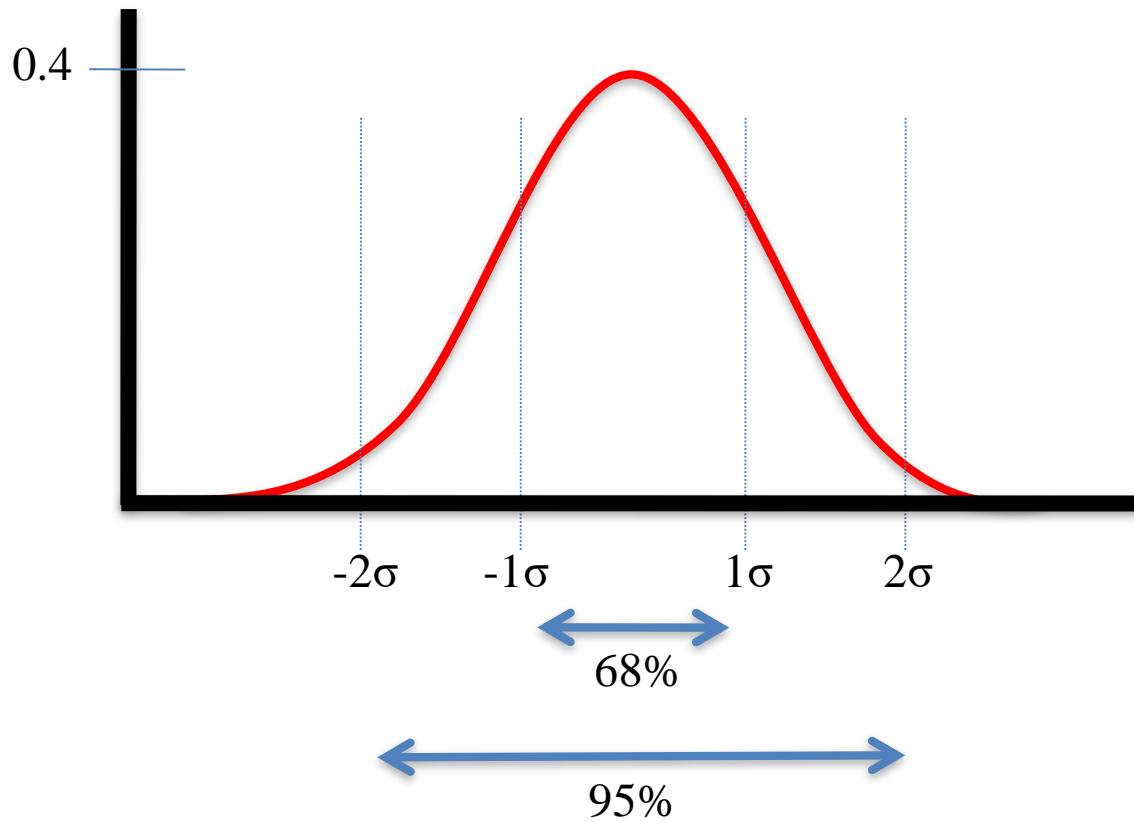
Background Color Variation



Background Color Variation



Use Gaussian to Explain Most Data



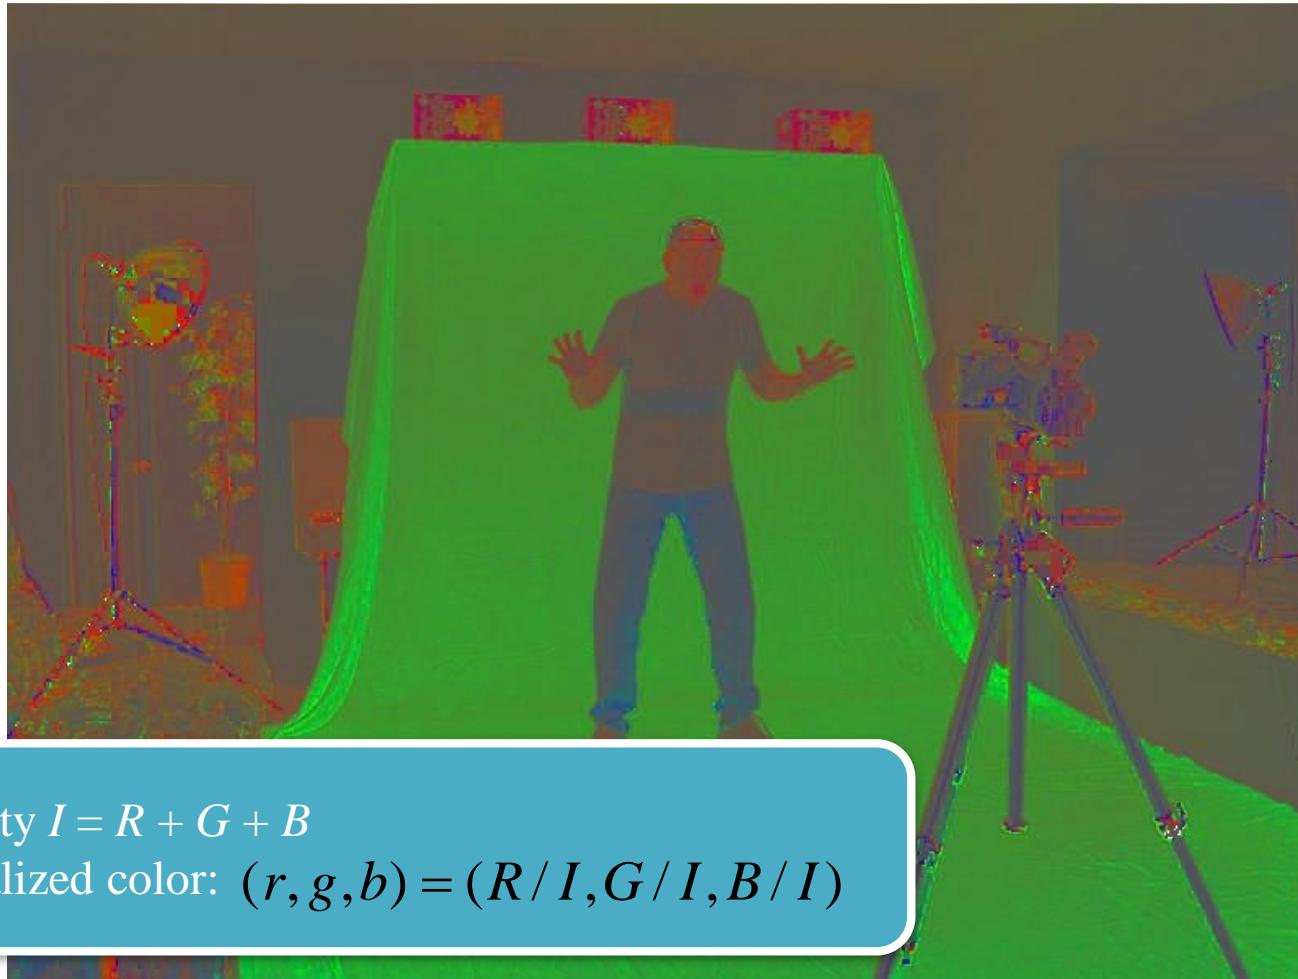
Green Variation in Reality



3D RGB
plot of just
BG

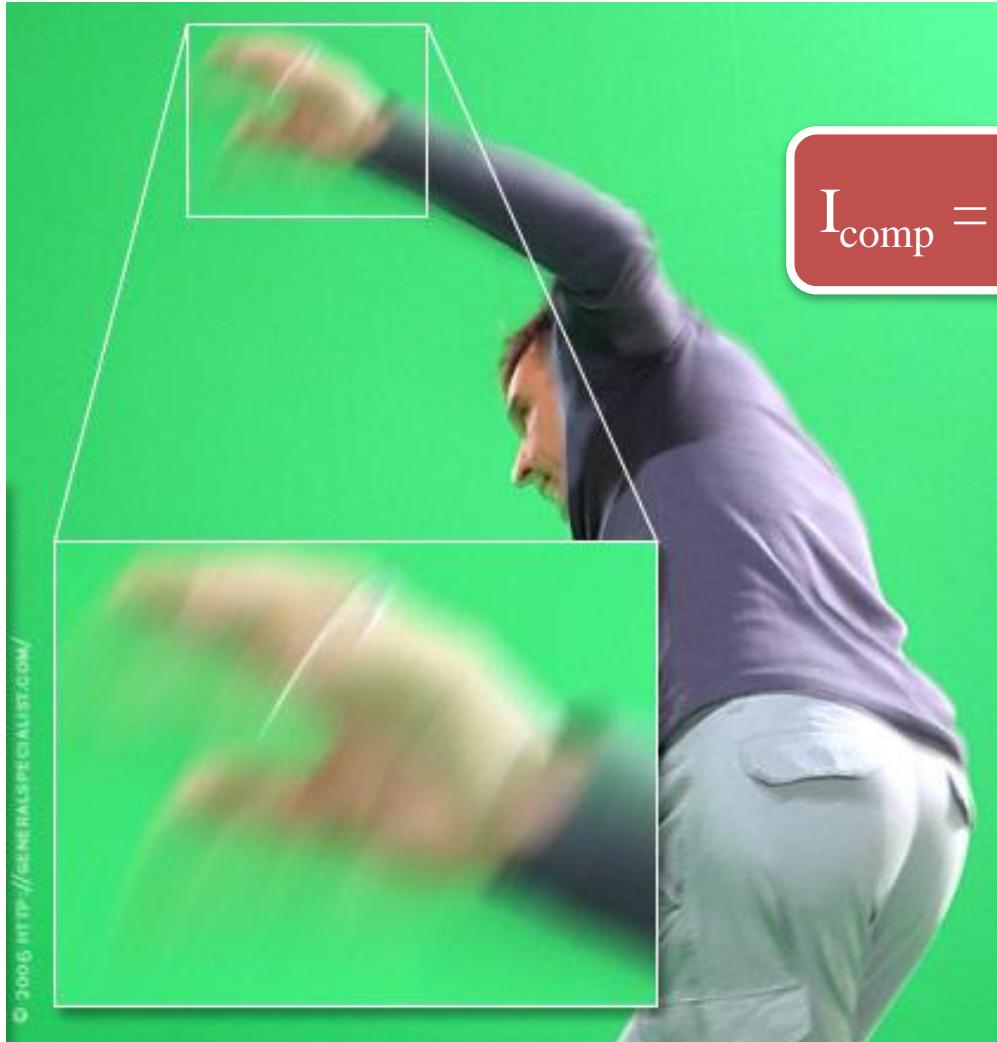
Making of “42 Story House”

Green Variation in Reality



Making of “42 Story House”

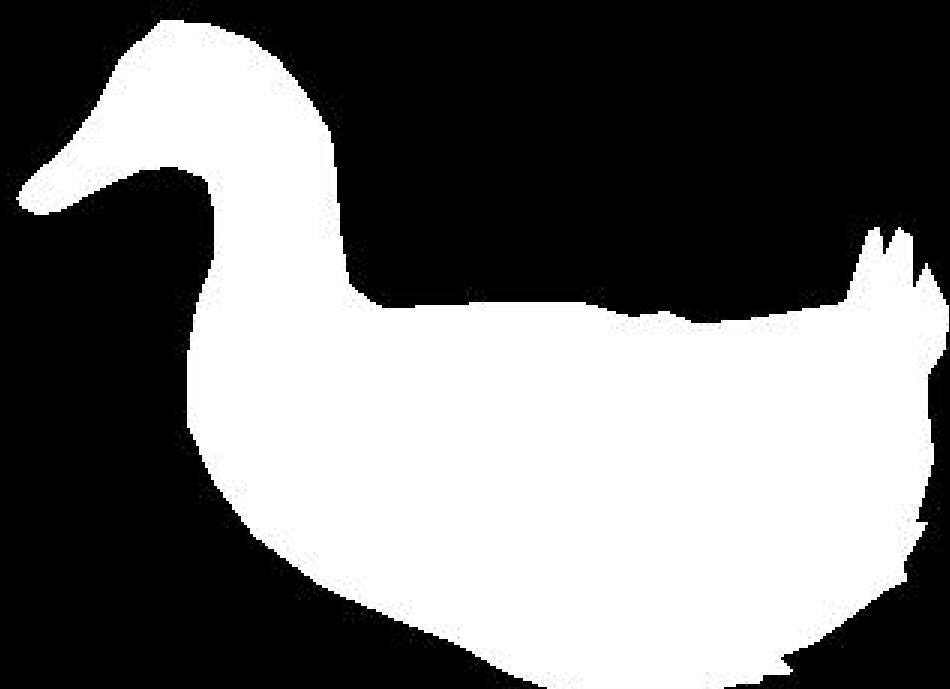
Mixed Pixels: Why I_α Should be Grayscale



Segmentation Performance

- To use automatic analysis, one needs to know the true classification of each test
- We need to do the segmentation by hand on some example images...

Ground truth



ROC Analysis

(ROC = Receiver operating characteristic)

- An ROC curve characterizes the performance of a binary classifier.
- A binary classifier distinguishes between two different types of thing. E.g.:
 - Healthy/afflicted patients – cancer screening
 - Pregnancy tests
 - Object detection
 - Foreground/background image pixels

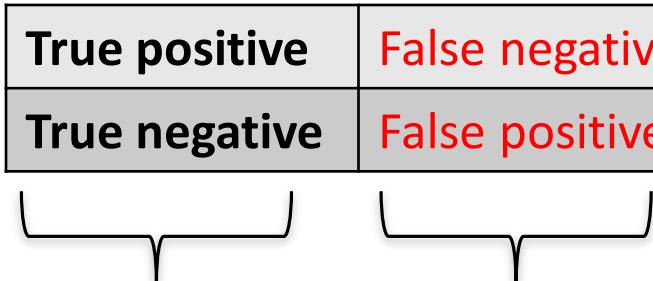
Classification Error

- Binary classifiers make errors
- Two types of input to a binary classifier:
 - Positives
 - Negatives
- Four possible outcomes in any test:

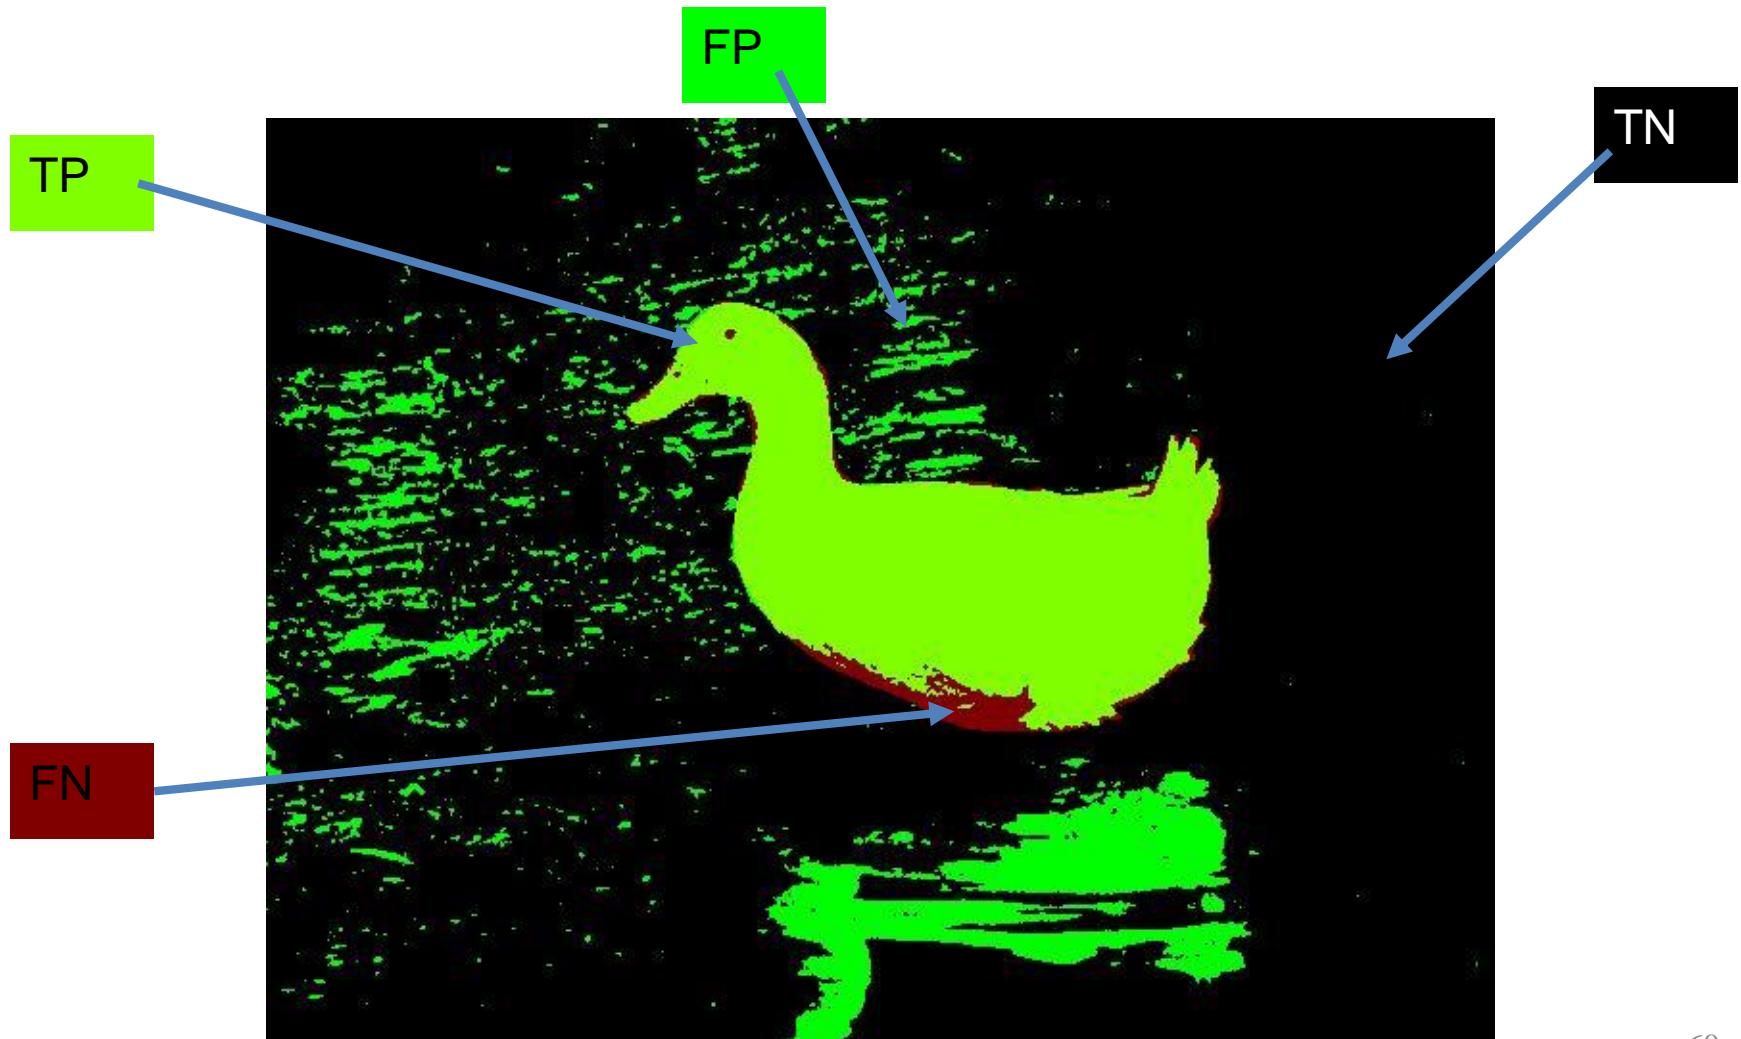
True positive	False negative
True negative	False positive

P: Total # positives N: Total # negatives

Classified: correctly incorrectly



Classification outcomes

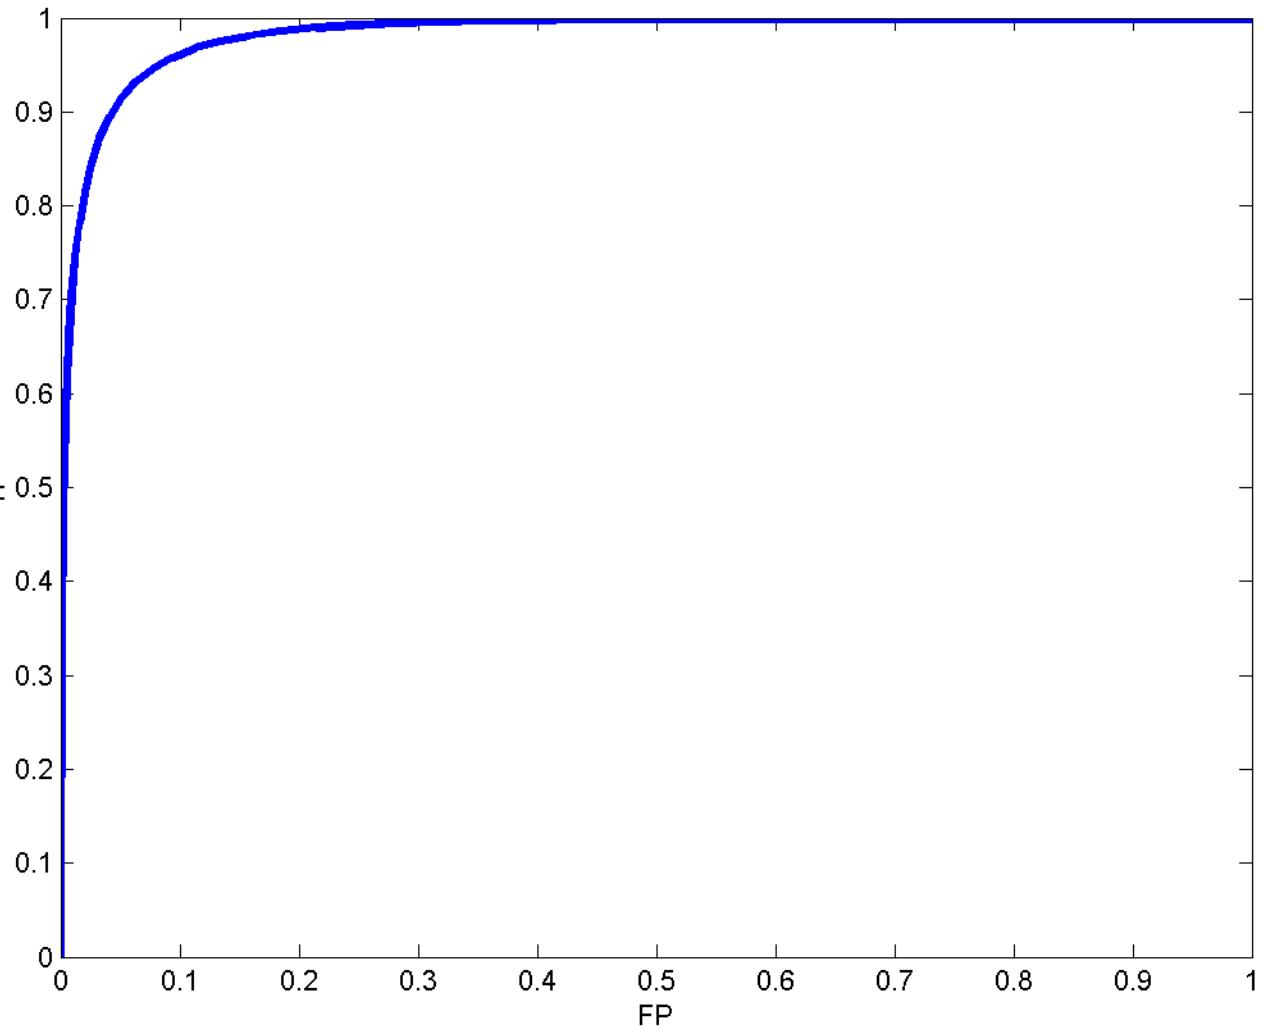


The ROC Curve

- Characterizes the error trade-off in binary classification tasks
- It plots the TP fraction against FP fraction
- TP fraction (*sensitivity*) is $\frac{\text{True positive count}}{P}$
- FP fraction (*1-specificity*) is $\frac{\text{False positive count}}{N}$

The ROC Curve

$$\frac{\text{True positive count}}{P} = \frac{TP}{TP + FN}$$



$$\frac{\text{False positive count}}{N} = \frac{FP}{FP + TN}$$

Properties of ROC curves

- An ROC curve always passes through $(0,0)$ and $(1,1)$
- What is the ROC curve of a perfect system?
- What if the ROC curve is a straight line from $(0,0)$ to $(1,1)$?

Threshold Sweep?

pdf

Foreground

Background

FN

FP

error
probability

Gray level

“MAP (Maximum A Posteriori) detector”

pdf

Foreground

Background

error
probability

Gray level

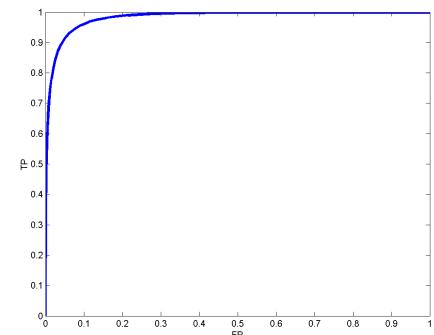
If different outcomes are associated with different costs:
more general “Bayes minimum risk detector”

Operating points

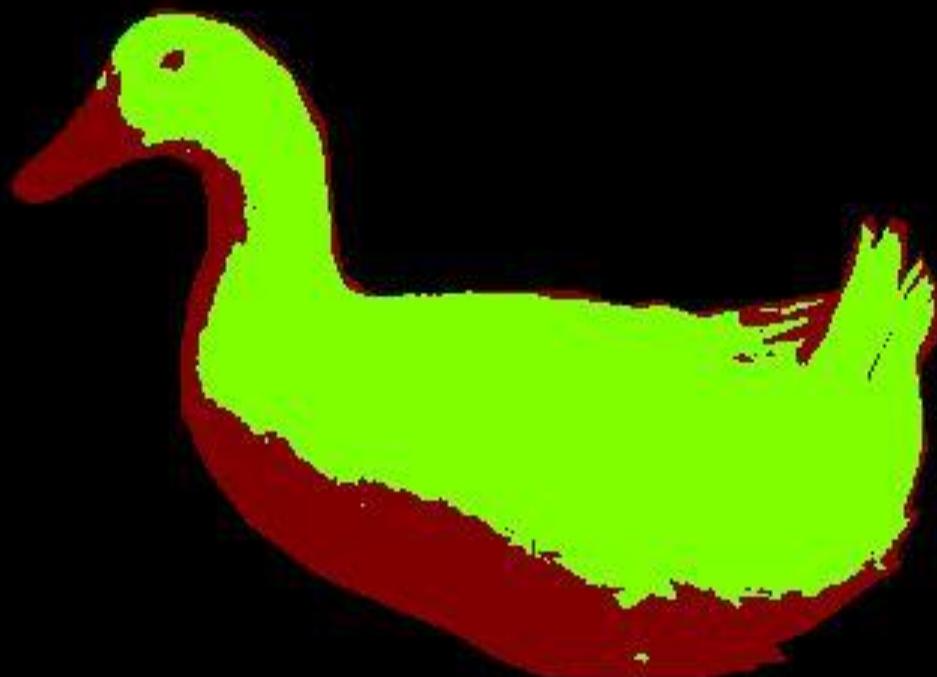
- Choose an *operating point* by assigning relative costs and values to each outcome:
 - V_{TN} – value of true negative
 - V_{TP} – value of true positive
 - C_{FN} – cost of false negative
 - C_{FP} – cost of false positive
- Choose the point on the ROC curve with **gradient**

$$\beta = \frac{N}{P} \frac{V_{TN} + C_{FP}}{V_{TP} + C_{FN}}$$

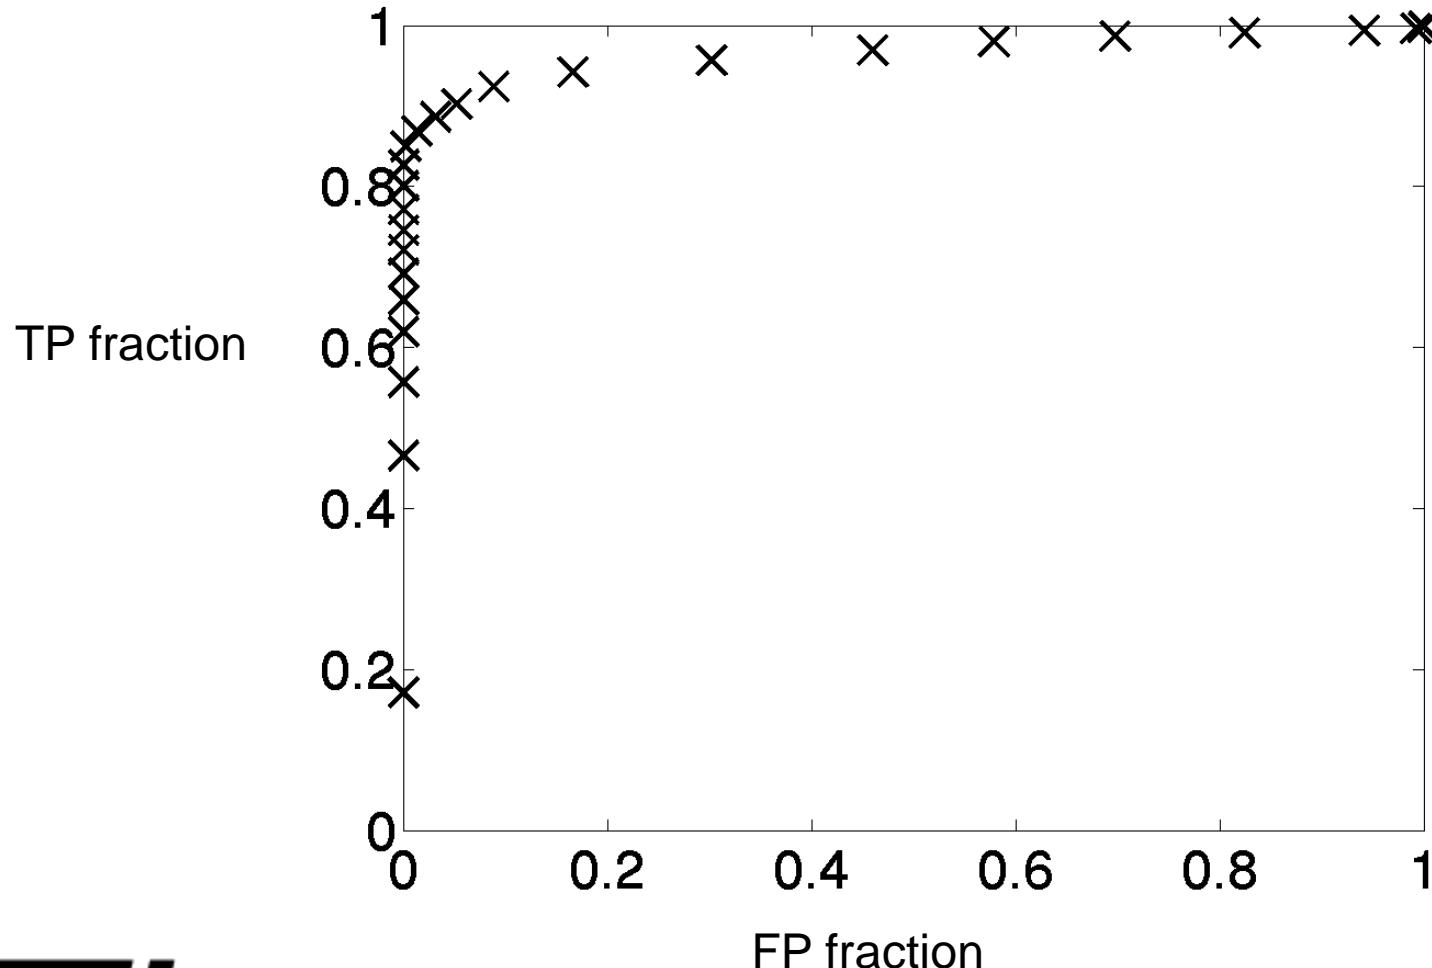
- For simplicity, we often set $V_{TN} = V_{TP} = 0$.



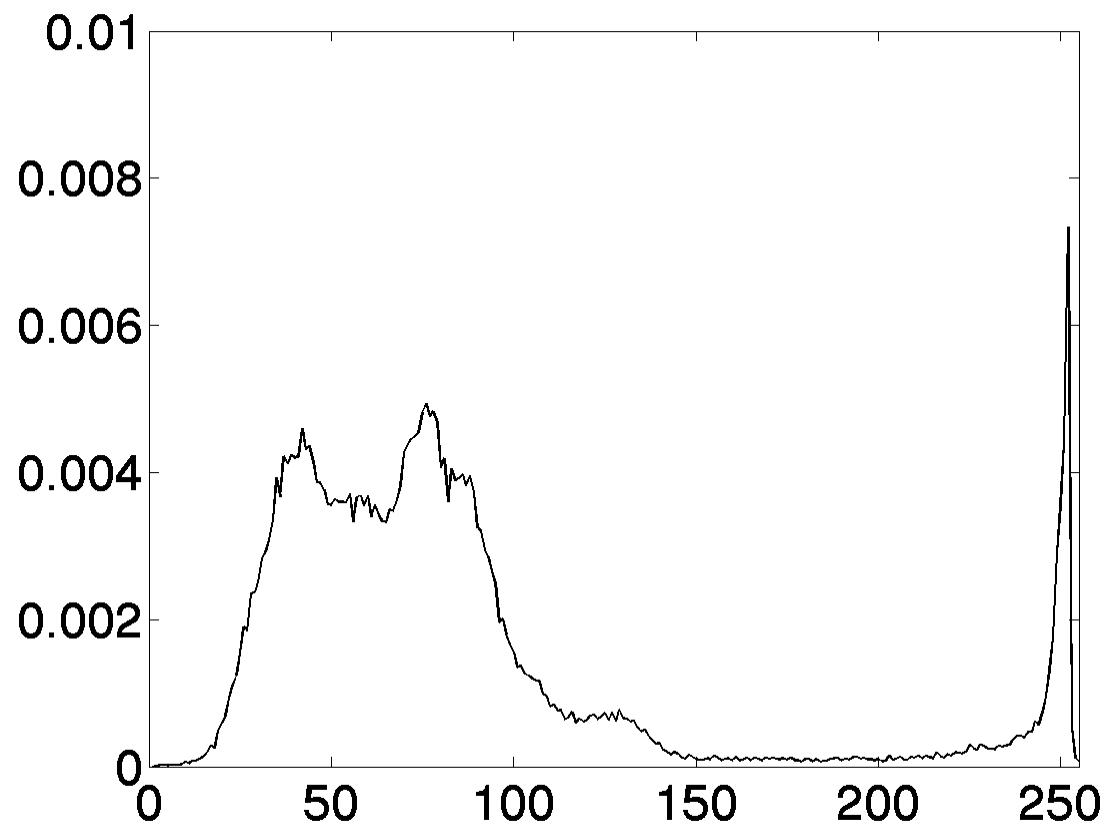
Classification outcomes



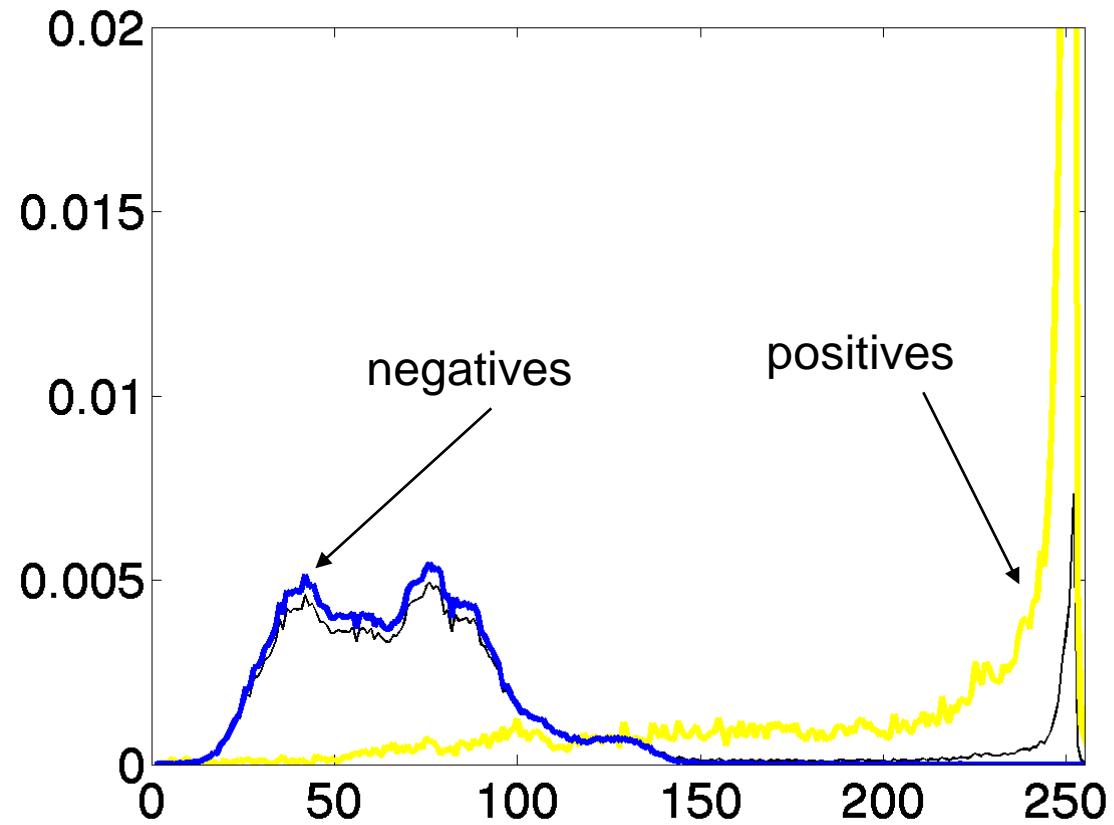
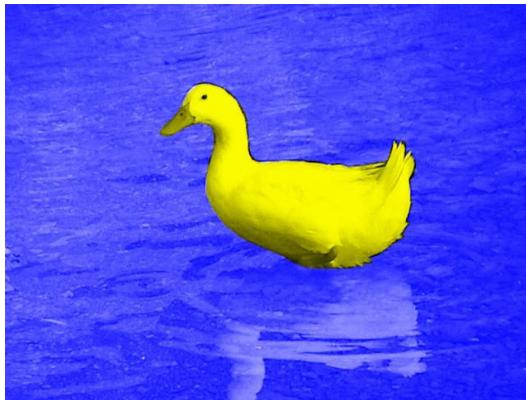
ROC curve



Greylevel Histograms



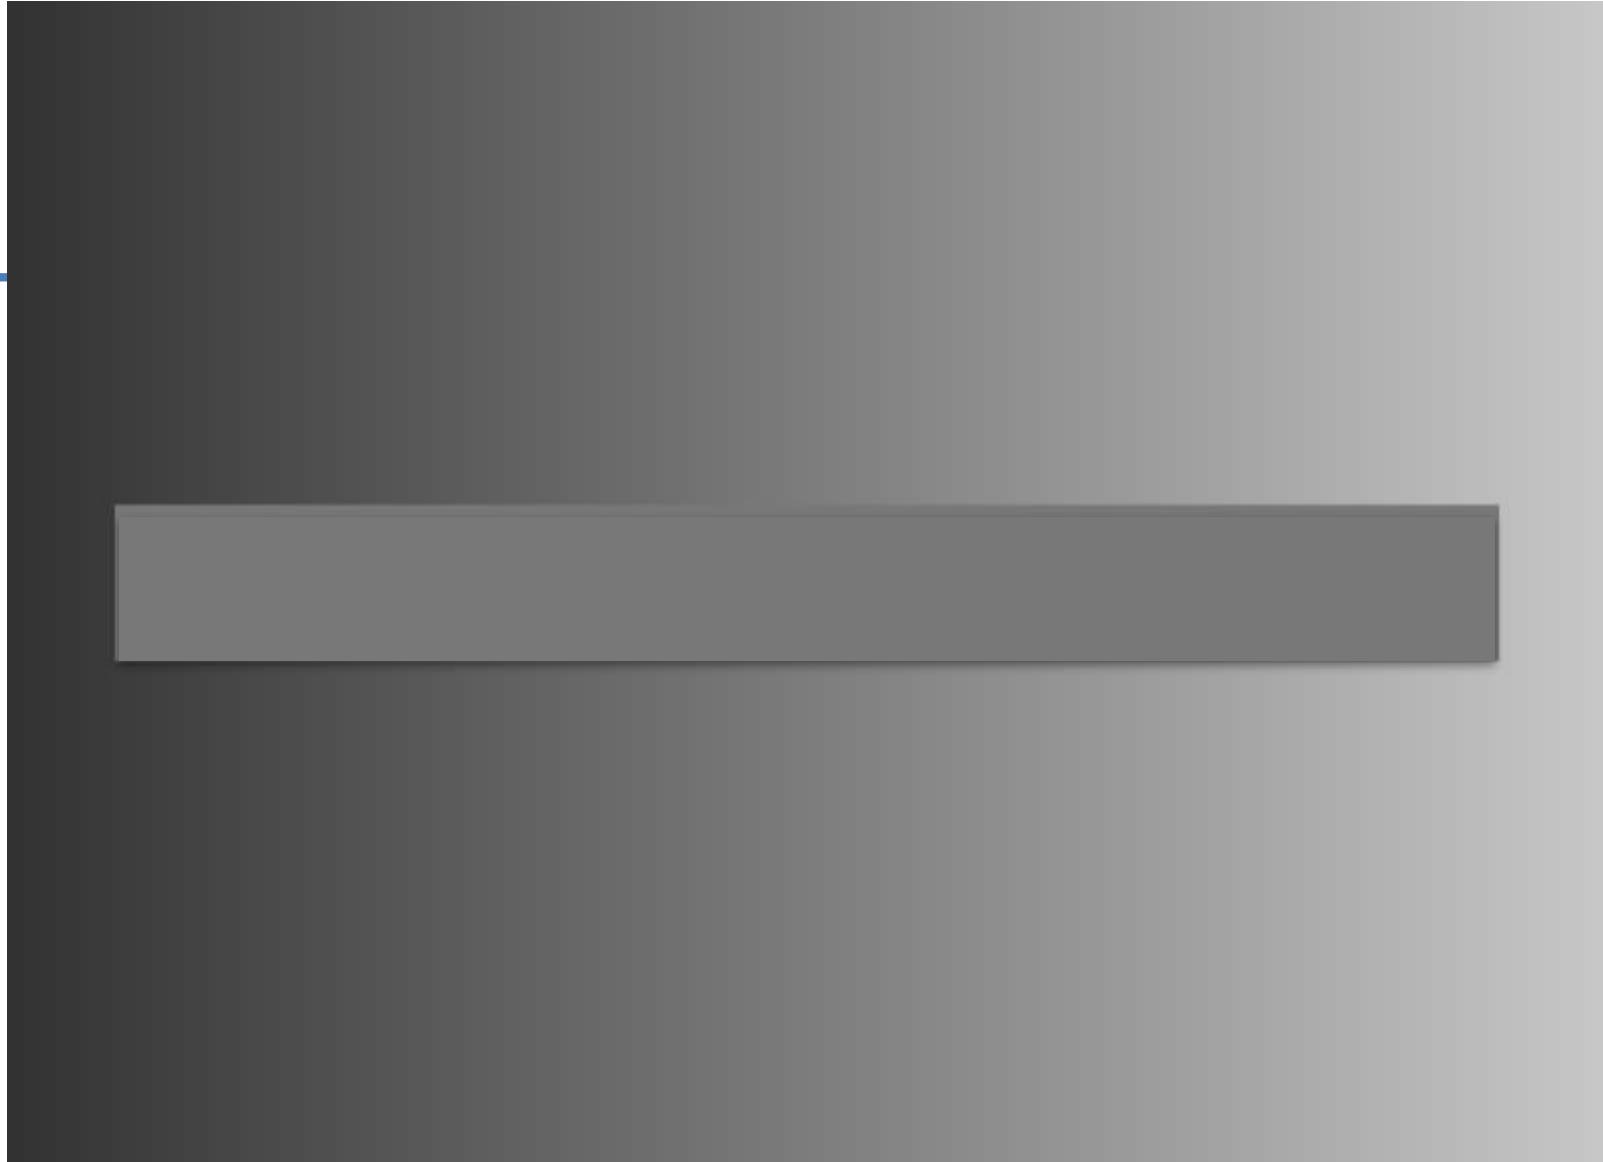
Positives and Negatives



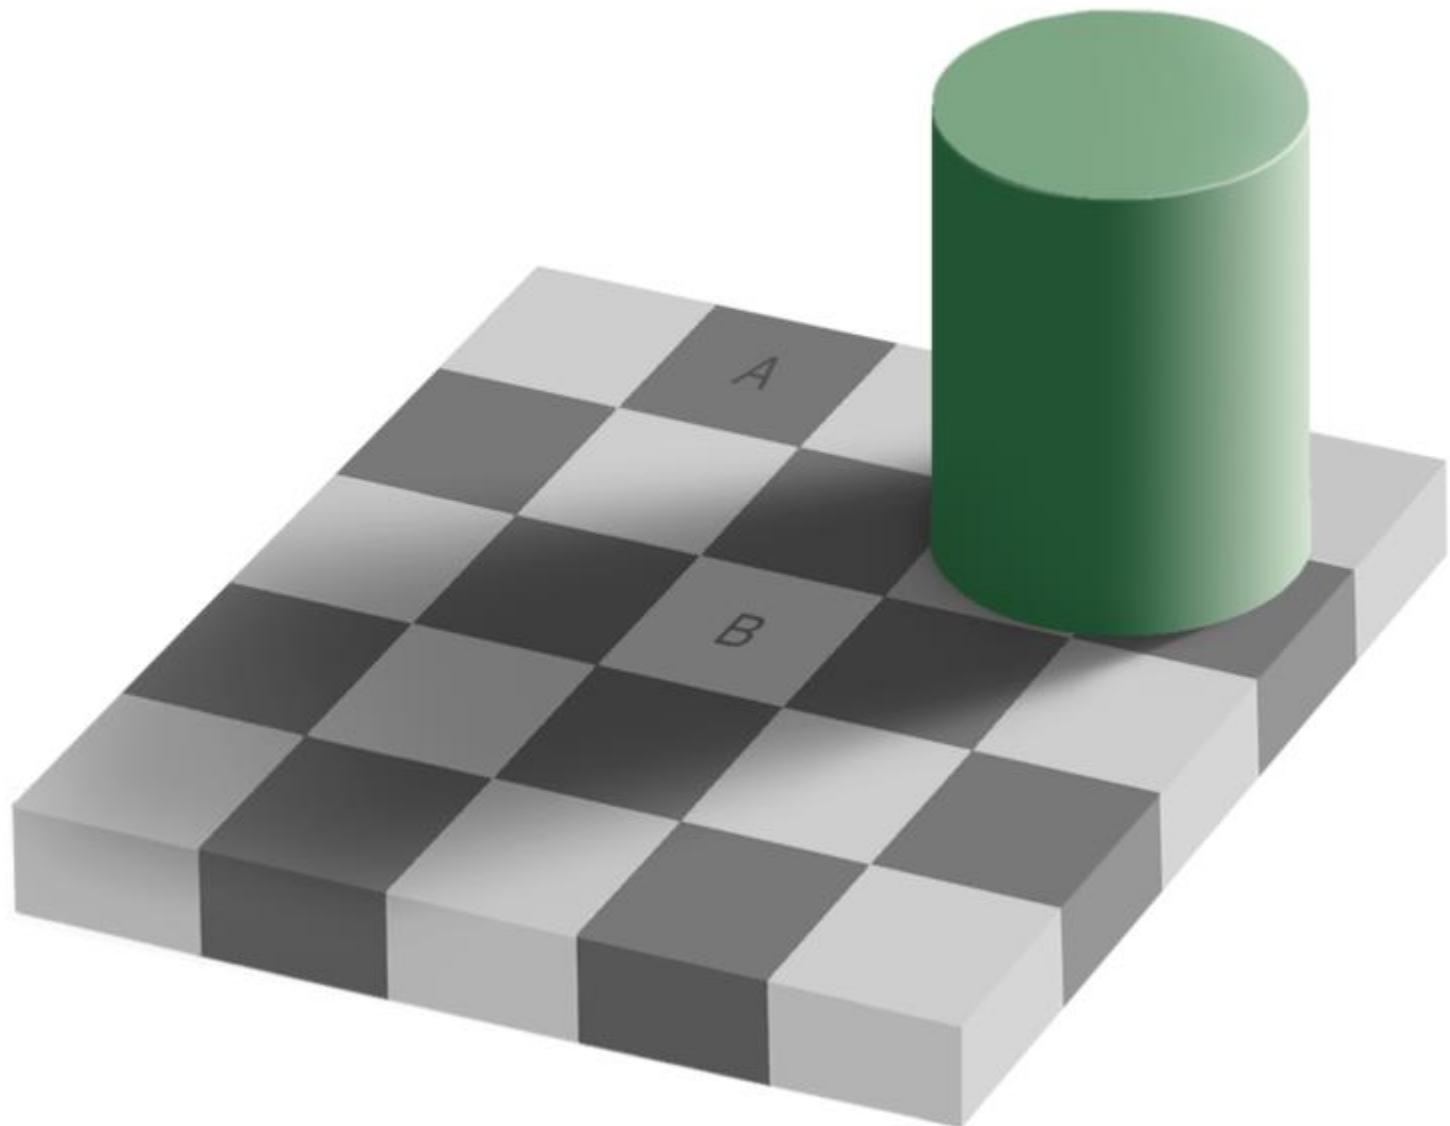
->22/09

Limitations of Thresholding

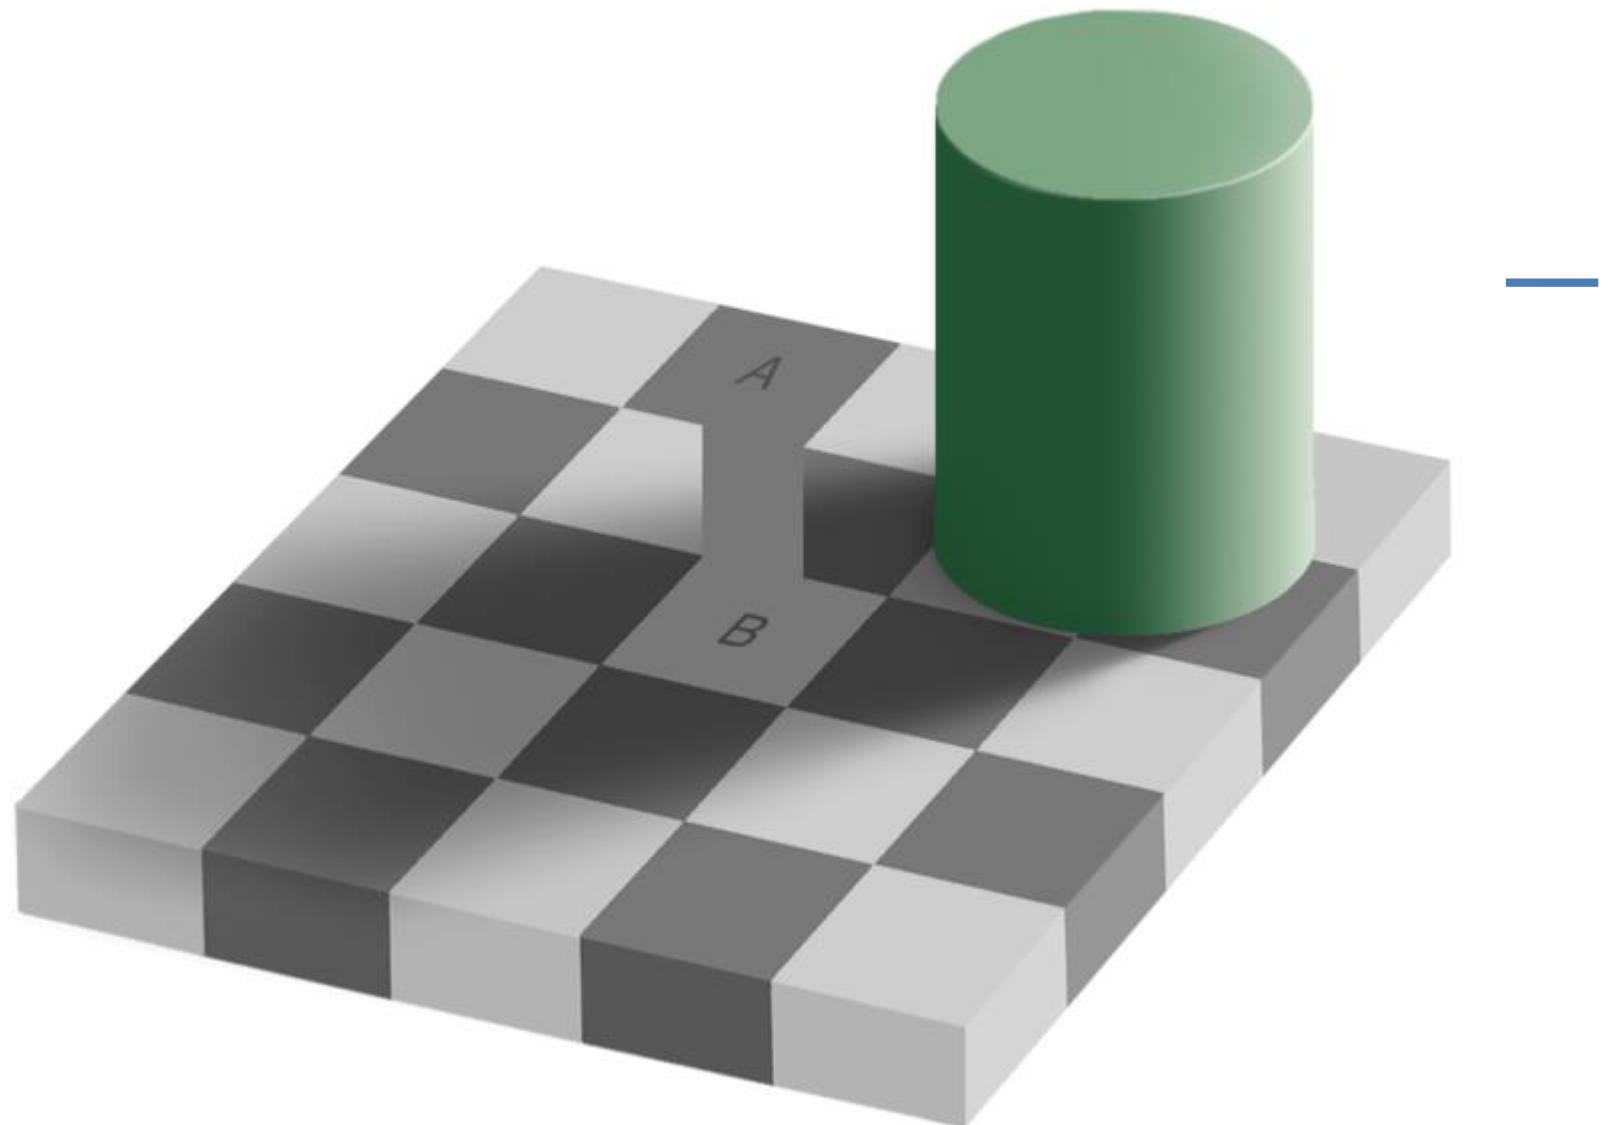
- Why can we segment images much better by eye than through thresholding processes?
- We might improve results by considering image **context**: Surface Coherence



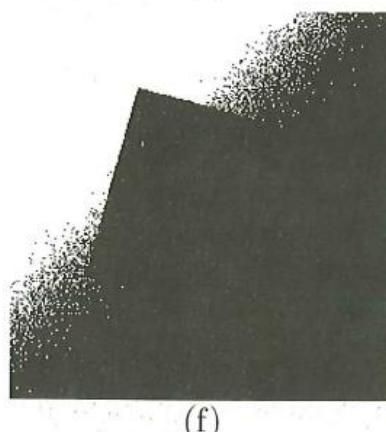
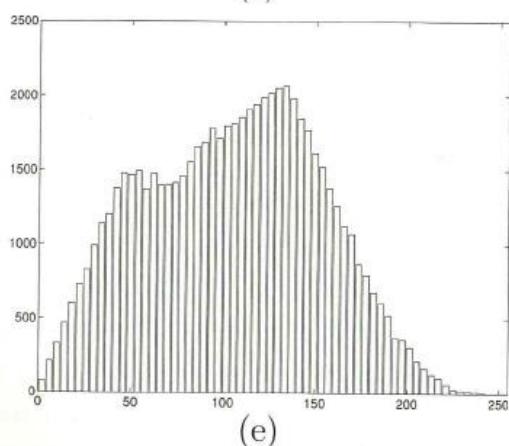
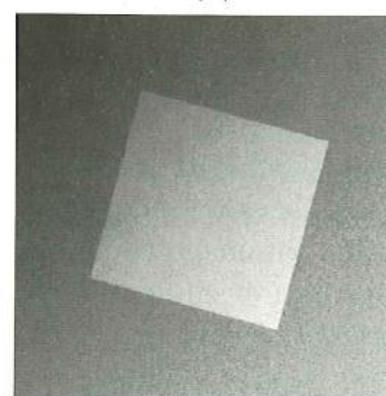
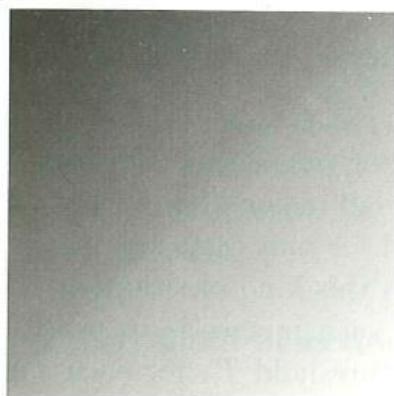
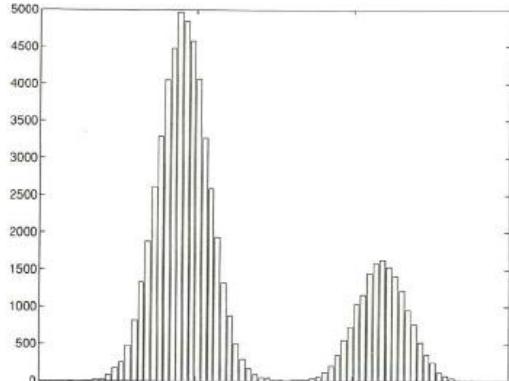
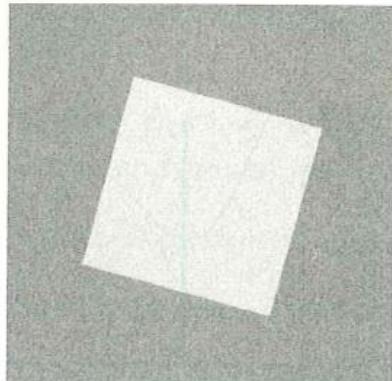
[Gradient.illusion.jpg](#)

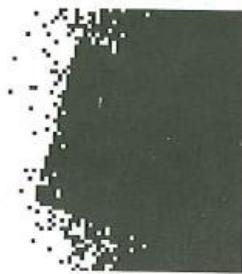
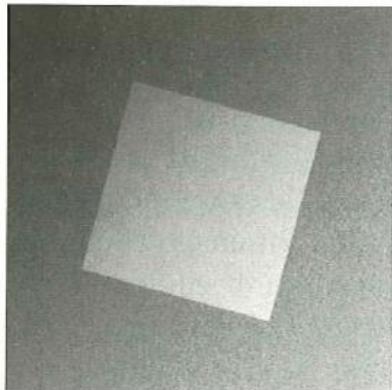


by [Adrian Pingstone](#), based on the [original](#) created by Edward H. Adelson

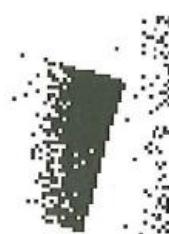


by [Adrian Pingstone](#), based on the [original](#) created by Edward H. Adelson

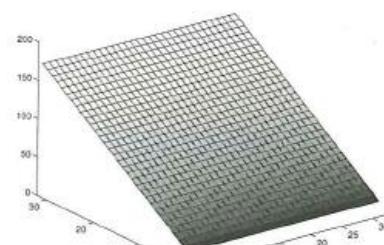




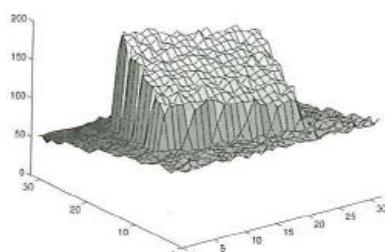
(d)



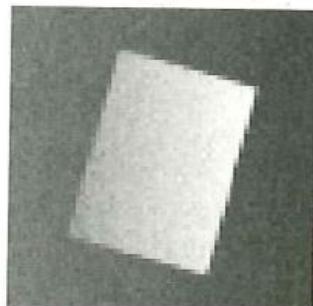
(e)



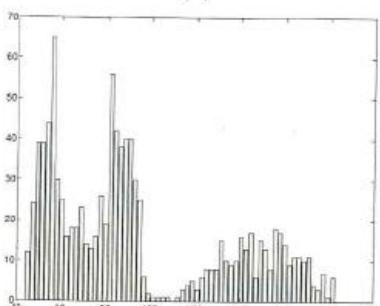
(f)



(g)



(h)



(i)



(j)

Note on Performance Assessment

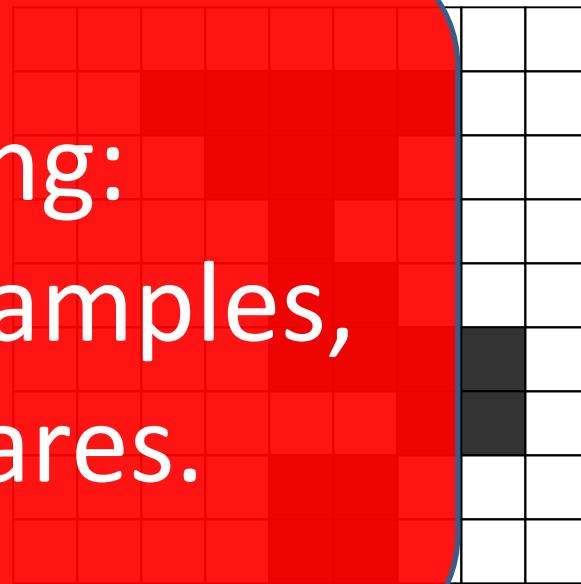
- In real-life, we use two or even three separate sets of test data:
 1. A ***training set***, for tuning the algorithm
 2. A ***validation*** set for tuning the performance score
 3. An unseen ***test set*** to get a final performance score on the tuned algorithm

Pixel connectivity

- We need to define which pixels are neighbors.
- Are the dark pixels in this image connected?

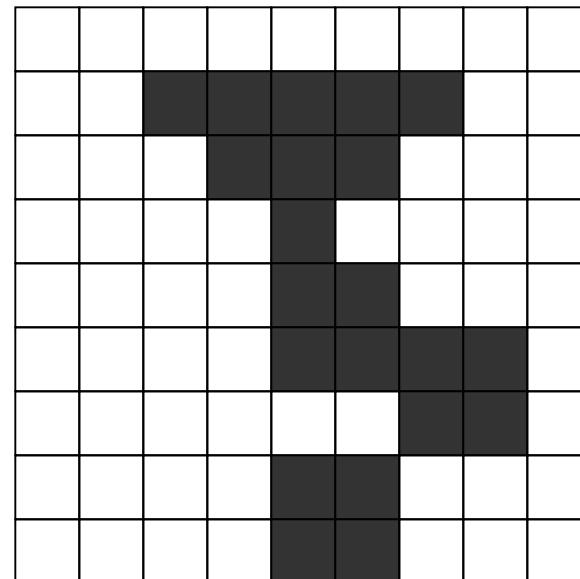
Warning:

**Pixels are samples,
not squares.**

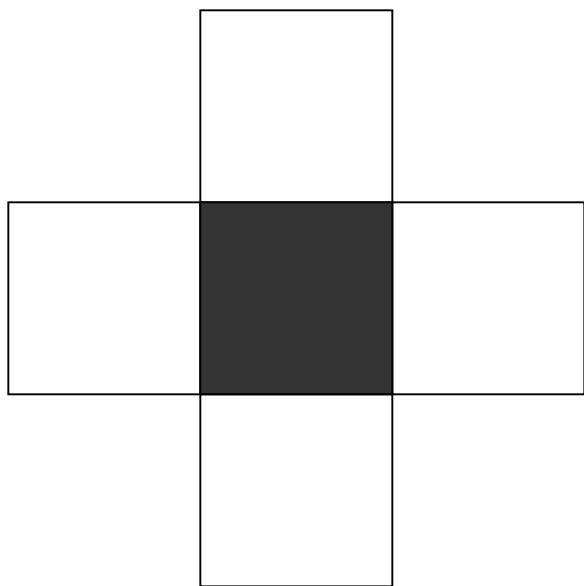


Pixel connectivity

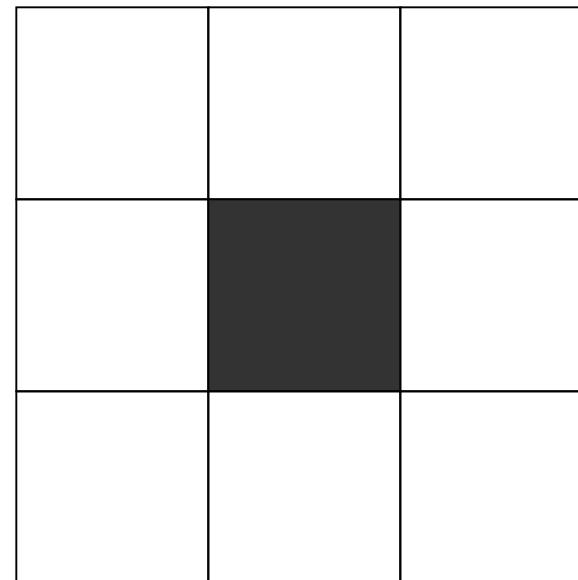
- We need to define which pixels are neighbors.
- Are the dark pixels in this array connected?



Pixel Neighborhoods



4-neighborhood



8-neighborhood

Pixel paths

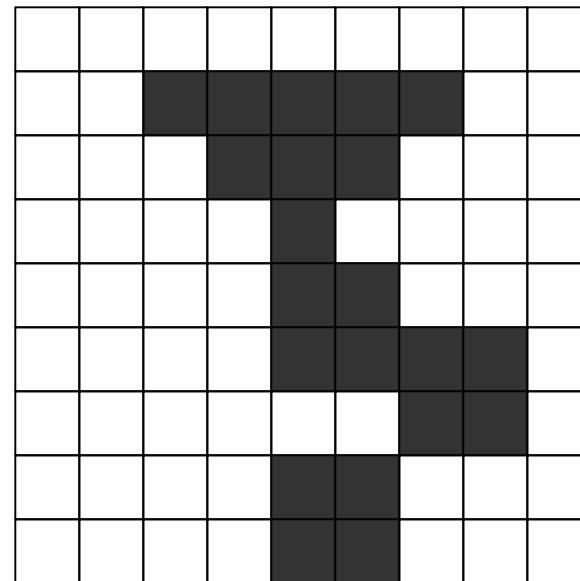
- A 4-connected path between pixels p_1 and p_n is a set of pixels $\{p_1, p_2, \dots, p_n\}$ such that p_i is a 4-neighbor of p_{i+1} , $i=1,\dots,n-1$.
- In an 8-connected path, p_i is an 8-neighbor of p_{i+1} .

Connected regions

- A region is 4-connected if it contains a 4-connected path between any two of its pixels.
- A region is 8-connected if it contains an 8-connected path between any two of its pixels.

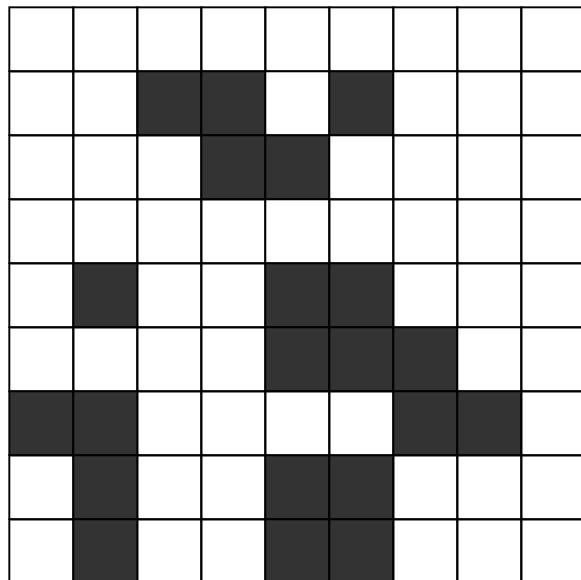
Connected regions

- Now what can we say about the dark pixels in this array?
- What about the light pixels?



Connected components labelling

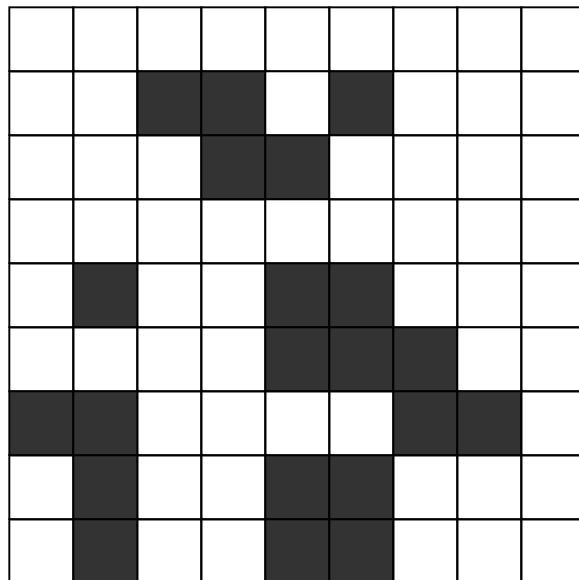
- Labels each connected component of a binary image with a separate number.



1	1	1	1	1	1	1	1	1	1
1	1	2	2	1	3	1	1	1	1
1	1	1	2	2	1	1	1	1	1
1	1	1	1	1	1	1	1	1	1
1	4	1	1	5	5	1	1	1	1
1	1	1	1	5	5	5	1	1	1
6	6	1	1	1	1	5	5	5	1
7	6	1	1	8	8	1	1	1	1
7	6	1	1	8	8	1	1	1	1

Foreground labelling

- Only extract the connected components of the foreground

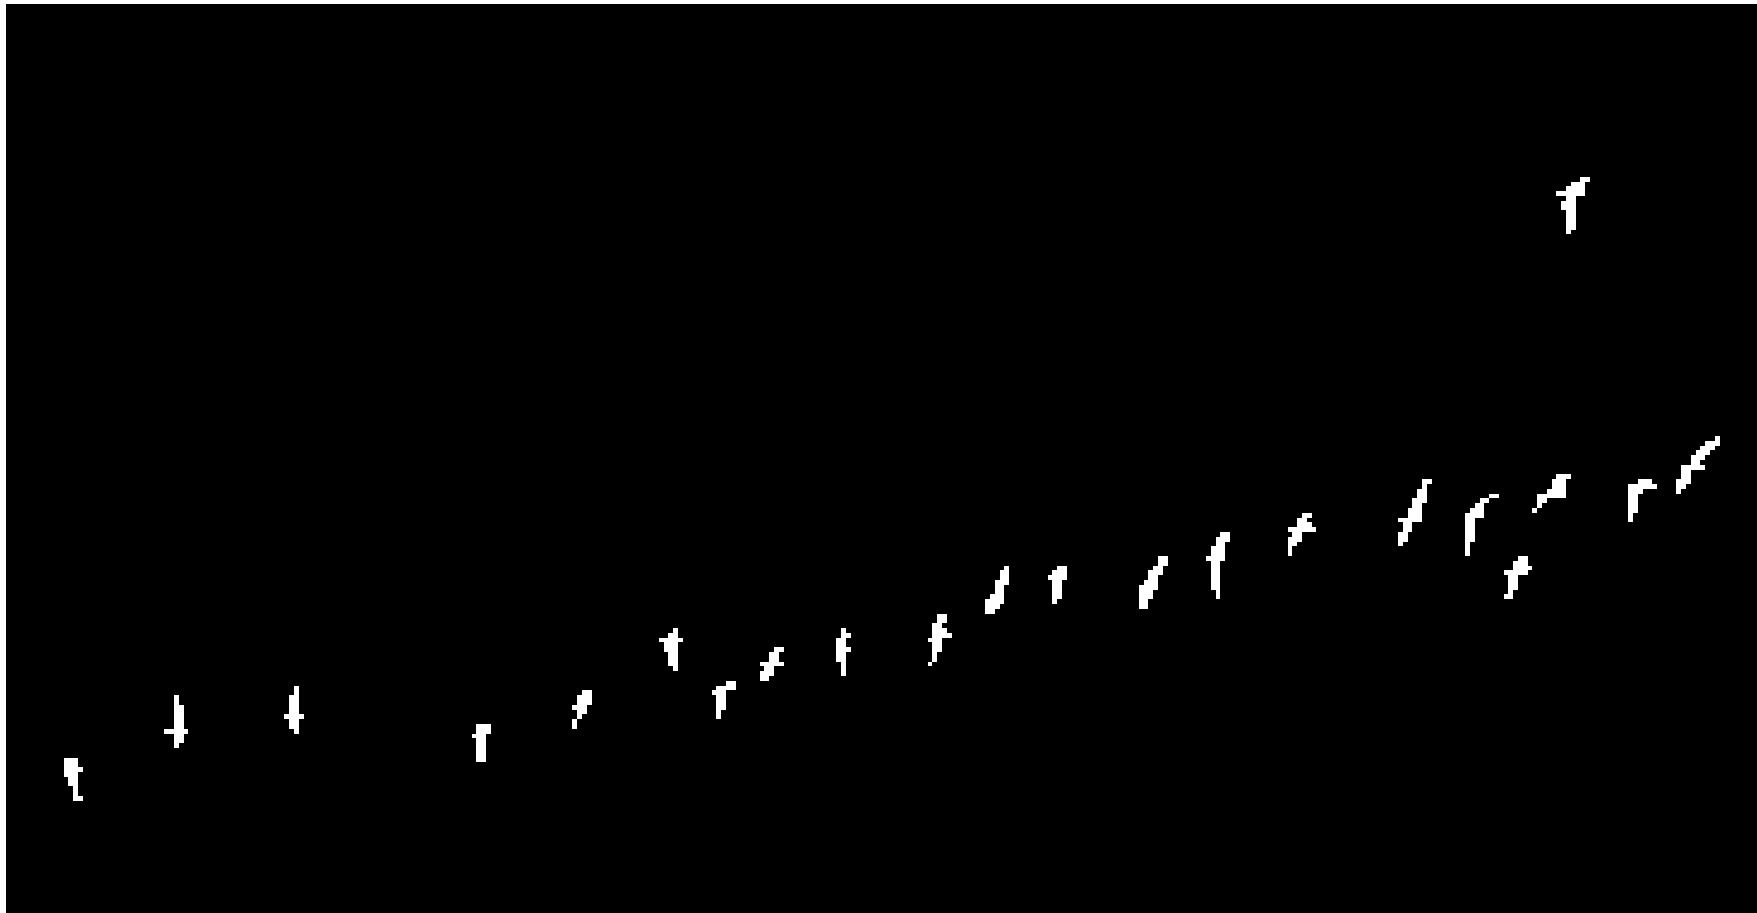


1	1	1	1	1	1	1	1	1	1
1	1	0	0	1	0	1	1	1	1
1	1	1	0	0	1	1	1	1	1
1	1	1	1	1	1	1	1	1	1
1	0	1	1	0	0	1	1	1	1
1	1	1	1	0	0	0	1	1	1
0	0	1	1	1	1	0	0	0	1
2	0	1	1	0	0	1	1	1	1
2	0	1	1	0	0	1	1	1	1

Goose detector



Goose detector



Region Growing

- Start from a seed point or region.
- Add neighboring pixels that satisfy the criteria defining a region.
- Repeat until we can include no more pixels.

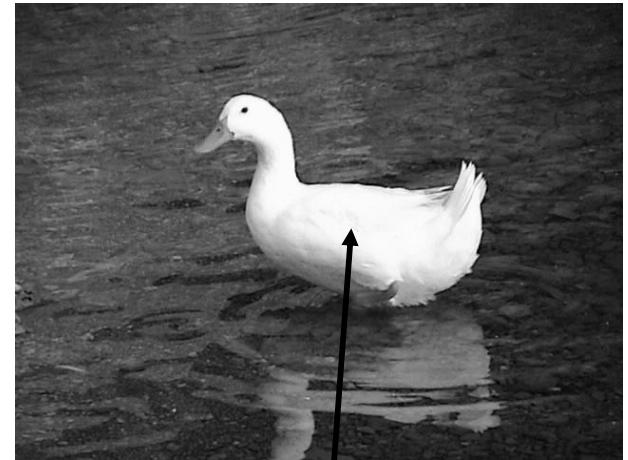
Region Growing

```
def regionGrow(I, seed):  
    X, Y = I.shape  
    visited = np.zeros((X,Y))  
    visited[seed] = 1  
    boundary = []  
    boundary.append(seed)  
    while len(boundary) > 0:  
        nextPoint = boundary.pop()  
        if include(nextPoint, seed):  
            visited[nextPoint] = 2  
            for (x, y) in neighbors(nextPoint):  
                if visited[x,y] == 0:  
                    boundary.append((x, y))  
                    visited[x,y] = 1
```

Region Growing example

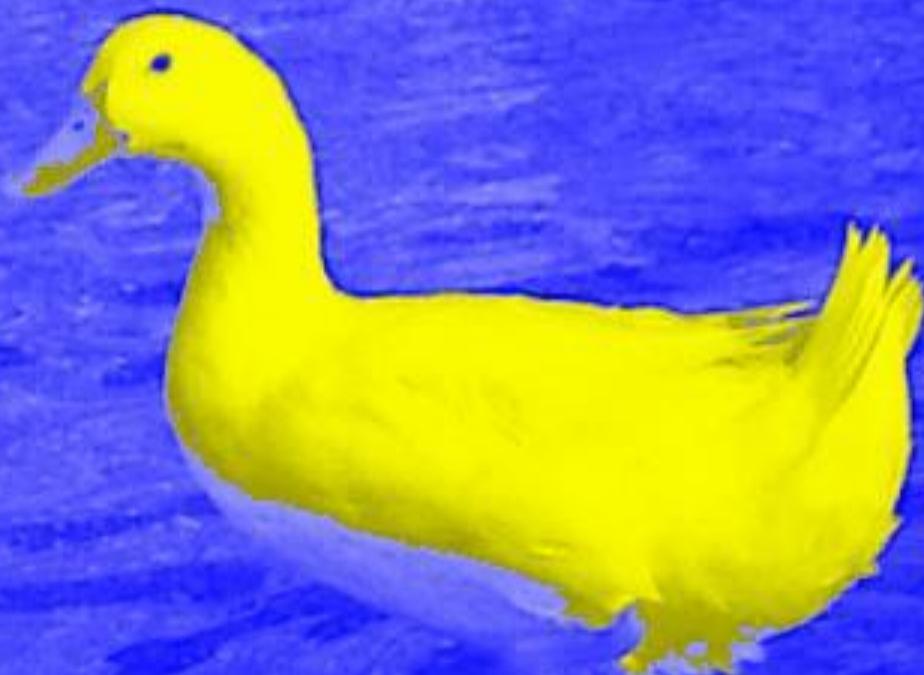
- Pick a single seed pixel
- Inclusion test is up to you:

```
def include(p, seed) :  
    test = ??  
    return test
```



Seed pixel

$T=150$



Variations

- Seed selection
- Inclusion criteria
- Boundary constraints and snakes

Seed selection

- Point and click seed point
- Seed region
 - By hand
 - Automatically, e.g., from a conservative thresholding.
- Multiple seeds
 - Automatically labels the regions

Inclusion criteria

- Greylevel thresholding
- Greylevel distribution model
 - Use mean μ and standard deviation σ in seed region:
 - Include if $(I(x, y) - \mu)^2 < (n\sigma)^2$. Eg: $n = 3$.
 - Can update the mean and standard deviation after every iteration.
- Color or texture information

Snakes

- A snake is an *active contour*
- It is a polygon, i.e., an ordered set of points joined up by lines
- Each point on the contour moves away from the seed while its image neighborhood satisfies an inclusion criterion
- Often the contour has smoothness constraints

Snakes

- The algorithm iteratively minimizes an energy function:
- $E = E_{\text{tension}} + E_{\text{stiffness}} + E_{\text{image}}$
- See Kass, Witkin, Terzopoulos, IJCV 1988

Example



Interim Summary

- Segmentation is hard
- But it is easier if you define the task carefully
 - Is the segmentation task binary or continuous?
 - What are the regions of interest?
 - How accurately must the algorithm locate the region boundaries?
- Research problems remain!

Thursday:
More segmentation