


Visual Computing: Optical Flow

Prof. Marc Pollefeys

Last lecture

- DCT (JPEG) and DWT (JPEG2000) compression



- JPEG2000 encoding pipeline

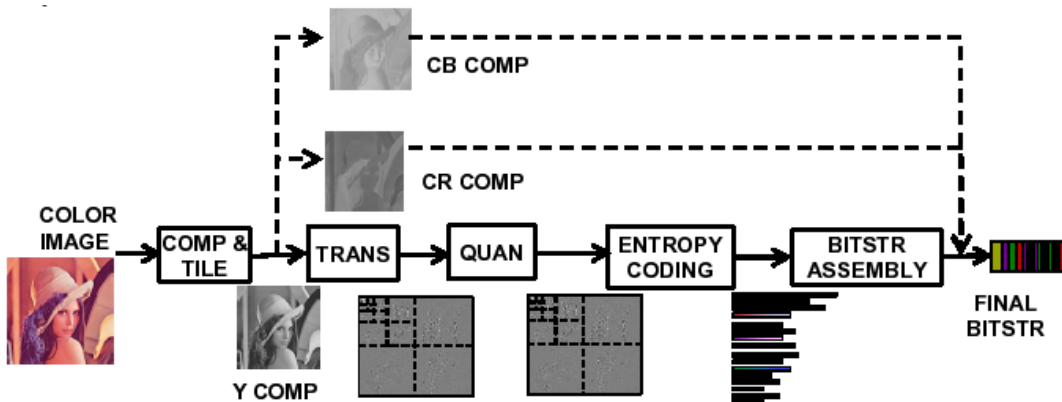




Figure 2 Operation flow of the JPEG 2000 standard.

multi-resolution analysis



biortogonal filters



Visual Computing: Optical Flow

Prof. Marc Pollefeys

Optical Flow


- Brightness Constancy
- The Aperture problem
- Regularization
- Lucas-Kanade
- Coarse-to-fine
- Parametric motion models
- Direct depth
- SSD tracking
- Robust flow
- Bayesian flow

Optical Flow: Where do pixels move to?



Motion is a basic cue

Motion can be the only cue for segmentation



Motion is a basic cue

Even impoverished motion data can elicit a strong percept



Applications

- tracking
- structure from motion
- motion segmentation
- stabilization
- compression
- Mosaicing
- ...

Optical Flow

- Brightness Constancy
- The Aperture problem
- Regularization
- Lucas-Kanade
- Coarse-to-fine
- Parametric motion models
- Direct depth
- SSD tracking
- Robust flow
- Bayesian flow

Definition of Optical Flow

OPTICAL FLOW = apparent motion of
brightness patterns

Ideally, the optical flow is the projection of the three-dimensional velocity vectors on the image

Caution required

Two examples :

1. Uniform, rotating sphere



$$\text{O.F.} = 0$$

2. No motion, but changing lighting



$$\text{O.F.} \neq 0$$

Caution required



Mathematical formulation

$I(x, y, t)$ = brightness at (x, y) at time t

Brightness constancy assumption:

$$I\left(x + \frac{dx}{dt} \delta t, y + \frac{dy}{dt} \delta t, t + \delta t\right) = I(x, y, t)$$

Optical flow constraint equation :

$$\frac{dI}{dt} = \frac{\partial I}{\partial x} \frac{dx}{dt} + \frac{\partial I}{\partial y} \frac{dy}{dt} + \frac{\partial I}{\partial t} = 0$$

Optical Flow

- Brightness Constancy
- The Aperture problem
- Regularization
- Lucas-Kanade
- Coarse-to-fine
- Parametric motion models
- Direct depth
- SSD tracking
- Robust flow
- Bayesian flow

The aperture problem

$$u = \frac{dx}{dt}, \quad v = \frac{dy}{dt}$$

$$I_x = \frac{\partial I}{\partial x}, \quad I_y = \frac{\partial I}{\partial y}, \quad I_t = \frac{\partial I}{\partial t}$$

$$I_x u + I_y v + I_t = 0$$

1 equation in 2 unknowns

The aperture problem






Figure 12-4. Local information on the brightness gradient and the rate of change of brightness with time provides only one constraint on the components of the optical flow vector. The flow velocity has to lie along a straight line perpendicular to the direction of the brightness gradient. We can only determine the component in the direction of the brightness gradient. Nothing is known about the flow component in the direction at right angles.

Aperture problem and Normal Flow




The aperture problem



Remarks



Apparently an aperture problem




What is Optic Flow, anyway?

- Estimate of observed projected motion field
- Not always well defined!
- Compare:
 - Motion Field (or Scene Flow)
projection of 3-D motion field
 - Normal Flow
observed tangent motion
 - Optic Flow
apparent motion of the brightness pattern
(hopefully equal to motion field)
- Consider Barber pole illusion



Planar motion examples

- Ideal motion of a plane



Scene Flow: \rightarrow

Normal Flow: undef

Optic Flow: ?, probably 0

Optical Flow

- Brightness Constancy
- The Aperture problem
- **Regularization**
- Lucas-Kanade
- Coarse-to-fine
- Parametric motion models
- Direct depth
- SSD tracking
- Robust flow
- Bayesian flow

Horn & Schunck algorithm

Additional smoothness constraint :

$$e_s = \iint ((u_x^2 + u_y^2) + (v_x^2 + v_y^2)) dx dy,$$

besides OF constraint equation term

$$e_c = \iint (I_x u + I_y v + I_t)^2 dx dy,$$

minimize $e_s + \lambda e_c$

Horn & Schunck

The Euler-Lagrange equations : $F_u - \frac{\partial}{\partial x} F_{u_x} - \frac{\partial}{\partial y} F_{u_y} = 0$

$$F_v - \frac{\partial}{\partial x} F_{v_x} - \frac{\partial}{\partial y} F_{v_y} = 0$$

In our case ,

$$F = (u_x^2 + u_y^2) + (v_x^2 + v_y^2) + \lambda(I_x u + I_y v + I_t)^2,$$

so the Euler-Lagrange equations are

$$\Delta u = \lambda(I_x u + I_y v + I_t)I_x,$$

$$\Delta v = \lambda(I_x u + I_y v + I_t)I_y,$$

$$\Delta = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \quad \text{is the Laplacian operator}$$

Horn & Schunk

Remarks :

1. Coupled PDEs solved using iterative methods and finite differences

$$\frac{\partial u}{\partial t} = \Delta u - \lambda(I_x u + I_y v + I_t)I_x,$$

$$\frac{\partial v}{\partial t} = \Delta v - \lambda(I_x u + I_y v + I_t)I_y,$$





2. More than two frames allow a better estimation of I_t
3. Information spreads from corner-type patterns



Horn & Schunk, remarks

1. Errors at boundaries
2. Example of *regularisation*
(selection principle for the solution of illposed problems)

Results of an enhanced system



Structure from motion with OF



Optical Flow

- Brightness Constancy
- The Aperture problem
- Regularization
- **Lucas-Kanade**
- Coarse-to-fine
- Parametric motion models
- SSD tracking
- Bayesian flow

Lucas-Kanade: Integrate over a Patch

Assume a single velocity for all pixels within an image patch

$$E(u, v) = \sum_{x, y \in \Omega} \left(I_x(x, y)u + I_y(x, y)v + I_t \right)^2$$

Solve with:

$$\begin{aligned} \frac{dE(u, v)}{du} &= \sum 2I_x(I_x u + I_y v + I_t) = 0 \\ \frac{dE(u, v)}{dv} &= \sum 2I_y(I_x u + I_y v + I_t) = 0 \end{aligned}$$

$$\begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} \begin{pmatrix} u \\ v \end{pmatrix} = - \begin{pmatrix} \sum I_x I_t \\ \sum I_y I_t \end{pmatrix}$$

On the LHS: sum of the 2x2 outer product
tensor of the gradient vector

$$\left(\sum \nabla I \nabla I^T \right) \vec{U} = - \sum \nabla I I_t$$

Lucas-Kanade: Singularities and the Aperture Problem

$$\text{Let } M = \sum (\nabla I)(\nabla I)^T \quad \text{and} \quad b = \begin{bmatrix} -\sum I_x I_t \\ -\sum I_y I_t \end{bmatrix}$$

- Algorithm: At each pixel compute U by solving $MU=b$
- M is singular if all gradient vectors point in the same direction
 - e.g., along an edge
 - of course, trivially singular if the summation is over a single pixel
 - i.e., only *normal flow* is available (aperture problem)
- Corners and textured areas are OK

KLT feature tracker:

see “Good Features to Track”, *Shi and Tomasi, CVPR'94*, 1994, pp. 593 - 600.


[video](#)

Iterative Refinement

- Estimate velocity at each pixel using one iteration of Lucas and Kanade estimation
- Warp one image toward the other using the estimated flow field
(easier said than done)
- Refine estimate by repeating the process

Motion and Gradients

Consider 1-d signal; assume linear function of x





$$\frac{dI}{dx} = - \frac{dI/dt}{u}$$

$$0 = I_x u + I_t$$


“shift by u to account
for I_x with I_t ”

$$u = - \frac{I_t}{I_x}$$

Iterative refinement



BUT!!



Optical Flow

- Brightness Constancy
- The Aperture problem
- Regularization
- Lucas-Kanade
- **Coarse-to-fine**
- Parametric motion models
- SSD tracking
- Bayesian flow



Limits of the (local) gradient method

1. Fails when intensity structure within window is poor
2. Fails when the displacement is large (typical operating range is motion of 1 pixel per iteration!)
 - *Linearization of brightness is suitable only for small displacements*

Also, brightness is not strictly constant in images


- *actually less problematic than it appears, since we can pre-filter images to make them look similar*

Pyramid / “Coarse-to-fine”



Coarse-to-Fine Estimation

$$I_x \cdot u + I_y \cdot v + I_t \approx 0 \implies \text{small } u \text{ and } v \dots$$

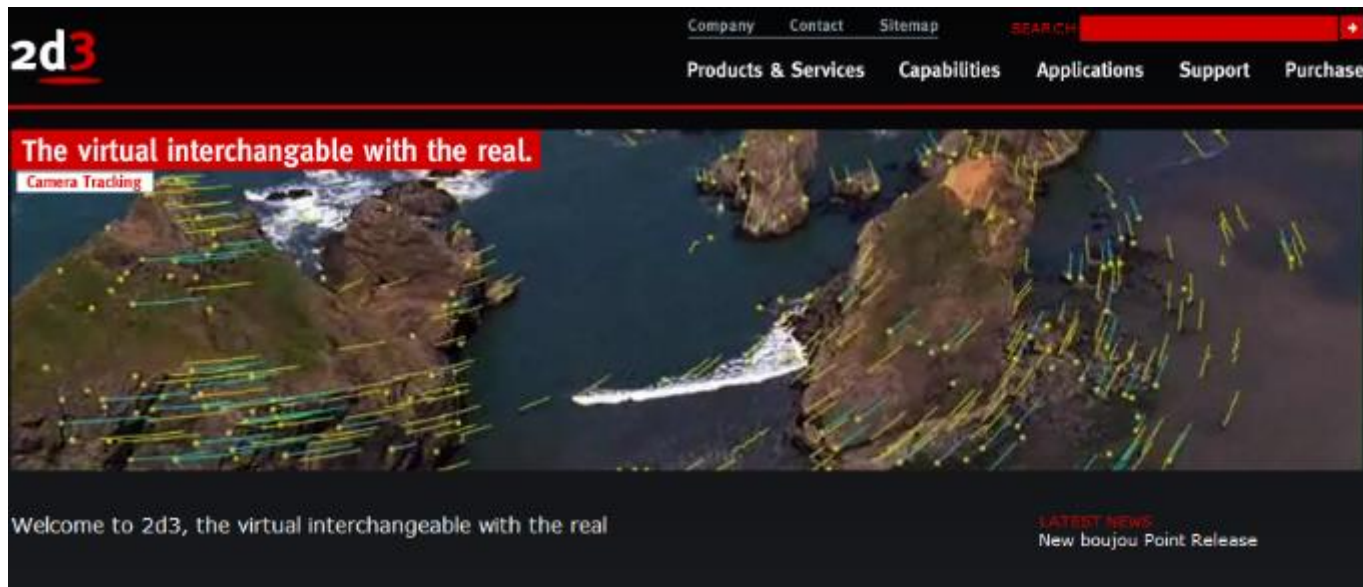


OF application: Image stabilization




[DeShaker](#)

OF application: MatchMoving



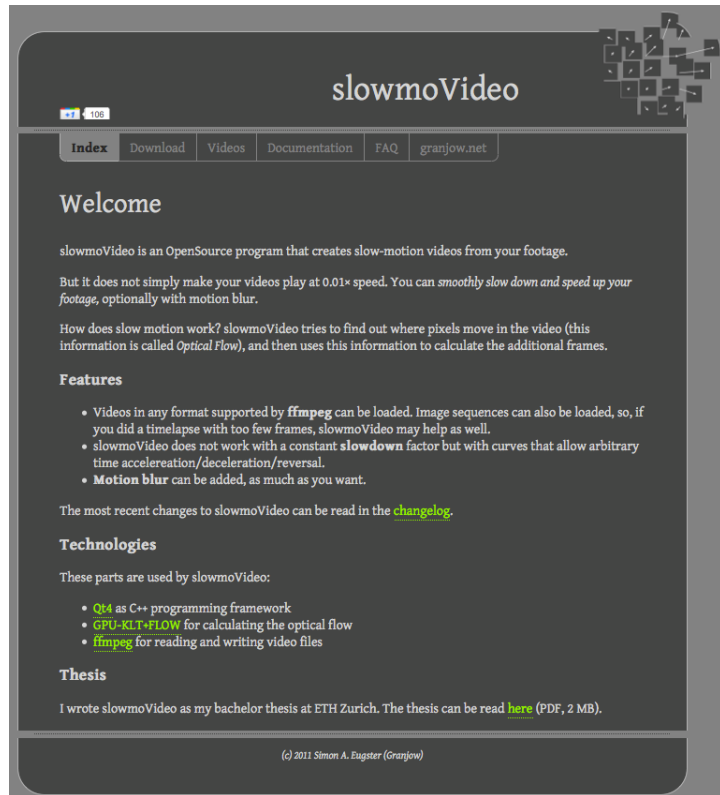
OF application: Slow motion

- Slow motion (generate intermediate frames)
- Technology is also key to 100Hz television



SlowMoVideo

Bachelor thesis Simon Eugster



The screenshot shows the homepage of the slowmoVideo project. At the top, the logo 'slowmoVideo' is displayed next to a small graphic of a film strip. Below the logo is a navigation bar with links: Index, Download, Videos, Documentation, FAQ, and granjow.net. The main content area starts with a 'Welcome' section, followed by a description of the program as an OpenSource tool for creating slow-motion videos. It then lists features such as support for various video formats, the ability to add motion blur, and the use of optical flow. A 'Technologies' section mentions the use of Qt4, GPU-KIT-Flow, and ffmpeg. Finally, a 'Thesis' section mentions the author's bachelor thesis at ETH Zurich. The footer contains the copyright notice: (c) 2011 Simon A. Eugster (Granjow).

slowmoVideo

Index Download Videos Documentation FAQ granjow.net

Welcome

slowmoVideo is an OpenSource program that creates slow-motion videos from your footage.

But it does not simply make your videos play at 0.01x speed. You can *smoothly slow down and speed up your footage*, optionally with motion blur.

How does slow motion work? slowmoVideo tries to find out where pixels move in the video (this information is called *Optical Flow*), and then uses this information to calculate the additional frames.

Features

- Videos in any format supported by **ffmpeg** can be loaded. Image sequences can also be loaded, so, if you did a timelapse with too few frames, slowmoVideo may help as well.
- slowmoVideo does not work with a constant **slowdown** factor but with curves that allow arbitrary time acceleration/deceleration/reversal.
- **Motion blur** can be added, as much as you want.

The most recent changes to slowmoVideo can be read in the [changelog](#).

Technologies

These parts are used by slowmoVideo:

- **Qt4** as C++ programming framework
- **GPU-KIT-FLOW** for calculating the optical flow
- **ffmpeg** for reading and writing video files

Thesis

I wrote slowmoVideo as my bachelor thesis at ETH Zurich. The thesis can be read [here](#) (PDF, 2 MB).

(c) 2011 Simon A. Eugster (Granjow)

<http://slowmovideo.granjow.net/>



Optical Flow

- Brightness Constancy
- The Aperture problem
- Regularization
- Lucas-Kanade
- Coarse-to-fine
- **Parametric motion models**
- SSD tracking
- Bayesian flow

Parametric (Global) Motion Models

Global motion models offer

- more constrained solutions than smoothness (Horn-Schunck)
- integration over a larger area than a translation-only model can accommodate (Lucas-Kanade)

Parametric (Global) Motion Models

2D Models:

(Translation)

Affine

Quadratic

Planar projective transform (Homography)

3D Models:

Instantaneous camera motion models

Homography+epipole

Plane+Parallax

$$E(\mathbf{h}) = \sum_{\mathbf{x} \in \mathbb{R}} [I(\mathbf{x} + \mathbf{h}) - I_0(\mathbf{x})]^2$$

- Transformations/warping of image

$$E(\mathbf{h}) = \sum_{\mathbf{x} \in R} [I(\mathbf{x} + \mathbf{h}) - I_0(\mathbf{x})]^2$$

Translations: $\mathbf{h} = \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$

What about other types of motion?


Generalization

- Transformations/warping of image

$$E(\mathbf{A}, \mathbf{h}) = \sum_{\mathbf{x} \in \mathbb{R}^2} [I(\mathbf{Ax} + \mathbf{h}) - I_0(\mathbf{x})]^2$$

$$\text{Affine:} \quad \mathbf{A} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \quad \mathbf{h} = \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

Generalization



Affine: $\mathbf{A} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \quad \mathbf{h} = \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$

Example: Affine Motion

$$\begin{aligned} u(x, y) &= a_1 + a_2x + a_3y \\ v(x, y) &= a_4 + a_5x + a_6y \end{aligned}$$

Substituting into the B.C. Equation:

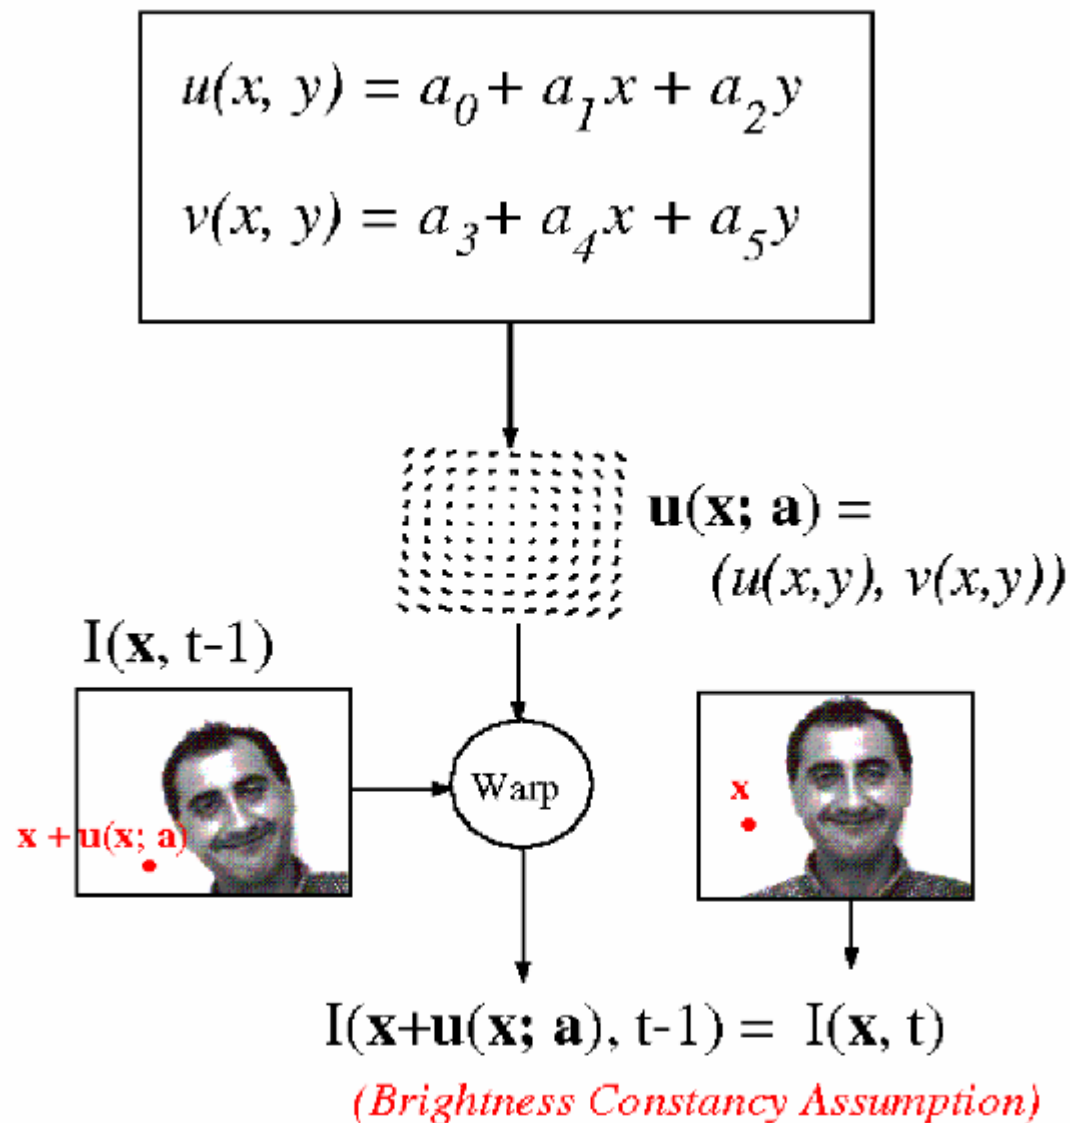
$$I_x \cdot u + I_y \cdot v + I_t \approx 0$$

$$I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \approx 0$$

Each pixel provides 1 linear constraint in 6 *global* unknowns
(*minimum 6 pixels necessary*)

Least Square Minimization (over all pixels):


$$Err(\vec{a}) = \sum \left[I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \right]^2$$



KLT: Good features to keep tracking



Figure 1: Three frame details from Woody Allen's *Manhattan*. The details are from the 1st, 11th, and 21st frames of a subsequence from the movie.



Simple displacement is sufficient between consecutive frames, but not to compare to reference template

Generalization


- Transformations/warping of image

$$E(\mathbf{A}) = \sum_{\mathbf{x} \in R} [I(\mathbf{A} \mathbf{x}) - I_0(\mathbf{x})]^2$$

$$\text{Planar perspective: } \mathbf{A} = \begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & 1 \end{bmatrix}$$

Generalization

Affine +



Planar perspective: $\mathbf{A} = \begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & 1 \end{bmatrix}$





Generalization

- Transformations/warping of image

$$E(\mathbf{h}) = \sum_{\mathbf{x} \in \mathbb{R}} [I(\mathbf{f}(\mathbf{x}, \mathbf{h})) - I_0(\mathbf{x})]^2$$

Other parametrized transformations

Generalization



Other parametrized transformations

2D Motion Models summary

Quadratic – instantaneous approximation to planar motion

$$\begin{aligned}u &= q_1 + q_2x + q_3y + q_7x^2 + q_8xy \\v &= q_4 + q_5x + q_6y + q_7xy + q_8y^2\end{aligned}$$

Projective – exact planar motion


$$\begin{aligned}x' &= \frac{h_1 + h_2x + h_3y}{h_7 + h_8x + h_9y} \\y' &= \frac{h_4 + h_5x + h_6y}{h_7 + h_8x + h_9y}\end{aligned}$$

and

$$u = x' - x, \quad v = y' - y$$

Advanced parametric model

- Optical flow constrained by non-rigid face model



Flexible flow for 3D nonrigid tracking and shape recovery,
Brand and Bhotika, CVPR2001.

3D Motion Models summary

Instantaneous camera motion:

Global parameters: $\Omega_x, \Omega_y, \Omega_z, T_x, T_y, T_z$

Local Parameter: $Z(x, y)$

$$u = -xy\Omega_x + (1+x^2)\Omega_y - y\Omega_z + (T_x - T_zx)/Z$$

$$v = -(1+y^2)\Omega_x + xy\Omega_y - x\Omega_z + (T_y - T_zx)/Z$$

Homography+Epipole

Global parameters: $h_1, \dots, h_9, t_1, t_2, t_3$

Local Parameter: $\gamma(x, y)$

$$x' = \frac{h_1x + h_2y + h_3 + \gamma t_1}{h_7x + h_8y + h_9 + \gamma t_3}$$

$$y' = \frac{h_4x + h_5y + h_6 + \gamma t_1}{h_7x + h_8y + h_9 + \gamma t_3}$$

$$\text{and : } u = x' - x, \quad v = y' - y$$

Residual Planar Parallax Motion


Global parameters: t_1, t_2, t_3

Local Parameter: $\gamma(x, y)$

$$u = x^w - x = \frac{\gamma}{1 + \gamma t_3} (t_3x - t_1)$$

$$v = y^w - x = \frac{\gamma}{1 + \gamma t_3} (t_3y - t_2)$$

Residual Planar Parallax Motion (Plane+Parallax)



Original sequence


Plane-aligned sequence

Recovered shape

Block sequence from [Kumar-Anandan-Hanna'94]

“Given two views where motion of points on a parametric surface has been compensated, the residual parallax is an epipolar field”

Residual Planar Parallax Motion



*The intersection of the two line constraints
uniquely defines the displacement.*




Figure 1: Model-based tracking is robust to degraded images and transient occlusions. Dots show flexed model in 3/4, frontal, and profile view. Dots on face show where the image is sampled. Dots on neck encode 3D motion parameters.



Optical Flow

- Brightness Constancy
- The Aperture problem
- Regularization
- Lucas-Kanade
- Coarse-to-fine
- Parametric motion models
- **SSD tracking**
- Bayesian flow



Correlation and SSD

- For large displacements, do template matching as was used in stereo disparity search.
 - Define a small area around a pixel as the template
 - Match the template against each pixel within a search area in next image.
 - Use a match measure such as correlation, normalized correlation, or sum-of-squares difference
 - Choose the maximum (or minimum) as the match
 - Sub-pixel interpolation also possible



SSD Surface – Textured area



SSD Surface -- Edge



SSD Surface – homogeneous area



Discrete Search vs. Gradient Based Estimation

Consider image I translated by u_0, v_0

$$\begin{aligned} I_0(x, y) &= I(x, y) \\ I_1(x + u_0, y + v_0) &= I(x, y) + \eta_1(x, y) \end{aligned}$$

$$\begin{aligned} E(u, v) &= \sum_{x, y} (I(x, y) - I_1(x + u, y + v))^2 \\ &= \sum_{x, y} (I(x, y) - I(x - u_0 + u, y - v_0 + v) - \eta_1(x, y))^2 \end{aligned}$$

Discrete search simply searches for the best estimate.


Gradient method linearizes the intensity function and solves for the estimate

Optical Flow

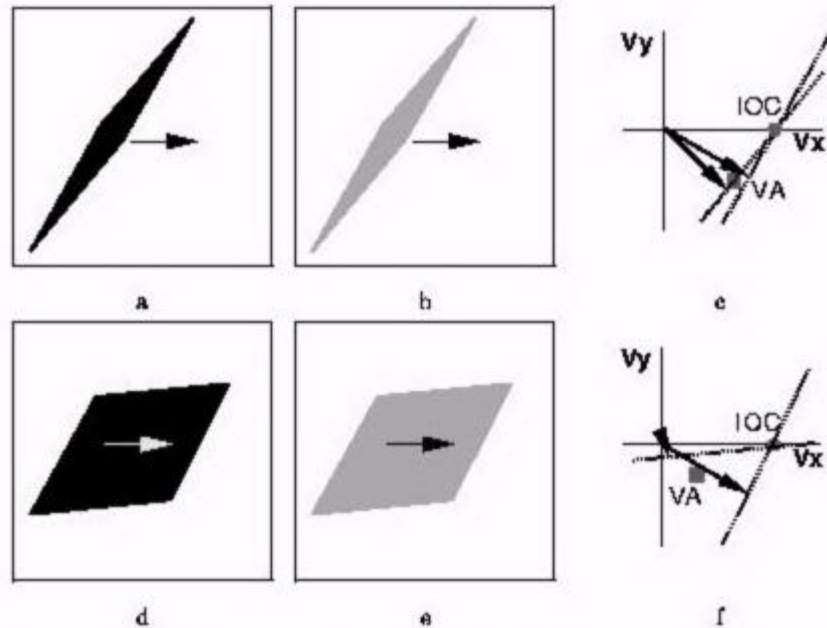
- Brightness Constancy
- The Aperture problem
- Regularization
- Lucas-Kanade
- Coarse-to-fine
- Parametric motion models
- SSD tracking
- Bayesian flow

Bayesian Optic Flow

- Some low-level human motion illusions can be explained by adding an uncertainty model to Lucas-Kanade tracking
- Theories from Psychology about normal flow fusion:
 - (VA) vector average (of normal motions)
 - (IOC) intersection of constraints (e.g., Lucas-Kanade):



Rhombus Displays



<http://www.cs.huji.ac.il/~yweiss/Rhombus/>

Brightness constancy with noise:

$$I(x,y,t) = I(x + v_x \Delta t, y + v_y \Delta t, t + \Delta t) + \eta$$

Assume Gaussian noise, smooth surfaces, locally constant; take first order linear approximation:

$$P(I(x_i, y_i, t) | v_i) \propto \exp \left(-\frac{1}{2\sigma^2} \int_{x,y} w_i(x,y) (I_x(x,y,t)v_x + I_y(x,y,t)v_y + I_t(x,y,t))^2 dx dy \right)$$

Prior favoring slow speeds:

$$P(v) \propto \exp(-\|v\|^2/2\sigma_p^2).$$

Assume noise is independent across location; apply Bayes:

$$P(v|I) \propto P(v) \prod_i P(I(x_i, y_i, t) | v),$$

With constant window $w=1$,

$$P(v|I) \propto \exp \left(-\|v\|^2/2\sigma_p^2 - \frac{1}{2\sigma^2} \int_{x,y} (I_x(x,y)v_x + I_y(x,y)v_y + I_t)^2 dx dy \right)$$

Form 'normal equations' to arrive at....


Lucas-Kanade with uncertainty:

$$v^* = - \begin{pmatrix} \Sigma I_x^2 + \frac{\sigma^2}{\sigma_p^2} & \Sigma I_x I_y \\ \Sigma I_x I_y & \Sigma I_y^2 + \frac{\sigma^2}{\sigma_p^2} \end{pmatrix}^{-1} \begin{pmatrix} \Sigma I_x I_t \\ \Sigma I_y I_t \end{pmatrix}$$

One parameter: ratio of observation and prior gaussian spread.

<http://www.cs.huji.ac.il/~yweiss/Rhombus>

[Weiss, Simoncelli, Adelson Nature Neuroscience 2002]






Figure 4: The response of the Bayesian estimator to a fat rhombus. (replotted from Weiss and Adelson 98)


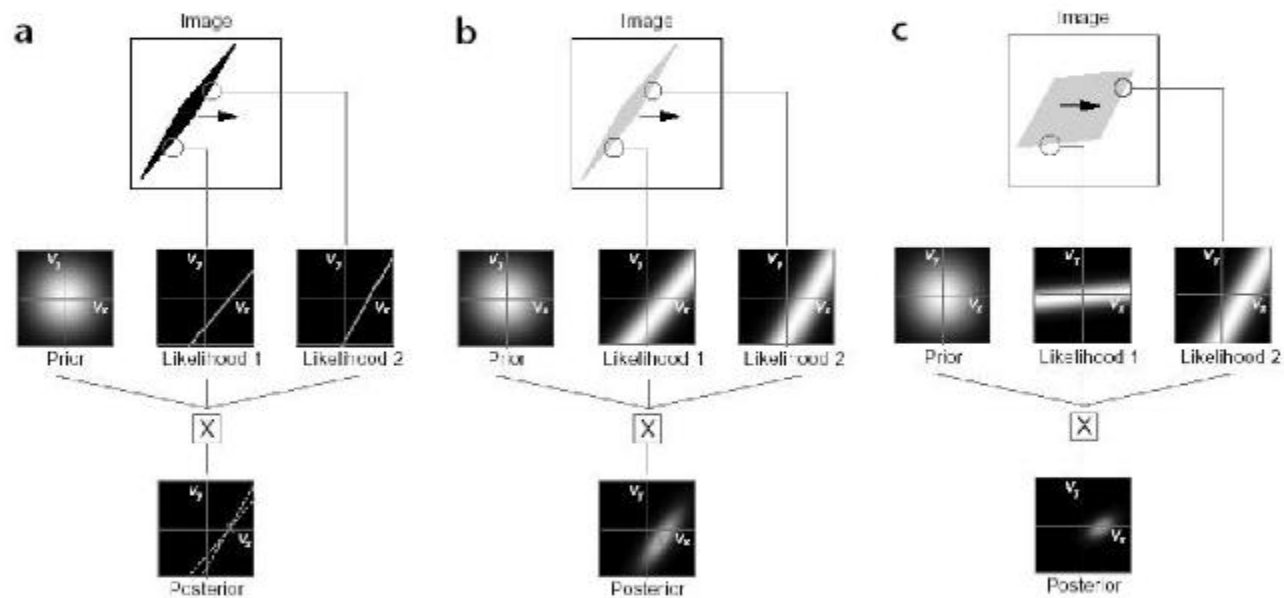


Figure 3: The response of the Bayesian estimator to a narrow rhombus. (replotted from Weiss and Adelson 98)

Effect of contrast



Thursday: video compression

You Tube Broadcast Yourself™
Worldwide | English

Home

Videos

Channels

Community

Videos being watched right now...

